

# model\_batting

April 27, 2020

```
[1]: import tensorflow as tf
import pandas as pd
import numpy as np
```

```
[2]: df = pd.read_csv('../core/output/batters.csv')
indexer = df.reset_index()[['index', 'retroID']].to_dict()['retroID']
y = df['Batting'].values
to_drop = ['debutYear', 'finalYear', 'G', '1B', 'AB', 'RBI', 'wOBA']
df.drop(columns=to_drop, inplace=True)
```

```
[3]: df
```

```
[3]:      retroID  weight  height  pos_1B  pos_2B  pos_3B  pos_C  pos_OF  \
0      aar001  0.569672   0.60      0      0      0      0      0
1      aar010  0.426230   0.45      0      0      0      0      1
2      aar011  0.467213   0.60      1      0      0      0      0
3      aar012  0.467213   0.60      0      0      0      0      0
4      aar013  0.442623   0.50      1      0      0      0      0
...      ...      ...      ...      ...      ...      ...      ...      ...
15288  zup001  0.590164   0.65      0      0      0      0      1
15289  zup010  0.434426   0.40      0      0      0      1      0
15290  zup011  0.487705   0.65      0      0      0      0      0
15291  zup012  0.397541   0.45      0      0      0      0      0
15292  zup013  0.467213   0.60      0      0      0      0      0

      pos_P  pos_SS  ...  SO  IBB  HBP  SH  SF  GIDP  NL  wRC+  WAR  \
0          1      0  ...   2    0    0   1   0    0   1  -100  -0.1
1          0      0  ... 1383  293  32  21 121  328   1   153 136.3
2          0      0  ...  145   3   0   9   6   36   1    76  -1.7
3          1      0  ...    3   0   0   0   0    0   1  -100  -0.1
4          0      0  ...    5   0   0   0   0    1   1    0  -0.4
...      ...      ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
15288      0      0  ...  137   3   6  20   8   15   0    74  -0.9
15289      0      0  ...    6   0   0   0   0    0   0    37  -0.2
15290      1      0  ...   39   0   0  16   0    3   1    0  -0.3
15291      0      1  ...   50   1   2  18   0    8   1   52  -2.2
15292      1      0  ...    0   0   0   0   0    0   0    0   0.0
```

	Batting
0	0.000358
1	0.350195
2	0.157131
3	0.000358
4	0.090001
...	...
15288	0.156062
15289	0.123425
15290	0.090088
15291	0.135118
15292	0.090195

[15293 rows x 31 columns]

## Building the Model

```
[4]: from sklearn.model_selection import train_test_split
```

```
[5]: X = df.drop(columns=['Batting']).values
     y = df[['retroID', 'Batting']].values
```

When we do our train-test split, since it's random in how it splits up the data, we need to keep track of the appropriate keys (retro IDs) for each data point.

```
[6]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=101)
X_train_keys = np.asarray([x[0] for x in X_train])
X_train = np.asarray([x[1:] for x in X_train])
X_test_keys = np.asarray([x[0] for x in X_test])
X_test = np.asarray([x[1:] for x in X_test])
y_train_keys = np.asarray([y[0] for y in y_train])
y_train = np.asarray([y[1] for y in y_train])
y_test_keys = np.asarray([y[0] for y in y_test])
y_test = np.asarray([y[1] for y in y_test])
```

```
[7]: import tensorflow as tf
     from tensorflow.keras.models import Sequential
     from tensorflow.keras.layers import Dense, Dropout
     from tensorflow.keras import regularizers
```

```
[8]: X_train.shape
```

```
[8]: (12234, 29)
```

```
[9]: from sklearn.preprocessing import MinMaxScaler
```

```

[10]: scaler = MinMaxScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)

[11]: def to_tensor_input(player):
      return scaler.transform(player.values.reshape(-1,29))[0]

[12]: tensor = df.drop(columns=['retroID', 'Batting'])
      player_tensor_inputs = tensor.apply(lambda player: to_tensor_input(player),
      ↪axis=1)

[13]: player_tensor_inputs

[13]: 0      [0.5696720000000001, 0.6, 0.0, 0.0, 0.0, 0.0, ...
      1      [0.42623, 0.45, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, ...
      2      [0.467213, 0.6, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
      3      [0.467213, 0.6, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, ...
      4      [0.44262299999999993, 0.5, 1.0, 0.0, 0.0, 0.0, ...
      ...
      15288   [0.590164, 0.65, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, ...
      15289   [0.43442600000000003, 0.4, 0.0, 0.0, 0.0, 1.0, ...
      15290   [0.487705, 0.65, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, ...
      15291   [0.397541, 0.45, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
      15292   [0.467213, 0.6, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, ...
      Length: 15293, dtype: object

[14]: tensor = pd.DataFrame(player_tensor_inputs.values.tolist())

[15]: tensor.to_csv('../core/tensors/t_batting.csv', index=False, float_format='%g')

[16]: epochs = 400
      batch_size = 128
      loss_param = 'mse'
      optimizer_param = 'adam'
      stop_monitor = 'val_loss'
      stop_patience = 6

[17]: model = Sequential()

      model.add(Dense(29, activation='relu', kernel_regularizer=regularizers.l2(0.
      ↪0001)))
      model.add(Dropout(0.5))

      model.add(Dense(58, activation='relu', kernel_regularizer=regularizers.l2(0.
      ↪0001)))
      model.add(Dropout(0.5))

```

```

model.add(Dense(116, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(232, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(116, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(58, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(29, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(14, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(7, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↳0001)))
model.add(Dropout(0.5))

model.add(Dense(units=1, activation='sigmoid'))

model.compile(loss=loss_param, optimizer=optimizer_param)

```

```
[18]: from tensorflow.keras.callbacks import EarlyStopping
```

```
[19]: early_stop = EarlyStopping(monitor=stop_monitor, patience=stop_patience)
```

```
[20]: results = model.fit(x=X_train, y=y_train,
                        epochs=epochs,
                        batch_size=batch_size,
                        validation_data=(X_test, y_test),
                        callbacks=[early_stop]
                    )
```

Train on 12234 samples, validate on 3059 samples

Epoch 1/400

12234/12234 [=====] - 2s 135us/sample - loss: 0.1964 -  
val\_loss: 0.1657

Epoch 2/400  
12234/12234 [=====] - 0s 37us/sample - loss: 0.1357 -  
val\_loss: 0.0755  
Epoch 3/400  
12234/12234 [=====] - 0s 34us/sample - loss: 0.0871 -  
val\_loss: 0.0368  
Epoch 4/400  
12234/12234 [=====] - 0s 33us/sample - loss: 0.0676 -  
val\_loss: 0.0287  
Epoch 5/400  
12234/12234 [=====] - 0s 34us/sample - loss: 0.0567 -  
val\_loss: 0.0233  
Epoch 6/400  
12234/12234 [=====] - 0s 35us/sample - loss: 0.0485 -  
val\_loss: 0.0191  
Epoch 7/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0424 -  
val\_loss: 0.0161  
Epoch 8/400  
12234/12234 [=====] - 1s 42us/sample - loss: 0.0372 -  
val\_loss: 0.0136  
Epoch 9/400  
12234/12234 [=====] - 1s 52us/sample - loss: 0.0328 -  
val\_loss: 0.0120  
Epoch 10/400  
12234/12234 [=====] - 0s 37us/sample - loss: 0.0294 -  
val\_loss: 0.0108  
Epoch 11/400  
12234/12234 [=====] - 0s 33us/sample - loss: 0.0272 -  
val\_loss: 0.0098  
Epoch 12/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0246 -  
val\_loss: 0.0091  
Epoch 13/400  
12234/12234 [=====] - 0s 34us/sample - loss: 0.0229 -  
val\_loss: 0.0087  
Epoch 14/400  
12234/12234 [=====] - 0s 33us/sample - loss: 0.0210 -  
val\_loss: 0.0081  
Epoch 15/400  
12234/12234 [=====] - 0s 34us/sample - loss: 0.0193 -  
val\_loss: 0.0078  
Epoch 16/400  
12234/12234 [=====] - 1s 43us/sample - loss: 0.0178 -  
val\_loss: 0.0076  
Epoch 17/400  
12234/12234 [=====] - 1s 44us/sample - loss: 0.0168 -  
val\_loss: 0.0073

```

Epoch 18/400
12234/12234 [=====] - 1s 58us/sample - loss: 0.0158 -
val_loss: 0.0072
Epoch 19/400
12234/12234 [=====] - 1s 58us/sample - loss: 0.0148 -
val_loss: 0.0070
Epoch 20/400
12234/12234 [=====] - 1s 42us/sample - loss: 0.0138 -
val_loss: 0.0068
Epoch 21/400
12234/12234 [=====] - 1s 47us/sample - loss: 0.0127 -
val_loss: 0.0067
Epoch 22/400
12234/12234 [=====] - 1s 42us/sample - loss: 0.0122 -
val_loss: 0.0066
Epoch 23/400
12234/12234 [=====] - 0s 36us/sample - loss: 0.0114 -
val_loss: 0.0065
Epoch 24/400
12234/12234 [=====] - 1s 51us/sample - loss: 0.0109 -
val_loss: 0.0064
Epoch 25/400
12234/12234 [=====] - 1s 73us/sample - loss: 0.0105 -
val_loss: 0.0064
Epoch 26/400
12234/12234 [=====] - 1s 66us/sample - loss: 0.0098 -
val_loss: 0.0063
Epoch 27/400
12234/12234 [=====] - 1s 58us/sample - loss: 0.0093 -
val_loss: 0.0062
Epoch 28/400
12234/12234 [=====] - 1s 50us/sample - loss: 0.0091 -
val_loss: 0.0062
Epoch 29/400
12234/12234 [=====] - 1s 47us/sample - loss: 0.0086 -
val_loss: 0.0061
Epoch 30/400
12234/12234 [=====] - ETA: 0s - loss: 0.008 - 1s
55us/sample - loss: 0.0083 - val_loss: 0.0061
Epoch 31/400
12234/12234 [=====] - 0s 33us/sample - loss: 0.0082 -
val_loss: 0.0060
Epoch 32/400
12234/12234 [=====] - 1s 55us/sample - loss: 0.0078 -
val_loss: 0.0060
Epoch 33/400
12234/12234 [=====] - 1s 46us/sample - loss: 0.0076 -
val_loss: 0.0059

```

Epoch 34/400  
12234/12234 [=====] - 1s 42us/sample - loss: 0.0074 -  
val\_loss: 0.0059

Epoch 35/400  
12234/12234 [=====] - 1s 46us/sample - loss: 0.0072 -  
val\_loss: 0.0058

Epoch 36/400  
12234/12234 [=====] - 1s 45us/sample - loss: 0.0070 -  
val\_loss: 0.0058

Epoch 37/400  
12234/12234 [=====] - 1s 44us/sample - loss: 0.0070 -  
val\_loss: 0.0058

Epoch 38/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0067 -  
val\_loss: 0.0058

Epoch 39/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0066 -  
val\_loss: 0.0058

Epoch 40/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0065 -  
val\_loss: 0.0058

Epoch 41/400  
12234/12234 [=====] - 0s 32us/sample - loss: 0.0064 -  
val\_loss: 0.0057

Epoch 42/400  
12234/12234 [=====] - 0s 30us/sample - loss: 0.0063 -  
val\_loss: 0.0057

Epoch 43/400  
12234/12234 [=====] - 0s 30us/sample - loss: 0.0062 -  
val\_loss: 0.0057

Epoch 44/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0061 -  
val\_loss: 0.0057

Epoch 45/400  
12234/12234 [=====] - 0s 35us/sample - loss: 0.0061 -  
val\_loss: 0.0057

Epoch 46/400  
12234/12234 [=====] - 1s 55us/sample - loss: 0.0059 -  
val\_loss: 0.0056

Epoch 47/400  
12234/12234 [=====] - 1s 41us/sample - loss: 0.0060 -  
val\_loss: 0.0056

Epoch 48/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0058 -  
val\_loss: 0.0056

Epoch 49/400  
12234/12234 [=====] - 0s 36us/sample - loss: 0.0058 -  
val\_loss: 0.0056

Epoch 50/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0058 -  
val\_loss: 0.0056

Epoch 51/400  
12234/12234 [=====] - 1s 50us/sample - loss: 0.0057 -  
val\_loss: 0.0056

Epoch 52/400  
12234/12234 [=====] - 0s 37us/sample - loss: 0.0058 -  
val\_loss: 0.0056

Epoch 53/400  
12234/12234 [=====] - 0s 36us/sample - loss: 0.0057 -  
val\_loss: 0.0056

Epoch 54/400  
12234/12234 [=====] - 0s 40us/sample - loss: 0.0057 -  
val\_loss: 0.0056

Epoch 55/400  
12234/12234 [=====] - 1s 47us/sample - loss: 0.0057 -  
val\_loss: 0.0056

Epoch 56/400  
12234/12234 [=====] - 1s 43us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 57/400  
12234/12234 [=====] - 1s 60us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 58/400  
12234/12234 [=====] - 1s 75us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 59/400  
12234/12234 [=====] - 1s 64us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 60/400  
12234/12234 [=====] - 1s 63us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 61/400  
12234/12234 [=====] - 1s 63us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 62/400  
12234/12234 [=====] - 1s 68us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 63/400  
12234/12234 [=====] - 1s 66us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 64/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0056 -  
val\_loss: 0.0056

Epoch 65/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0056 -  
val\_loss: 0.0056



Epoch 66/400  
12234/12234 [=====] - 0s 32us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 67/400  
12234/12234 [=====] - 0s 32us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 68/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 69/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 70/400  
12234/12234 [=====] - 0s 32us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 71/400  
12234/12234 [=====] - 0s 31us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 72/400  
12234/12234 [=====] - 0s 33us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 73/400  
12234/12234 [=====] - 0s 35us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 74/400  
12234/12234 [=====] - 1s 44us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 75/400  
12234/12234 [=====] - 1s 42us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 76/400  
12234/12234 [=====] - 0s 40us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 77/400  
12234/12234 [=====] - 0s 38us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 78/400  
12234/12234 [=====] - 0s 40us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 79/400  
12234/12234 [=====] - 1s 51us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 80/400  
12234/12234 [=====] - 0s 35us/sample - loss: 0.0056 -  
val\_loss: 0.0056  
Epoch 81/400  
12234/12234 [=====] - 0s 33us/sample - loss: 0.0056 -  
val\_loss: 0.0056

```
Epoch 82/400
12234/12234 [=====] - 0s 32us/sample - loss: 0.0056 -
val_loss: 0.0056
Epoch 83/400
12234/12234 [=====] - 0s 33us/sample - loss: 0.0056 -
val_loss: 0.0056
```

```
[21]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	multiple	870
dropout (Dropout)	multiple	0
dense_1 (Dense)	multiple	1740
dropout_1 (Dropout)	multiple	0
dense_2 (Dense)	multiple	6844
dropout_2 (Dropout)	multiple	0
dense_3 (Dense)	multiple	27144
dropout_3 (Dropout)	multiple	0
dense_4 (Dense)	multiple	27028
dropout_4 (Dropout)	multiple	0
dense_5 (Dense)	multiple	6786
dropout_5 (Dropout)	multiple	0
dense_6 (Dense)	multiple	1711
dropout_6 (Dropout)	multiple	0
dense_7 (Dense)	multiple	420
dropout_7 (Dropout)	multiple	0
dense_8 (Dense)	multiple	105
dropout_8 (Dropout)	multiple	0

```

-----
dense_9 (Dense)           multiple           8
=====
Total params: 72,656
Trainable params: 72,656
Non-trainable params: 0
-----

```

```
[22]: import os
```

```
[23]: losses = model.history.history
losses['loss'] = np.asarray(losses['loss'])
losses['val_loss'] = np.asarray(losses['val_loss'])
final_number_of_epochs = len(losses['loss'])
min_loss = losses['loss'].min()
mean_loss = losses['loss'].mean()
final_loss = losses['loss'][-1]
min_val_loss = losses['val_loss'].min()
mean_val_loss = losses['val_loss'].mean()
final_val_loss = losses['val_loss'][-1]

def get_model_summary():
    output = []
    model.summary(print_fn=lambda line: output.append(line))
    return str(output).strip('[]')

summary = get_model_summary()

record = {
    'Epochs': final_number_of_epochs,
    'Batch_Size': batch_size,
    'Loss_Func': loss_param,
    'Optimizer': optimizer_param,
    'Early_Stop_Monitor': stop_monitor,
    'Early_Stop_Patience': stop_patience,
    'Min_Loss': min_loss,
    'Mean_Loss': mean_loss,
    'Final_Loss': final_loss,
    'Min_Val_Loss': min_val_loss,
    'Mean_Val_Loss': mean_val_loss,
    'Final_Val_Loss': final_val_loss,
    'Model': summary
}

new_data = pd.DataFrame(record, index=[0])

if os.path.exists('../core/records/batting_results.csv'):

```

```

df_records = pd.read_csv('../core/records/batting_results.csv')
df_records.append(new_data)
else:
    df_records = pd.DataFrame(new_data)

df_records.to_csv('../core/records/batting_results.csv', index=False,
    ↪float_format='%g')


```

## Model Evaluation

```
[24]: losses = pd.DataFrame(model.history.history)
```

```
[25]: losses.plot()
```

```
[25]: <matplotlib.axes._subplots.AxesSubplot at 0x14011ce10>
```



```
[26]: predictions = model.predict(X_test)
predictions = [pred for sublist in predictions for pred in sublist]
```

```
[27]: test_player_ratings = dict(zip(X_test_keys, predictions))
```

```
[28]: player_key = df['retroID']
```

```
[29]: player_key
```

```
[29]: 0      aar001
      1      aar0101
      2      aar0101
      3      aar001
      4      abada001
      ...
      15288   zupcb001
      15289   zupof101
      15290   zuveg101
      15291   zuveg001
      15292   zyht001
      Name: retroID, Length: 15293, dtype: object
```

```
[30]: results = model.predict(tensor.to_numpy())
```

```
[31]: len(results)
```

```
[31]: 15293
```

```
[32]: results.mean()
```

```
[32]: 0.12198973
```

```
[33]: df['Batting'].shape
```

```
[33]: (15293,)
```

```
[34]: results.shape
```


```
[34]: (15293, 1)
```

```
[35]: results = [pred for sublist in results for pred in sublist]
```

```
[36]: diff = df['Batting'] - results
```

```
[37]: diff.plot()
```

```
[37]: <matplotlib.axes._subplots.AxesSubplot at 0x140c786d0>
```



```
[38]: diff.mean()
```

```
[38]: -0.00040076900953220845
```

```
[ ]:
```