

CNN

May 17, 2020

```
[1]: import os
import pandas as pd
import numpy as np
```

Test on a single year to ensure the method works.

```
[2]: df = pd.read_csv('../core/tensors/games/2015.csv', header=None)
```

```
[3]: df
```

```
[3]:
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 531 \ |
|------|----------|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|----------|
| 0 | 0.446721 | 0.40 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... | 0.136311 |
| 1 | 0.467213 | 0.40 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | ... | 0.212168 |
| 2 | 0.409836 | 0.30 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | ... | 0.298806 |
| 3 | 0.487705 | 0.45 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | ... | 0.185214 |
| 4 | 0.508197 | 0.55 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | ... | 0.112437 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2423 | 0.500000 | 0.50 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... | 0.001155 |
| 2424 | 0.508197 | 0.50 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.041201 |
| 2425 | 0.385246 | 0.40 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | ... | 0.006931 |
| 2426 | 0.487705 | 0.30 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... | 0.033115 |
| 2427 | 0.590164 | 0.60 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | ... | 0.118598 |

| | 532 | 533 | 534 | 535 | 536 | 537 | 538 \ |
|------|----------|----------|----------|----------|----------|-----|----------|
| 0 | 0.008721 | 0.140351 | 0.089219 | 0.101562 | 0.164557 | 1.0 | 0.164568 |
| 1 | 0.007267 | 0.098246 | 0.070632 | 0.085938 | 0.065823 | 1.0 | 0.160971 |
| 2 | 0.015988 | 0.049123 | 0.237918 | 0.414062 | 0.225316 | 1.0 | 0.167266 |
| 3 | 0.021802 | 0.129825 | 0.052045 | 0.250000 | 0.096203 | 1.0 | 0.178058 |
| 4 | 0.011628 | 0.056140 | 0.070632 | 0.085938 | 0.091139 | 1.0 | 0.163669 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2423 | 0.000000 | 0.000000 | 0.007435 | 0.000000 | 0.000000 | 1.0 | 0.103417 |
| 2424 | 0.000000 | 0.000000 | 0.085502 | 0.000000 | 0.020253 | 1.0 | 0.114209 |
| 2425 | 0.000000 | 0.000000 | 0.022305 | 0.000000 | 0.000000 | 1.0 | 0.107014 |
| 2426 | 0.000000 | 0.000000 | 0.115242 | 0.000000 | 0.002532 | 1.0 | 0.064748 |
| 2427 | 0.000000 | 0.003509 | 0.271375 | 0.015625 | 0.002532 | 1.0 | 0.075540 |

| | 539 | 540 |
|---|----------|-----|
| 0 | 0.116185 | 1.0 |

```

1      0.084971  1.0
2      0.124277  1.0
3      0.150867  0.0
4      0.076879  1.0
...      ...      ...
2423   0.049711  0.0
2424   0.062428  1.0
2425   0.051445  1.0
2426   0.046243  1.0
2427   0.047399  0.0

```

```
[2428 rows x 541 columns]
```

```
[4]: game0 = list(df.iloc[0])
```

```
[5]: len(game0)
```

```
[5]: 541
```

```
[6]: y0 = int(game0[-1])
```

```
[7]: y0
```

```
[7]: 1
```

```
[8]: type(game0)
```

```
[8]: list
```

```
[9]: game0 = game0[:-1]
```

```
[10]: game0 = np.reshape(game0, (18, 30, 1))
```

```
[11]: game0.shape
```

```
[11]: (18, 30, 1)
```

```
[12]: game1 = list(df.iloc[1])
```

```
[13]: y1 = int(game1[-1])
```

```
[14]: game1 = game1[:-1]
```

```
[15]: game1 = np.reshape(game1, (18, 30, 1))
```

```
[16]: games = [game0, game1]
```

```
[17]: rows = len(games)
      cols = len(games[0])
      fors = len(games[0][0])
      last = len(games[0][0][0])
```

```
[18]: (rows, cols, fors, last)
```

```
[18]: (2, 18, 30, 1)
```

```
[19]: game2 = list(df.iloc[2])
      y2 = int(game2[-1])
      game2 = game2[:-1]
      game2 = np.reshape(game2, (18, 30))
```

```
[20]: games.append(game2)
```

```
[21]: rows = len(games)
      cols = len(games[0])
      fors = len(games[0][0])
      (rows, cols, fors)
```

```
[21]: (3, 18, 30)
```

We've shown that this is an effective method of getting games as 3D tables from the .csv files, as well as combining them into a list of tables. We can scale this to get a list of every game, and this will be our CNN input.

```
[22]: df.shape
```

```
[22]: (2428, 541)
```

```
[23]: X = []
      y = []
      for index in range(0, df.shape[0]):
          game = list(df.iloc[index])
          y_sample = int(game[-1])
          game = game[:-1]
          game = np.reshape(game, ([18, 30, 1]))
          X.append(game)
          y.append(y_sample)
      X = np.array(X)
      y = np.array(y)
```

```
[24]: y = y[... , np.newaxis]
```

```
[25]: X.shape
```

```
[25]: (2428, 18, 30, 1)
```

```
[26]: y.shape
```

```
[26]: (2428, 1)
```

```
[27]: rows = len(games)
      cols = len(games[0])
      fors = len(games[0][0])
      last = len(games[0][0][0])
      (rows, cols, fors, last)
```

```
[27]: (3, 18, 30, 1)
```

```
[28]: len(y)
```

```
[28]: 2428
```

```
[29]: input_shape = (18, 30, 1)
```

Now we can extrapolate this to pull all games from all years into one list. This will be the last step before running our predictive model.

```
[30]: X = []
      y = []
      for year in range(1919, 2020):
          df = pd.read_csv('../core/tensors/games/{}.csv'.format(year), header=None)
          for index in range(0, df.shape[0]):
              game = list(df.iloc[index])
              y_sample = int(game[-1])
              game = game[:-1]
              game = np.reshape(game, ([18, 30, 1]))
              X.append(game)
              y.append(y_sample)
```

```
[31]: X = np.array(X)
      y = np.array(y)
```

```
[32]: y = y[..., np.newaxis]
```

```
[33]: rows = len(X)
      cols = len(X[0])
      fors = len(X[0][0])
      last = len(X[0][0][0])
      (rows, cols, fors, last)
```

```
[33]: (176687, 18, 30, 1)
```

```
[34]: len(y)
```

```
[34]: 176687
```

```
[35]: from sklearn.model_selection import train_test_split
```

```
[36]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
↳ random_state=42)
```

```
[37]: len(X_train)
```

```
[37]: 141349
```

```
[38]: X_train.shape
```

```
[38]: (141349, 18, 30, 1)
```

```
[39]: len(X_test)
```

```
[39]: 35338
```

```
[40]: len(y_train)
```

```
[40]: 141349
```

```
[41]: y_train.shape
```

```
[41]: (141349, 1)
```

```
[42]: len(y_test)
```

```
[42]: 35338
```

```
[43]: y_train
```

```
[43]: array([[0],  
        [0],  
        [0],  
        ...,  
        [0],  
        [1],  
        [1]])
```

Building the CNN

```
[44]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Dropout, Flatten
```

```
[45]: image_shape=(18, 30, 1)
```

```
[46]: epochs = 4000
batch_size = 16
loss_param = 'binary_crossentropy'
optimizer_param = 'adam'
stop_monitor = 'val_loss'
metric = 'accuracy'
stop_patience = 20
```

```
[47]: model = Sequential()

model.add(Conv2D(filters=32, kernel_size=(1, 1),
                 input_shape=image_shape, activation='tanh'))
model.add(MaxPool2D(pool_size=(1, 1)))

model.add(Conv2D(filters=32, kernel_size=(3, 3),
                 input_shape=image_shape, activation='tanh'))
model.add(MaxPool2D(pool_size=(3, 3)))
model.add(Dropout(0.5))

model.add(Conv2D(filters=64, kernel_size=(3, 3),
                 activation='relu', padding="same"))
model.add(MaxPool2D(pool_size=(2, 2)))

model.add(Conv2D(filters=128, kernel_size=(3, 3),
                 activation='relu', padding="same"))
model.add(MaxPool2D(pool_size=(2, 2)))
model.add(Dropout(0.5))

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))

model.add(Dense(1, activation='sigmoid'))

model.compile(loss=loss_param, optimizer=optimizer_param,
              metrics=[metric])
```

```
[48]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------------|---------|
| conv2d (Conv2D) | (None, 18, 30, 32) | 64 |

```

-----
max_pooling2d (MaxPooling2D) (None, 18, 30, 32)      0
-----
conv2d_1 (Conv2D)          (None, 16, 28, 32)      9248
-----
max_pooling2d_1 (MaxPooling2 (None, 5, 9, 32)      0
-----
dropout (Dropout)          (None, 5, 9, 32)      0
-----
conv2d_2 (Conv2D)          (None, 5, 9, 64)      18496
-----
max_pooling2d_2 (MaxPooling2 (None, 2, 4, 64)      0
-----
conv2d_3 (Conv2D)          (None, 2, 4, 128)      73856
-----
max_pooling2d_3 (MaxPooling2 (None, 1, 2, 128)      0
-----
dropout_1 (Dropout)        (None, 1, 2, 128)      0
-----
flatten (Flatten)          (None, 256)          0
-----
dense (Dense)              (None, 256)          65792
-----
dropout_2 (Dropout)        (None, 256)          0
-----
dense_1 (Dense)            (None, 1)            257
=====
Total params: 167,713
Trainable params: 167,713
Non-trainable params: 0
-----

```

```
[49]: from tensorflow.keras.callbacks import EarlyStopping
```

```
[50]: early_stop = EarlyStopping(monitor=stop_monitor, patience=stop_patience)
```

```
[51]: results = model.fit(X_train, y_train, epochs=epochs,
                           validation_data=(X_test, y_test),
                           callbacks=[early_stop]
                           )
```

Train on 141349 samples, validate on 35338 samples

Epoch 1/4000

141349/141349 [=====] - 189s 1ms/sample - loss: 0.6857
- accuracy: 0.5518 - val_loss: 0.6808 - val_accuracy: 0.5633

Epoch 2/4000

141349/141349 [=====] - 200s 1ms/sample - loss: 0.6839
- accuracy: 0.5568 - val_loss: 0.6831 - val_accuracy: 0.5560

Epoch 3/4000
141349/141349 [=====] - 195s 1ms/sample - loss: 0.6835
- accuracy: 0.5562 - val_loss: 0.6825 - val_accuracy: 0.5595

Epoch 4/4000
141349/141349 [=====] - 179s 1ms/sample - loss: 0.6836
- accuracy: 0.5566 - val_loss: 0.6811 - val_accuracy: 0.5591

Epoch 5/4000
141349/141349 [=====] - 177s 1ms/sample - loss: 0.6834
- accuracy: 0.5576 - val_loss: 0.6816 - val_accuracy: 0.5570

Epoch 6/4000
141349/141349 [=====] - 183s 1ms/sample - loss: 0.6835
- accuracy: 0.5579 - val_loss: 0.6823 - val_accuracy: 0.5572

Epoch 7/4000
141349/141349 [=====] - 180s 1ms/sample - loss: 0.6832
- accuracy: 0.5580 - val_loss: 0.6825 - val_accuracy: 0.5627

Epoch 8/4000
141349/141349 [=====] - 184s 1ms/sample - loss: 0.6831
- accuracy: 0.5585 - val_loss: 0.6816 - val_accuracy: 0.5615

Epoch 9/4000
141349/141349 [=====] - 180s 1ms/sample - loss: 0.6832
- accuracy: 0.5582 - val_loss: 0.6836 - val_accuracy: 0.5626

Epoch 10/4000
141349/141349 [=====] - 186s 1ms/sample - loss: 0.6832
- accuracy: 0.5589 - val_loss: 0.6819 - val_accuracy: 0.5623

Epoch 11/4000
141349/141349 [=====] - 183s 1ms/sample - loss: 0.6829
- accuracy: 0.5583 - val_loss: 0.6825 - val_accuracy: 0.5622

Epoch 12/4000
141349/141349 [=====] - 186s 1ms/sample - loss: 0.6829
- accuracy: 0.5588 - val_loss: 0.6813 - val_accuracy: 0.5596

Epoch 13/4000
141349/141349 [=====] - 187s 1ms/sample - loss: 0.6829
- accuracy: 0.5566 - val_loss: 0.6816 - val_accuracy: 0.5620

Epoch 14/4000
141349/141349 [=====] - 186s 1ms/sample - loss: 0.6830
- accuracy: 0.5577 - val_loss: 0.6818 - val_accuracy: 0.5612

Epoch 15/4000
141349/141349 [=====] - 178s 1ms/sample - loss: 0.6830
- accuracy: 0.5586 - val_loss: 0.6831 - val_accuracy: 0.5614

Epoch 16/4000
141349/141349 [=====] - 178s 1ms/sample - loss: 0.6828
- accuracy: 0.5584 - val_loss: 0.6828 - val_accuracy: 0.5634

Epoch 17/4000
141349/141349 [=====] - 180s 1ms/sample - loss: 0.6829
- accuracy: 0.5576 - val_loss: 0.6812 - val_accuracy: 0.5614

Epoch 18/4000
141349/141349 [=====] - 177s 1ms/sample - loss: 0.6828
- accuracy: 0.5585 - val_loss: 0.6817 - val_accuracy: 0.5595


```

Epoch 19/4000
141349/141349 [=====] - 180s 1ms/sample - loss: 0.6829
- accuracy: 0.5584 - val_loss: 0.6834 - val_accuracy: 0.5610
Epoch 20/4000
141349/141349 [=====] - 181s 1ms/sample - loss: 0.6829
- accuracy: 0.5584 - val_loss: 0.6830 - val_accuracy: 0.5531
Epoch 21/4000
141349/141349 [=====] - 178s 1ms/sample - loss: 0.6829
- accuracy: 0.5593 - val_loss: 0.6813 - val_accuracy: 0.5588

```

```

[52]: losses = model.history.history
losses['loss'] = np.asarray(losses['loss'])
losses['val_loss'] = np.asarray(losses['val_loss'])
final_number_of_epochs = len(losses['loss'])
min_loss = losses['loss'].min()
mean_loss = losses['loss'].mean()
final_loss = losses['loss'][-1]
min_val_loss = losses['val_loss'].min()
mean_val_loss = losses['val_loss'].mean()
final_val_loss = losses['val_loss'][-1]

def get_model_summary():
    output = []
    model.summary(print_fn=lambda line: output.append(line))
    return str(output).strip('[]')

summary = get_model_summary()

record = {
    'Epochs': final_number_of_epochs,
    'Batch_Size': batch_size,
    'Loss_Func': loss_param,
    'Optimizer': optimizer_param,
    'Early_Stop_Monitor': stop_monitor,
    'Early_Stop_Patience': stop_patience,
    'Min_Loss': min_loss,
    'Mean_Loss': mean_loss,
    'Final_Loss': final_loss,
    'Min_Val_Loss': min_val_loss,
    'Mean_Val_Loss': mean_val_loss,
    'Final_Val_Loss': final_val_loss,
    'Model': summary
}

new_data = pd.DataFrame(record, index=[0])

```

```

if os.path.exists('../core/records/game_predictions.csv'):
    df_records = pd.read_csv('../core/records/game_predictions.csv')
    df_records = df_records.append(new_data)
else:
    df_records = pd.DataFrame(new_data)

df_records.to_csv('../core/records/game_predictions.csv',
                  index=False, float_format='%g')

model.save('../core/models/model_games.h5')


```

Model Evaluation

```
[53]: losses = pd.DataFrame(model.history.history)
```

```
[54]: losses.plot()
```

```
[54]: <matplotlib.axes._subplots.AxesSubplot at 0x197e53250>
```



```
[ ]:
```