

batting_pre

March 9, 2020

```
[120]: import math
import numpy as np
import pandas as pd

# We're going to be reassigning some columns, so we'll turn off this warning -␣
→we know what we're doing!
pd.options.mode.chained_assignment = None # default='warn'
```

```
[121]: # This will be exported to a separate module
ids = pd.read_csv('../data/lahman/mlb_data/People.csv')
ids = ids[['playerID', 'retroID']]
id_dict = ids.set_index('playerID').to_dict()['retroID']

def get_retroid(id):
    return id_dict[id] if id_dict is not None else id
```

```
[122]: df = pd.read_csv('../data/lahman/mlb_data/Batting.csv').sort_values('playerID')
```

```
[123]: df['playerID'] = df['playerID'].apply(get_retroid)
```

```
[124]: df.rename(columns={'playerID': 'retroID'}, inplace=True)
```

```
[125]: df[df['retroID'] == None]
```

```
[125]: Empty DataFrame
Columns: [retroID, yearID, stint, teamID, lgID, G, AB, R, H, 2B, 3B, HR, RBI,
SB, CS, BB, SO, IBB, HBP, SH, SF, GIDP]
Index: []

[0 rows x 22 columns]
```

Cleaning the Data - Missing Values

Print percentages of missing data in each column of the batting table

```
[126]: 100 * df.isnull().sum() / len(df)
```

```
[126]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID     0.000000
      lgID       0.000000
      G          0.000000
      AB         0.000000
      R          0.000000
      H          0.000000
      2B         0.000000
      3B         0.000000
      HR         0.000000
      RBI        0.000000
      SB         0.000000
      CS         8.221708
      BB         0.000000
      SO         0.000000
      IBB        21.711883
      HBP        0.000000
      SH         0.000000
      SF         21.090864
      GIDP       9.839985
      dtype: float64
```

Since this data is by season, it's likely that we have entries for a player for one season with no data in these fields but there is data for other seasons. Since we're taking aggregate sums for each player, we have two options: set these null values to zero so they don't add to the sum, or set them to the average for that player. We'll have to test the theory to see which is more viable.

We're going to start with IBB rather than CS, since it's a more significant chunk of the dataset.

Handling missing IBB data

```
[127]: df[(df['IBB'].isnull())]
```

```
[127]:      retroID  yearID  stint  teamID  lgID   G  AB  R  H  2B  ...  RBI  \
19269  aaroh101   1954      1    ML1    NL  122  468  58  131  27  ...   69
18684  abera101   1953      2    DET    AL   17   23   2   3   0  ...    2
16858  abera101   1950      1    CLE    AL    1    2   0   0   0  ...    0
19270  abera101   1954      1    DET    AL   32   39   3   5   0  ...    3
18683  abera101   1953      1    CLE    AL    6    0   0   0   0  ...    0
...      ...      ...      ...      ...      ...  ...  ...  ..  ...  ..  ...  ...
15128  zubeb101   1946      1    NYA    AL    3    2   0   0   0  ...    0
14447  zubeb101   1945      1    NYA    AL   21   42   1   7   0  ...    3
13868  zubeb101   1944      1    NYA    AL   22   31   1   4   0  ...    1
19843  zuveg101   1954      1    CIN    NL    2    2   1   1   0  ...    0
19844  zuveg101   1954      2    DET    AL   35   64   1   8   1  ...    3
```

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 15128 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |
| 19844 | 0 | 1.0 | 1 | 14 | NaN | 0 | 9 | 0.0 | 2.0 |

[19159 rows x 22 columns]

```
[128]: df[(df['retroID'] == 'abera101')]
```

```
[128]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|-------|----------|--------|-------|--------|------|----|----|---|---|----|-----|-----|----|---|
| 19846 | abera101 | 1955 | 1 | DET | AL | 39 | 17 | 0 | 1 | 0 | ... | 0 | 0 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | 0 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | 0 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | 0 | |
| 21123 | abera101 | 1957 | 2 | KC1 | AL | 3 | 1 | 0 | 1 | 0 | ... | 0 | 0 | |
| 20501 | abera101 | 1956 | 1 | DET | AL | 42 | 10 | 0 | 3 | 0 | ... | 0 | 0 | |
| 21122 | abera101 | 1957 | 1 | DET | AL | 28 | 8 | 0 | 1 | 0 | ... | 1 | 0 | |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|----|----|-----|-----|----|-----|------|
| 19846 | 0.0 | 0 | 9 | 0.0 | 0 | 2 | 0.0 | 1.0 |
| 18684 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| 21123 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 20501 | 0.0 | 1 | 4 | 0.0 | 0 | 2 | 0.0 | 0.0 |
| 21122 | 0.0 | 1 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 |

[8 rows x 22 columns]

```
[129]: df[(df['retroID'] == 'zubeb101')]
```

```
[129]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|-------|----------|--------|-------|--------|------|----|----|---|---|----|-----|-----|----|---|
| 9445 | zubeb101 | 1936 | 1 | CLE | AL | 2 | 5 | 1 | 1 | 0 | ... | 0 | 0 | |
| 10501 | zubeb101 | 1938 | 1 | CLE | AL | 15 | 7 | 0 | 0 | 0 | ... | 0 | 0 | |
| 11080 | zubeb101 | 1939 | 1 | CLE | AL | 16 | 5 | 0 | 1 | 0 | ... | 0 | 0 | |
| 11621 | zubeb101 | 1940 | 1 | CLE | AL | 17 | 3 | 0 | 1 | 0 | ... | 0 | 0 | |
| 12203 | zubeb101 | 1941 | 1 | WS1 | AL | 36 | 26 | 0 | 0 | 0 | ... | 0 | 0 | |

| | | | | | | | | | | | | | |
|-------|----------|------|---|-----|----|----|----|---|---|---|-----|---|---|
| 12742 | zubeb101 | 1942 | 1 | WS1 | AL | 37 | 39 | 5 | 6 | 3 | ... | 3 | 0 |
| 13299 | zubeb101 | 1943 | 1 | NYA | AL | 20 | 38 | 1 | 7 | 1 | ... | 2 | 0 |
| 15711 | zubeb101 | 1947 | 1 | BOS | AL | 20 | 13 | 0 | 2 | 0 | ... | 0 | 0 |
| 15129 | zubeb101 | 1946 | 2 | BOS | AL | 15 | 18 | 1 | 2 | 0 | ... | 2 | 0 |
| 15128 | zubeb101 | 1946 | 1 | NYA | AL | 3 | 2 | 0 | 0 | 0 | ... | 0 | 0 |
| 14447 | zubeb101 | 1945 | 1 | NYA | AL | 21 | 42 | 1 | 7 | 0 | ... | 3 | 0 |
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 | 0 |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|----|----|-----|-----|----|-----|------|
| 9445 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | NaN |
| 10501 | 0.0 | 0 | 1 | NaN | 0 | 1 | NaN | NaN |
| 11080 | 0.0 | 0 | 2 | NaN | 0 | 0 | NaN | 0.0 |
| 11621 | 0.0 | 0 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| 12203 | 0.0 | 1 | 8 | NaN | 0 | 2 | NaN | 0.0 |
| 12742 | 0.0 | 1 | 7 | NaN | 0 | 3 | NaN | 2.0 |
| 13299 | 0.0 | 4 | 14 | NaN | 0 | 5 | NaN | 2.0 |
| 15711 | 0.0 | 2 | 3 | NaN | 0 | 2 | NaN | 2.0 |
| 15129 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 15128 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |

[12 rows x 22 columns]

First let's look at IBB, intentional bases on balls. It seems like most of the missing data is from early in the dataset - it could be that IBB was not recorded then, and/or not considered a trackable play?

```
[130]: df[(df['IBB'].isnull())]['yearID'].max()
```

```
[130]: 1954
```

```
[131]: df[(df['IBB'].isnull())]['yearID'].min()
```

```
[131]: 1919
```

```
[132]: df[(df['IBB'].isnull())]
```

```
[132]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|-----|-----|-----|-----|-----|---|
| 19269 | aaroh101 | 1954 | 1 | ML1 | NL | 122 | 468 | 58 | 131 | 27 | ... | 69 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15128 | zubeb101 | 1946 | 1 | NYA | AL | 3 | 2 | 0 | 0 | 0 | ... | 0 | |
| 14447 | zubeb101 | 1945 | 1 | NYA | AL | 21 | 42 | 1 | 7 | 0 | ... | 3 | |

| | | | | | | | | | | | | |
|-------|----------|------|---|-----|----|----|----|---|---|---|-----|---|
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 |
| 19843 | zuveg101 | 1954 | 1 | CIN | NL | 2 | 2 | 1 | 1 | 0 | ... | 0 |
| 19844 | zuveg101 | 1954 | 2 | DET | AL | 35 | 64 | 1 | 8 | 1 | ... | 3 |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 15128 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |
| 19844 | 0 | 1.0 | 1 | 14 | NaN | 0 | 9 | 0.0 | 2.0 |

[19159 rows x 22 columns]

We have 19159 total rows where there is no data for IBB, and we know none of those rows goes past the year 1954...

```
[133]: df[(df['yearID'] < 1955)]
```

```
[133]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|----|-----|----|-----|-----|---|
| 19269 | aaroh101 | 1954 | 1 | ML1 | NL | 122 | 468 | 58 | 131 | 27 | ... | 69 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | .. | ... | .. | ... | ... | |
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 | |
| 18682 | zuveg101 | 1952 | 1 | CLE | AL | 2 | 0 | 1 | 0 | 0 | ... | 0 | |
| 19843 | zuveg101 | 1954 | 1 | CIN | NL | 2 | 2 | 1 | 1 | 0 | ... | 0 | |
| 19844 | zuveg101 | 1954 | 2 | DET | AL | 35 | 64 | 1 | 8 | 1 | ... | 3 | |
| 18050 | zuveg101 | 1951 | 1 | CLE | AL | 16 | 0 | 0 | 0 | 0 | ... | 0 | |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 18682 | 0 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |

```
19844    0  1.0   1  14  NaN    0   9  0.0   2.0
18050    0  0.0   0   0  0.0    0   0  0.0   0.0
```

```
[19845 rows x 22 columns]
```

And we have 19845 total rows up to the year 1954. That means...

```
[134]: 19159 / 19845
```

```
[134]: 0.9654320987654321
```

Over 96% of the data before 1955 is missing IBB. I think this gives justification to just setting all of those NaNs to 0.

```
[135]: df['IBB'].fillna(value=0, inplace=True)
```

```
[136]: 100 * df.isnull().sum() / len(df)
```

```
[136]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID      0.000000
      lgID        0.000000
      G           0.000000
      AB          0.000000
      R           0.000000
      H           0.000000
      2B          0.000000
      3B          0.000000
      HR          0.000000
      RBI         0.000000
      SB          0.000000
      CS          8.221708
      BB          0.000000
      SO          0.000000
      IBB         0.000000
      HBP         0.000000
      SH          0.000000
      SF          21.090864
      GIDP        9.839985
      dtype: float64
```

Our IBB issue is solved. Let's move on to SF (sacrifice flies). We'll check the years and rows again to see if we're justified in using the same method to eliminate nulls.

Handling missing SF data

```
[137]: df[(df['SF'].isnull())]['yearID'].max()
```

```
[137]: 1953
```

```
[138]: df[(df['SF'].isnull())]['yearID'].min()
```

```
[138]: 1919
```

```
[139]: df[(df['SF'].isnull())].shape[0]
```

```
[139]: 18611
```

```
[140]: df[(df['yearID'] < 1954)].shape[0]
```

```
[140]: 19269
```

```
[141]: 18611/19269
```

```
[141]: 0.965851886449738
```

Almost the same percentage, and one less year covered. I think we can fill those missing values with 0.

```
[142]: df['SF'].fillna(value=0, inplace=True)
```

```
[143]: 100 * df.isnull().sum() / len(df)
```

```
[143]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID      0.000000
      lgID        0.000000
      G           0.000000
      AB          0.000000
      R           0.000000
      H           0.000000
      2B          0.000000
      3B          0.000000
      HR          0.000000
      RBI         0.000000
      SB          0.000000
      CS          8.221708
      BB          0.000000
      SO          0.000000
      IBB         0.000000
      HBP         0.000000
      SH          0.000000
      SF          0.000000
      GIDP        9.839985
      dtype: float64
```

Two more to go, let's move on to CS (caught stealing)

Handling missing CS data

```
[144]: df[(df['CS'].isnull())]['yearID'].max()
```

```
[144]: 1950
```

```
[145]: df[(df['CS'].isnull())]['yearID'].min()
```

```
[145]: 1919
```

```
[146]: df[(df['CS'].isnull())]
```

```
[146]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|-----|------|-----|-----|-----|---|
| 14448 | abewr101 | 1946 | 1 | NY1 | NL | 15 | 8 | 0 | 0 | 0 | ... | 0 | |
| 16285 | aberc101 | 1949 | 1 | CHN | NL | 4 | 7 | 0 | 0 | 0 | ... | 0 | |
| 15712 | aberc101 | 1948 | 1 | CHN | NL | 12 | 32 | 1 | 6 | 1 | ... | 6 | |
| 15131 | aberc101 | 1947 | 1 | CHN | NL | 47 | 140 | 24 | 39 | 6 | ... | 20 | |
| 16286 | abrac101 | 1949 | 1 | BRO | NL | 8 | 24 | 6 | 2 | 1 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | |
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | |
| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP | | | | |
| 14448 | 0 | NaN | 0 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 | | | | |
| 16285 | 0 | NaN | 0 | 2 | 0.0 | 0 | 0 | 0.0 | 1.0 | | | | |
| 15712 | 0 | NaN | 5 | 10 | 0.0 | 0 | 0 | 0.0 | 0.0 | | | | |
| 15131 | 0 | NaN | 20 | 32 | 0.0 | 0 | 0 | 0.0 | 5.0 | | | | |
| 16286 | 1 | NaN | 7 | 6 | 0.0 | 0 | 0 | 0.0 | 1.0 | | | | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | | | | |
| 532 | 2 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN | | | | |
| 4782 | 9 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN | | | | |
| 5312 | 13 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN | | | | |
| 533 | 0 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN | | | | |
| 5842 | 4 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN | | | | |

[7255 rows x 22 columns]

```
[147]: df[(df['retroID'] == 'zitzb101')]
```

```
[147]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|------|----------|--------|-------|--------|------|-----|-----|----|----|----|-----|-----|----|---|
| 4241 | zitzb101 | 1926 | 1 | CIN | NL | 53 | 94 | 21 | 23 | 2 | ... | 3 | 3 | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | 2 | |
| 3715 | zitzb101 | 1925 | 1 | CIN | NL | 104 | 301 | 53 | 76 | 13 | ... | 21 | 11 | |

| | | | | | | | | | | | | | |
|------|----------|------|---|-----|----|-----|-----|----|----|----|-----|----|----|
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | 9 |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | 13 |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | 0 |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | 4 |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|------|------|----|----|-----|-----|----|-----|------|
| 4241 | NaN | 6 | 7 | 0.0 | 2 | 3 | 0.0 | NaN |
| 532 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN |
| 3715 | 11.0 | 35 | 22 | 0.0 | 6 | 2 | 0.0 | NaN |
| 4782 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN |
| 5312 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN |
| 533 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN |
| 5842 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN |

[7 rows x 22 columns]

```
[148]: df[(df['CS'].isnull())].shape[0]
```

```
[148]: 7255
```

```
[149]: df[(df['yearID'] < 1951)].shape[0]
```

```
[149]: 17435
```

```
[150]: 7255/17435
```

```
[150]: 0.4161170060223688
```

There isn't a great solution for this. If we drop all missing rows with NaN for CS, we're going to lose over 41% of the data prior to 1951. It doesn't encompass enough of the data to just fill in values like we did before, we can't drop rows, and we don't want to drop the column since it isn't missing any data after 1950. One idea, and this may be controversial, is to find the average ratio between SB (stolen bases) and CS and fill in with values based on that ratio.

First, we'll get all data without missing CS values

```
[151]: df_temp = df[(df['CS'].notnull())]
# df_temp
```

```
[152]: total_sb = df_temp['SB'].sum()
total_sb
```

```
[152]: 182622
```

```
[153]: total_cs = df_temp['CS'].sum()
total_cs
```

```
[153]: 94186.0
```

```
[154]: total_sb/total_cs
```

```
[154]: 1.9389505871360924
```

So on average, players are almost twice as likely to steal a base as they are to get caught. This is easy math that we're going to round to make it even easier. It's probably not the best method of solving this issue but at least we still have over 60 years of clean data!

```
[155]: df[(df['CS'].isnull())]
```

```
[155]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|-----------|--------|-------|--------|------|-----|-----|-----|-----|-----|-----|-----|---|
| 14448 | aberrw101 | 1946 | 1 | NY1 | NL | 15 | 8 | 0 | 0 | 0 | ... | 0 | |
| 16285 | aberc101 | 1949 | 1 | CHN | NL | 4 | 7 | 0 | 0 | 0 | ... | 0 | |
| 15712 | aberc101 | 1948 | 1 | CHN | NL | 12 | 32 | 1 | 6 | 1 | ... | 6 | |
| 15131 | aberc101 | 1947 | 1 | CHN | NL | 47 | 140 | 24 | 39 | 6 | ... | 20 | |
| 16286 | abrac101 | 1949 | 1 | BRO | NL | 8 | 24 | 6 | 2 | 1 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | |
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|-----|-----|-----|-----|-----|-----|-----|------|
| 14448 | 0 | NaN | 0 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 16285 | 0 | NaN | 0 | 2 | 0.0 | 0 | 0 | 0.0 | 1.0 |
| 15712 | 0 | NaN | 5 | 10 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 15131 | 0 | NaN | 20 | 32 | 0.0 | 0 | 0 | 0.0 | 5.0 |
| 16286 | 1 | NaN | 7 | 6 | 0.0 | 0 | 0 | 0.0 | 1.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 532 | 2 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN |
| 4782 | 9 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN |
| 5312 | 13 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN |
| 533 | 0 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN |
| 5842 | 4 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN |

```
[7255 rows x 22 columns]
```

```
[156]: df[(df['CS'].isnull())].apply(lambda x: x['SB'] / 2, axis=1)
```

```
[156]:
```

| | |
|-------|-----|
| 14448 | 0.0 |
| 16285 | 0.0 |
| 15712 | 0.0 |
| 15131 | 0.0 |
| 16286 | 0.5 |
| ... | ... |
| 532 | 1.0 |

```
4782      4.5
5312      6.5
533       0.0
5842      2.0
Length: 7255, dtype: float64
```

```
[157]: df[(df['CS']).isnull()].apply(lambda x: x['SB'] / 2, axis=1).value_counts()
```

```
[157]: 0.0      4494
      0.5      794
      1.0      438
      1.5      303
      2.0      257
      2.5      165
      3.0      139
      3.5      131
      4.0       91
      4.5       81
      5.0       51
      5.5       50
      6.5       38
      6.0       36
      7.5       28
      7.0       22
      8.0       21
      9.0       19
      8.5       16
      9.5       11
     11.5        8
     10.0        8
     10.5        8
     11.0        7
     13.0        6
     14.0        5
     12.0        5
     14.5        3
     18.5        3
     12.5        2
     17.5        2
     13.5        2
     16.5        2
     16.0        2
     20.0        1
     15.0        1
     15.5        1
     24.0        1
     18.0        1
```

```
17.0      1
21.5      1
dtype: int64
```

I don't love the max of 24, but overall these values look good and we definitely don't have many of the higher values. So we're going to apply this to our missing CS data

First I'm going to test it out on a copy

```
[158]: df_temp = df[(df['CS']).isnull()]
```

```
[159]: df_temp['CS'] = df_temp.apply(lambda x: x['SB'] / 2, axis=1)
```

```
[160]: df_temp[(df_temp['retroID'] == 'zitzb101')]
```

```
[160]:      retroID  yearID  stint teamID lgID   G  AB  R  H  2B  ...  RBI  SB  \
4241  zitzb101    1926      1    CIN  NL   53  94  21  23   2  ...   3   3
532   zitzb101    1919      1    PIT  NL   11  26   5   5   1  ...   2   2
4782  zitzb101    1927      1    CIN  NL   88 232  47  66  10  ...  24   9
5312  zitzb101    1928      1    CIN  NL  101 266  53  79   9  ...  33  13
533   zitzb101    1919      2    CIN  NL    2   1   0   0   0  ...   0   0
5842  zitzb101    1929      1    CIN  NL   47  84  18  19   3  ...   6   4
```

```
      CS  BB  SO  IBB  HBP  SH  SF  GIDP
4241  1.5   6   7  0.0    2   3  0.0   NaN
532   1.0   0   6  0.0    0   1  0.0   NaN
4782  4.5  20  18  0.0    4  17  0.0   NaN
5312  6.5  13  22  0.0    3  14  0.0   NaN
533   0.0   0   0  0.0    0   0  0.0   NaN
5842  2.0   9  10  0.0    1   2  0.0   NaN
```

```
[6 rows x 22 columns]
```

We know from before that this guy had NaNs for his CS and now it's all filled in, so our plan worked. Let's do it for the actual data

I don't know how to reassign values to a subset of a DataFrame based on a predicate (or if it's possible), so we'll get a little hacky and apply a function with a conditional. Here's what I tried originally:

```
df[(df['CS']).isnull()]['CS'] = df.apply(lambda x: x['SB'] / 2, axis=1)
```

```
[161]: def fill_cs(data):
        if math.isnan(data['CS']):
            return data['SB'] / 2
        else:
            return data['CS']
```

```
[162]: df['CS'] = df.apply(lambda x: fill_cs(x), axis=1)
```

```
[163]: df[(df['retroID'] == 'zitzb101')]
```

```
[163]:      retroID  yearID  stint teamID lgID   G  AB  R  H  2B  ...  RBI  SB  \
4241  zitzb101   1926     1    CIN   NL   53  94  21  23   2  ...   3   3
532   zitzb101   1919     1    PIT   NL   11  26   5   5   1  ...   2   2
3715  zitzb101   1925     1    CIN   NL  104 301  53  76  13  ...  21  11
4782  zitzb101   1927     1    CIN   NL   88 232  47  66  10  ...  24   9
5312  zitzb101   1928     1    CIN   NL  101 266  53  79   9  ...  33  13
533   zitzb101   1919     2    CIN   NL    2   1   0   0   0  ...   0   0
5842  zitzb101   1929     1    CIN   NL   47  84  18  19   3  ...   6   4
```

```
      CS  BB  SO  IBB  HBP  SH  SF  GIDP
4241  1.5   6   7  0.0    2   3  0.0   NaN
532   1.0   0   6  0.0    0   1  0.0   NaN
3715 11.0  35  22  0.0    6   2  0.0   NaN
4782  4.5  20  18  0.0    4  17  0.0   NaN
5312  6.5  13  22  0.0    3  14  0.0   NaN
533   0.0   0   0  0.0    0   0  0.0   NaN
5842  2.0   9  10  0.0    1   2  0.0   NaN
```

```
[7 rows x 22 columns]
```

```
[164]: 100 * df.isnull().sum() / len(df)
```

```
[164]: retroID      0.000000
yearID      0.000000
stint       0.000000
teamID      0.000000
lgID        0.000000
G           0.000000
AB          0.000000
R           0.000000
H           0.000000
2B          0.000000
3B          0.000000
HR          0.000000
RBI         0.000000
SB          0.000000
CS          0.000000
BB          0.000000
SO          0.000000
IBB         0.000000
HBP         0.000000
SH          0.000000
SF          0.000000
GIDP        9.839985
dtype: float64
```

Handling missing GIDP data

```
[165]: df[(df['GIDP'].isnull())]
```

```
[165]:      retroID  yearID  stint  teamID lgID   G  AB  R  H  2B  ...  RBI  \
2082  abrag101   1923     1    CIN   NL    3   1  0  1  0  ...    0
534   acosj101   1920     1    WS1   AL   17  25  2  6  1  ...    1
1569  acosj101   1922     1    CHA   AL    5   5  0  1  0  ...    0
1049  acosj101   1921     1    WS1   AL   33  30  2  2  0  ...    0
6374  adaij102   1931     1    CHN   NL   18  76  9  21  3  ...    3
...      ...      ...    ...    ...   ...  ...  ...  ...  ...  ...  ...
5312  zitzb101   1928     1    CIN   NL  101 266 53 79  9  ...   33
533   zitzb101   1919     2    CIN   NL    2   1  0  0  0  ...    0
5842  zitzb101   1929     1    CIN   NL   47  84 18 19  3  ...    6
9445  zubeb101   1936     1    CLE   AL    2   5  1  1  0  ...    0
10501 zubeb101   1938     1    CLE   AL   15   7  0  0  0  ...    0
```

```
      SB  CS  BB  SO  IBB  HBP  SH  SF  GIDP
2082   0  0.0  0   0  0.0   0   0  0.0  NaN
534    0  0.0  4   7  0.0   0   2  0.0  NaN
1569   0  0.0  1   1  0.0   0   0  0.0  NaN
1049   1  0.0  6  14  0.0   0   1  0.0  NaN
6374   1  0.5  1   8  0.0   0   2  0.0  NaN
...    ..  ...  ..  ..  ...  ...  ..  ...  ...
5312  13  6.5 13  22  0.0   3  14  0.0  NaN
533    0  0.0  0   0  0.0   0   0  0.0  NaN
5842   4  2.0  9  10  0.0   1   2  0.0  NaN
9445   0  0.0  0   1  0.0   0   0  0.0  NaN
10501  0  0.0  0   1  0.0   0   1  0.0  NaN
```

[8683 rows x 22 columns]

```
[166]: df[(df['GIDP'].isnull())]['yearID'].max()
```

```
[166]: 1938
```

```
[167]: df[(df['yearID'] < 1939)].shape[0]
```

```
[167]: 10502
```

```
[168]: df[(df['GIDP'].isnull())].shape[0]
```

```
[168]: 8683
```

```
[169]: 8683/10502
```

```
[169]: 0.8267948962102457
```

Over 82% of records before 1939 are missing GIDP, but it doesn't extend beyond that. I think we can once again just fill the values in with 0

```
[170]: df['GIDP'].fillna(value=0, inplace=True)
```

```
[171]: 100 * df.isnull().sum() / len(df)
```

```
[171]: retroID    0.0
      yearID    0.0
      stint    0.0
      teamID    0.0
      lgID      0.0
      G         0.0
      AB        0.0
      R         0.0
      H         0.0
      2B        0.0
      3B        0.0
      HR        0.0
      RBI       0.0
      SB        0.0
      CS        0.0
      BB        0.0
      SO        0.0
      IBB       0.0
      HBP       0.0
      SH        0.0
      SF        0.0
      GIDP      0.0
      dtype: float64
```

We've handled all missing data in the batting database

Data Integration

Now we need to eliminate an columns that we don't want (if any) and convert the ones we keep to numerical values.

```
[172]: df.head()
```

```
[172]:      retroID  yearID  stint  teamID  lgID   G  AB  R  H  2B  ...  RBI  SB  \
79400  aardd001   2013     1    NYN   NL  43   0  0  0  0  ...   0  0
82244  aardd001   2015     1    ATL   NL  33   1  0  0  0  ...   0  0
69712  aardd001   2006     1    CHN   NL  45   2  0  0  0  ...   0  0
73859  aardd001   2009     1    SEA   AL  73   0  0  0  0  ...   0  0
71089  aardd001   2007     1    CHA   AL  25   0  0  0  0  ...   0  0

      CS  BB  SO  IBB  HBP  SH  SF  GIDP
79400  0.0  0  0  0.0   0  0  0.0  0.0
```

```

82244  0.0  0  1  0.0  0  0  0.0  0.0
69712  0.0  0  0  0.0  0  1  0.0  0.0
73859  0.0  0  0  0.0  0  0  0.0  0.0
71089  0.0  0  0  0.0  0  0  0.0  0.0

```

[5 rows x 22 columns]

```
[173]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 88242 entries, 79400 to 86706
Data columns (total 22 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   retroID    88242 non-null  object
 1   yearID     88242 non-null  int64
 2   stint      88242 non-null  int64
 3   teamID     88242 non-null  object
 4   lgID       88242 non-null  object
 5   G          88242 non-null  int64
 6   AB         88242 non-null  int64
 7   R          88242 non-null  int64
 8   H          88242 non-null  int64
 9   2B         88242 non-null  int64
10   3B         88242 non-null  int64
11   HR         88242 non-null  int64
12   RBI        88242 non-null  int64
13   SB         88242 non-null  int64
14   CS         88242 non-null  float64
15   BB         88242 non-null  int64
16   SO         88242 non-null  int64
17   IBB        88242 non-null  float64
18   HBP        88242 non-null  int64
19   SH         88242 non-null  int64
20   SF         88242 non-null  float64
21   GIDP       88242 non-null  float64
dtypes: float64(4), int64(15), object(3)
memory usage: 15.5+ MB

```

We will handle the metadata columns later and only worry about numerical columns for now

```
[174]: df['lgID'].value_counts()
```

```

[174]: NL      44129
      AL      44113
      Name: lgID, dtype: int64

```

```
[175]: pd.get_dummies(df['lgID'], drop_first=True)
```



```
[175]:      NL
79400    1
82244    1
69712    1
73859    0
71089    0
...     ..
20499    0
18050    0
83729    0
85212    0
86706    0
```

```
[88242 rows x 1 columns]
```

This one will be easy - there are only two leagues in the dataset, so we can just transform that into a single boolean column. Of course that column will be NL, the superior league.

```
[176]: df['NL'] = pd.get_dummies(df['lgID'], drop_first=True)
df.drop(columns=['lgID'], inplace=True)
```

```
[177]: df
```

```
[177]:      retroID  yearID  stint  teamID   G  AB  R  H  2B  3B  ...  SB   CS  BB  \
79400  aardd001   2013      1    NYN   43   0  0  0   0   0  ...   0  0.0   0
82244  aardd001   2015      1    ATL   33   1  0  0   0   0  ...   0  0.0   0
69712  aardd001   2006      1    CHN   45   2  0  0   0   0  ...   0  0.0   0
73859  aardd001   2009      1    SEA   73   0  0  0   0   0  ...   0  0.0   0
71089  aardd001   2007      1    CHA   25   0  0  0   0   0  ...   0  0.0   0
...      ...      ...      ...      ...  ..  ..  ..  ..  ..  ..  ...  ..  ...  ..
20499  zuveg101   1955      2    BAL   28  23  1  5   1   0  ...   0  0.0   1
18050  zuveg101   1951      1    CLE   16   0  0  0   0   0  ...   0  0.0   0
83729  zycht001   2015      1    SEA   13   0  0  0   0   0  ...   0  0.0   0
85212  zycht001   2016      1    SEA   12   0  0  0   0   0  ...   0  0.0   0
86706  zycht001   2017      1    SEA   45   0  0  0   0   0  ...   0  0.0   0

      SO  IBB  HBP  SH  SF  GIDP  NL
79400   0  0.0   0   0  0.0   0.0   1
82244   1  0.0   0   0  0.0   0.0   1
69712   0  0.0   0   1  0.0   0.0   1
73859   0  0.0   0   0  0.0   0.0   0
71089   0  0.0   0   0  0.0   0.0   0
...     ..  ...  ...  ..  ...  ...  ..
20499   5  0.0   0   1  0.0   1.0   0
18050   0  0.0   0   0  0.0   0.0   0
83729   0  0.0   0   0  0.0   0.0   0
85212   0  0.0   0   0  0.0   0.0   0
```

```
86706    0  0.0    0  0  0.0    0.0    0
```

```
[88242 rows x 22 columns]
```

Now we need to figure out how to handle the teamID column.

```
[178]: df['teamID'].nunique()
```

```
[178]: 45
```

Since we have more than 30 team IDs, to keep things consistent I'm just going to map them to franchise ID.

```
[179]: # This will be exported to a separate module
teams = pd.read_csv('../data/lahman/mlb_data/Teams.csv')
teams = teams[['teamID', 'franchID']]
team_dict = teams.set_index('teamID').to_dict()['franchID']

def get_team(team):
    return team_dict[team] if id_dict is not None else team
```

```
[180]: df['teamID'] = df['teamID'].apply(get_team)
```

```
[181]: df['teamID'].nunique()
```

```
[181]: 30
```

We're now all set with team IDs as strings

```
[182]: df.head()
```

```
[182]:
```

| | retroID | yearID | stint | teamID | G | AB | R | H | 2B | 3B | ... | SB | CS | BB | \ |
|-------|----------|--------|-------|--------|----|----|---|---|----|----|-----|----|-----|----|---|
| 79400 | aardd001 | 2013 | 1 | NYM | 43 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 82244 | aardd001 | 2015 | 1 | ATL | 33 | 1 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 69712 | aardd001 | 2006 | 1 | CHC | 45 | 2 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 73859 | aardd001 | 2009 | 1 | SEA | 73 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 71089 | aardd001 | 2007 | 1 | CHW | 25 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |

| | SO | IBB | HBP | SH | SF | GIDP | NL |
|-------|----|-----|-----|----|-----|------|----|
| 79400 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 82244 | 1 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 69712 | 0 | 0.0 | 0 | 1 | 0.0 | 0.0 | 1 |
| 73859 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 71089 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

```
[5 rows x 22 columns]
```

```
[183]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 88242 entries, 79400 to 86706
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  -
0   retroID     88242 non-null  object
1   yearID      88242 non-null  int64
2   stint       88242 non-null  int64
3   teamID      88242 non-null  object
4   G           88242 non-null  int64
5   AB          88242 non-null  int64
6   R           88242 non-null  int64
7   H           88242 non-null  int64
8   2B          88242 non-null  int64
9   3B          88242 non-null  int64
10  HR          88242 non-null  int64
11  RBI         88242 non-null  int64
12  SB          88242 non-null  int64
13  CS          88242 non-null  float64
14  BB          88242 non-null  int64
15  SO          88242 non-null  int64
16  IBB         88242 non-null  float64
17  HBP         88242 non-null  int64
18  SH          88242 non-null  int64
19  SF          88242 non-null  float64
20  GIDP        88242 non-null  float64
21  NL          88242 non-null  uint8
dtypes: float64(4), int64(15), object(2), uint8(1)
memory usage: 14.9+ MB

```

```
[184]: df = df.sort_index()
```

```
[185]: df.head()
```

```

[185]:   retroID  yearID  stint  teamID   G  AB  R  H  2B  3B  ...  SB  CS  BB  \
0  adamb104    1919      1    PIT   34  92  2  17  2   1  ...   0  0.0  6
1  adamb106    1919      1    PHI   78 232 14  54  7   2  ...   4  2.0  6
2  adamw101    1919      1    OAK    1   2   0   0   0   0  ...   0  0.0  0
3  agnes101    1919      1    MIN   42  98  6  23  7   0  ...   1  0.5 10
4  ainse101    1919      1    DET  114 364 42  99 17  12  ...   9  4.5 45

      SO  IBB  HBP  SH  SF  GIDP  NL
0   13  0.0   0   3  0.0  0.0   1
1   27  0.0   0   3  0.0  0.0   1
2    1  0.0   0   0  0.0  0.0   0
3    8  0.0   1   9  0.0  0.0   0
4   30  0.0   1  12  0.0  0.0   0

```

[5 rows x 22 columns]

We need some sort of dictionary to associate a player's retroID with an index. The following steps care of that. This is so we can later associate the correct retroID with our data.

```
[186]: df.reset_index(inplace=True)
```

```
[187]: metadata_column_labels = ['index', 'yearID', 'stint', 'teamID']
```

```
[188]: metadata = df[metadata_column_labels].set_index(df['retroID']).reset_index()
```

```
[189]: metadata.head()
```

```
[189]:
```

| | retroID | index | yearID | stint | teamID |
|---|----------|-------|--------|-------|--------|
| 0 | adamb104 | 0 | 1919 | 1 | PIT |
| 1 | adamb106 | 1 | 1919 | 1 | PHI |
| 2 | adamw101 | 2 | 1919 | 1 | OAK |
| 3 | agnes101 | 3 | 1919 | 1 | MIN |
| 4 | ainse101 | 4 | 1919 | 1 | DET |

The metadata table will eventually be expanded with information from Players.csv to hold all relevant player information that isn't used for the neural network.

```
[190]: indexer = metadata.drop_duplicates('retroID').set_index('index').T.  
      →to_dict('retroID')[0]
```

```
[191]: df = df.drop(columns=metadata_column_labels)
```

```
[192]: df.head()
```

```
[192]:
```

| | retroID | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | SO | IBB | HBP | SH | \ |
|---|----------|-----|-----|----|----|----|----|----|-----|----|-----|----|----|-----|-----|----|---|
| 0 | adamb104 | 34 | 92 | 2 | 17 | 2 | 1 | 0 | 4 | 0 | 0.0 | 6 | 13 | 0.0 | 0 | 3 | |
| 1 | adamb106 | 78 | 232 | 14 | 54 | 7 | 2 | 1 | 17 | 4 | 2.0 | 6 | 27 | 0.0 | 0 | 3 | |
| 2 | adamw101 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 1 | 0.0 | 0 | 0 | |
| 3 | agnes101 | 42 | 98 | 6 | 23 | 7 | 0 | 0 | 10 | 1 | 0.5 | 10 | 8 | 0.0 | 1 | 9 | |
| 4 | ainse101 | 114 | 364 | 42 | 99 | 17 | 12 | 3 | 32 | 9 | 4.5 | 45 | 30 | 0.0 | 1 | 12 | |

| | SF | GIDP | NL |
|---|-----|------|----|
| 0 | 0.0 | 0.0 | 1 |
| 1 | 0.0 | 0.0 | 1 |
| 2 | 0.0 | 0.0 | 0 |
| 3 | 0.0 | 0.0 | 0 |
| 4 | 0.0 | 0.0 | 0 |

Now that the metadata is gone, we just have the ID and the numerical batting information. We can group by the ID and just sum every other column to get player career totals.

```
[193]: df = df.groupby('retroID').sum().reset_index()
```

```
[194]: df
```

```
[194]:
```

| | retroID | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | \ |
|-------|----------|------|-------|------|------|-----|-----|-----|------|-----|------|------|---|
| 0 | aardd001 | 331 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |
| 1 | aaroh101 | 3298 | 12364 | 2174 | 3771 | 624 | 98 | 755 | 2297 | 240 | 73.0 | 1402 | |
| 2 | aarot101 | 437 | 944 | 102 | 216 | 42 | 6 | 13 | 94 | 9 | 8.0 | 86 | |
| 3 | aased001 | 448 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |
| 4 | abada001 | 15 | 21 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 1.0 | 4 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15187 | zupcb001 | 319 | 795 | 99 | 199 | 47 | 4 | 7 | 80 | 7 | 5.0 | 57 | |
| 15188 | zupof101 | 16 | 18 | 3 | 3 | 1 | 0 | 0 | 0 | 0 | 0.0 | 2 | |
| 15189 | zuveg101 | 266 | 142 | 5 | 21 | 2 | 1 | 0 | 7 | 0 | 1.0 | 9 | |
| 15190 | zuvep001 | 209 | 491 | 41 | 109 | 17 | 2 | 2 | 20 | 2 | 0.0 | 34 | |
| 15191 | zycht001 | 70 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |

| | S0 | IBB | HBP | SH | SF | GIDP | NL |
|-------|------|-------|-----|-----|-------|-------|-----|
| 0 | 2 | 0.0 | 0 | 1 | 0.0 | 0.0 | 4 |
| 1 | 1383 | 293.0 | 32 | 21 | 121.0 | 328.0 | 21 |
| 2 | 145 | 3.0 | 0 | 9 | 6.0 | 36.0 | 7 |
| 3 | 3 | 0.0 | 0 | 0 | 0.0 | 0.0 | 2 |
| 4 | 5 | 0.0 | 0 | 0 | 0.0 | 1.0 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 15187 | 137 | 3.0 | 6 | 20 | 8.0 | 15.0 | 0 |
| 15188 | 6 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 15189 | 39 | 0.0 | 0 | 16 | 0.0 | 3.0 | 1 |
| 15190 | 50 | 1.0 | 2 | 18 | 0.0 | 8.0 | 4 |
| 15191 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

```
[15192 rows x 19 columns]
```

Since we summed everything, we just need to change the NL column back. We can divide each value by itself to get either 1 or 0 like we had before.

```
[195]: df['NL'] = np.where(df['NL'] > 0, 1, 0)
```

```
[196]: tensor = df.drop(columns=['retroID'])
```

```
[197]: tensor
```

```
[197]:
```

| | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | S0 | \ |
|---|------|-------|------|------|-----|----|-----|------|-----|------|------|------|---|
| 0 | 331 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 2 | |
| 1 | 3298 | 12364 | 2174 | 3771 | 624 | 98 | 755 | 2297 | 240 | 73.0 | 1402 | 1383 | |
| 2 | 437 | 944 | 102 | 216 | 42 | 6 | 13 | 94 | 9 | 8.0 | 86 | 145 | |
| 3 | 448 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 3 | |
| 4 | 15 | 21 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 1.0 | 4 | 5 | |

| | | | | | | | | | | | | |
|-------|-----|-----|-----|-----|-----|----|-----|-----|-----|-----|-----|-----|
| ... | ... | ... | ... | ... | ... | .. | ... | ... | ... | ... | ... | ... |
| 15187 | 319 | 795 | 99 | 199 | 47 | 4 | 7 | 80 | 7 | 5.0 | 57 | 137 |
| 15188 | 16 | 18 | 3 | 3 | 1 | 0 | 0 | 0 | 0 | 0.0 | 2 | 6 |
| 15189 | 266 | 142 | 5 | 21 | 2 | 1 | 0 | 7 | 0 | 1.0 | 9 | 39 |
| 15190 | 209 | 491 | 41 | 109 | 17 | 2 | 2 | 20 | 2 | 0.0 | 34 | 50 |
| 15191 | 70 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 |

| | | | | | | |
|-------|-------|-----|----|-------|-------|----|
| | IBB | HBP | SH | SF | GIDP | NL |
| 0 | 0.0 | 0 | 1 | 0.0 | 0.0 | 1 |
| 1 | 293.0 | 32 | 21 | 121.0 | 328.0 | 1 |
| 2 | 3.0 | 0 | 9 | 6.0 | 36.0 | 1 |
| 3 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 4 | 0.0 | 0 | 0 | 0.0 | 1.0 | 1 |
| ... | ... | ... | .. | ... | ... | .. |
| 15187 | 3.0 | 6 | 20 | 8.0 | 15.0 | 0 |
| 15188 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 15189 | 0.0 | 0 | 16 | 0.0 | 3.0 | 1 |
| 15190 | 1.0 | 2 | 18 | 0.0 | 8.0 | 1 |
| 15191 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

[15192 rows x 18 columns]

```
[198]: tensor.to_csv('../output/tensor.csv')
        metadata.to_csv('../output/metadata.csv')
```

We now have a tensor with only relevant information, an indexing dictionary to get the player for each row, and a (soon to be expanded) metadata table to get more information on each player.