

Data Aggregation and Preprocessing and Tensor Setup

Data Preprocessing

Data cleaning, normalization and processing and addition of advanced statistics pulled from external sources.

batting_pre

March 9, 2020

```
[120]: import math
import numpy as np
import pandas as pd

# We're going to be reassigning some columns, so we'll turn off this warning -␣
→we know what we're doing!
pd.options.mode.chained_assignment = None # default='warn'
```

```
[121]: # This will be exported to a separate module
ids = pd.read_csv('../data/lahman/mlb_data/People.csv')
ids = ids[['playerID', 'retroID']]
id_dict = ids.set_index('playerID').to_dict()['retroID']

def get_retroid(id):
    return id_dict[id] if id_dict is not None else id
```

```
[122]: df = pd.read_csv('../data/lahman/mlb_data/Batting.csv').sort_values('playerID')
```

```
[123]: df['playerID'] = df['playerID'].apply(get_retroid)
```

```
[124]: df.rename(columns={'playerID': 'retroID'}, inplace=True)
```

```
[125]: df[df['retroID'] == None]
```

```
[125]: Empty DataFrame
Columns: [retroID, yearID, stint, teamID, lgID, G, AB, R, H, 2B, 3B, HR, RBI,
SB, CS, BB, SO, IBB, HBP, SH, SF, GIDP]
Index: []

[0 rows x 22 columns]
```

Cleaning the Data - Missing Values

Print percentages of missing data in each column of the batting table

```
[126]: 100 * df.isnull().sum() / len(df)
```

```
[126]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID      0.000000
      lgID        0.000000
      G           0.000000
      AB          0.000000
      R           0.000000
      H           0.000000
      2B          0.000000
      3B          0.000000
      HR          0.000000
      RBI         0.000000
      SB          0.000000
      CS          8.221708
      BB          0.000000
      SO          0.000000
      IBB         21.711883
      HBP         0.000000
      SH          0.000000
      SF          21.090864
      GIDP        9.839985
      dtype: float64
```

Since this data is by season, it's likely that we have entries for a player for one season with no data in these fields but there is data for other seasons. Since we're taking aggregate sums for each player, we have two options: set these null values to zero so they don't add to the sum, or set them to the average for that player. We'll have to test the theory to see which is more viable.

We're going to start with IBB rather than CS, since it's a more significant chunk of the dataset.

Handling missing IBB data

```
[127]: df[(df['IBB'].isnull())]
```

```
[127]:      retroID  yearID  stint  teamID  lgID   G  AB  R  H  2B  ...  RBI  \
19269  aaroh101   1954      1    ML1    NL  122  468  58  131  27  ...   69
18684  abera101   1953      2    DET    AL   17   23   2   3   0  ...    2
16858  abera101   1950      1    CLE    AL    1    2   0   0   0  ...    0
19270  abera101   1954      1    DET    AL   32   39   3   5   0  ...    3
18683  abera101   1953      1    CLE    AL    6    0   0   0   0  ...    0
...      ...      ...      ...      ...      ...  ...  ...  ..  ...  ..  ...  ...
15128  zubeb101   1946      1    NYA    AL    3    2   0   0   0  ...    0
14447  zubeb101   1945      1    NYA    AL   21   42   1   7   0  ...    3
13868  zubeb101   1944      1    NYA    AL   22   31   1   4   0  ...    1
19843  zuveg101   1954      1    CIN    NL    2    2   1   1   0  ...    0
19844  zuveg101   1954      2    DET    AL   35   64   1   8   1  ...    3
```

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 15128 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |
| 19844 | 0 | 1.0 | 1 | 14 | NaN | 0 | 9 | 0.0 | 2.0 |

[19159 rows x 22 columns]

```
[128]: df[(df['retroID'] == 'abera101')]
```

```
[128]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|-------|----------|--------|-------|--------|------|----|----|---|---|----|-----|-----|----|---|
| 19846 | abera101 | 1955 | 1 | DET | AL | 39 | 17 | 0 | 1 | 0 | ... | 0 | 0 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | 0 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | 0 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | 0 | |
| 21123 | abera101 | 1957 | 2 | KC1 | AL | 3 | 1 | 0 | 1 | 0 | ... | 0 | 0 | |
| 20501 | abera101 | 1956 | 1 | DET | AL | 42 | 10 | 0 | 3 | 0 | ... | 0 | 0 | |
| 21122 | abera101 | 1957 | 1 | DET | AL | 28 | 8 | 0 | 1 | 0 | ... | 1 | 0 | |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|----|----|-----|-----|----|-----|------|
| 19846 | 0.0 | 0 | 9 | 0.0 | 0 | 2 | 0.0 | 1.0 |
| 18684 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| 21123 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 20501 | 0.0 | 1 | 4 | 0.0 | 0 | 2 | 0.0 | 0.0 |
| 21122 | 0.0 | 1 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 |

[8 rows x 22 columns]

```
[129]: df[(df['retroID'] == 'zubeb101')]
```

```
[129]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|-------|----------|--------|-------|--------|------|----|----|---|---|----|-----|-----|----|---|
| 9445 | zubeb101 | 1936 | 1 | CLE | AL | 2 | 5 | 1 | 1 | 0 | ... | 0 | 0 | |
| 10501 | zubeb101 | 1938 | 1 | CLE | AL | 15 | 7 | 0 | 0 | 0 | ... | 0 | 0 | |
| 11080 | zubeb101 | 1939 | 1 | CLE | AL | 16 | 5 | 0 | 1 | 0 | ... | 0 | 0 | |
| 11621 | zubeb101 | 1940 | 1 | CLE | AL | 17 | 3 | 0 | 1 | 0 | ... | 0 | 0 | |
| 12203 | zubeb101 | 1941 | 1 | WS1 | AL | 36 | 26 | 0 | 0 | 0 | ... | 0 | 0 | |

| | | | | | | | | | | | | | |
|-------|----------|------|---|-----|----|----|----|---|---|---|-----|---|---|
| 12742 | zubeb101 | 1942 | 1 | WS1 | AL | 37 | 39 | 5 | 6 | 3 | ... | 3 | 0 |
| 13299 | zubeb101 | 1943 | 1 | NYA | AL | 20 | 38 | 1 | 7 | 1 | ... | 2 | 0 |
| 15711 | zubeb101 | 1947 | 1 | BOS | AL | 20 | 13 | 0 | 2 | 0 | ... | 0 | 0 |
| 15129 | zubeb101 | 1946 | 2 | BOS | AL | 15 | 18 | 1 | 2 | 0 | ... | 2 | 0 |
| 15128 | zubeb101 | 1946 | 1 | NYA | AL | 3 | 2 | 0 | 0 | 0 | ... | 0 | 0 |
| 14447 | zubeb101 | 1945 | 1 | NYA | AL | 21 | 42 | 1 | 7 | 0 | ... | 3 | 0 |
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 | 0 |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|----|----|-----|-----|----|-----|------|
| 9445 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | NaN |
| 10501 | 0.0 | 0 | 1 | NaN | 0 | 1 | NaN | NaN |
| 11080 | 0.0 | 0 | 2 | NaN | 0 | 0 | NaN | 0.0 |
| 11621 | 0.0 | 0 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| 12203 | 0.0 | 1 | 8 | NaN | 0 | 2 | NaN | 0.0 |
| 12742 | 0.0 | 1 | 7 | NaN | 0 | 3 | NaN | 2.0 |
| 13299 | 0.0 | 4 | 14 | NaN | 0 | 5 | NaN | 2.0 |
| 15711 | 0.0 | 2 | 3 | NaN | 0 | 2 | NaN | 2.0 |
| 15129 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 15128 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |

[12 rows x 22 columns]

First let's look at IBB, intentional bases on balls. It seems like most of the missing data is from early in the dataset - it could be that IBB was not recorded then, and/or not considered a trackable play?

```
[130]: df[(df['IBB'].isnull())]['yearID'].max()
```

```
[130]: 1954
```

```
[131]: df[(df['IBB'].isnull())]['yearID'].min()
```

```
[131]: 1919
```

```
[132]: df[(df['IBB'].isnull())]
```

```
[132]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|-----|-----|-----|-----|-----|---|
| 19269 | aaroh101 | 1954 | 1 | ML1 | NL | 122 | 468 | 58 | 131 | 27 | ... | 69 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15128 | zubeb101 | 1946 | 1 | NYA | AL | 3 | 2 | 0 | 0 | 0 | ... | 0 | |
| 14447 | zubeb101 | 1945 | 1 | NYA | AL | 21 | 42 | 1 | 7 | 0 | ... | 3 | |

| | | | | | | | | | | | | |
|-------|----------|------|---|-----|----|----|----|---|---|---|-----|---|
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 |
| 19843 | zuveg101 | 1954 | 1 | CIN | NL | 2 | 2 | 1 | 1 | 0 | ... | 0 |
| 19844 | zuveg101 | 1954 | 2 | DET | AL | 35 | 64 | 1 | 8 | 1 | ... | 3 |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 15128 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 14447 | 0 | 0.0 | 1 | 13 | NaN | 0 | 2 | NaN | 1.0 |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |
| 19844 | 0 | 1.0 | 1 | 14 | NaN | 0 | 9 | 0.0 | 2.0 |

[19159 rows x 22 columns]

We have 19159 total rows where there is no data for IBB, and we know none of those rows goes past the year 1954...

```
[133]: df[(df['yearID'] < 1955)]
```

```
[133]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|----|-----|----|-----|-----|---|
| 19269 | aaroh101 | 1954 | 1 | ML1 | NL | 122 | 468 | 58 | 131 | 27 | ... | 69 | |
| 18684 | abera101 | 1953 | 2 | DET | AL | 17 | 23 | 2 | 3 | 0 | ... | 2 | |
| 16858 | abera101 | 1950 | 1 | CLE | AL | 1 | 2 | 0 | 0 | 0 | ... | 0 | |
| 19270 | abera101 | 1954 | 1 | DET | AL | 32 | 39 | 3 | 5 | 0 | ... | 3 | |
| 18683 | abera101 | 1953 | 1 | CLE | AL | 6 | 0 | 0 | 0 | 0 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | .. | ... | .. | ... | ... | |
| 13868 | zubeb101 | 1944 | 1 | NYA | AL | 22 | 31 | 1 | 4 | 0 | ... | 1 | |
| 18682 | zuveg101 | 1952 | 1 | CLE | AL | 2 | 0 | 1 | 0 | 0 | ... | 0 | |
| 19843 | zuveg101 | 1954 | 1 | CIN | NL | 2 | 2 | 1 | 1 | 0 | ... | 0 | |
| 19844 | zuveg101 | 1954 | 2 | DET | AL | 35 | 64 | 1 | 8 | 1 | ... | 3 | |
| 18050 | zuveg101 | 1951 | 1 | CLE | AL | 16 | 0 | 0 | 0 | 0 | ... | 0 | |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|----|-----|----|----|-----|-----|----|-----|------|
| 19269 | 2 | 2.0 | 28 | 39 | NaN | 3 | 6 | 4.0 | 13.0 |
| 18684 | 0 | 0.0 | 1 | 6 | NaN | 0 | 1 | NaN | 0.0 |
| 16858 | 0 | 0.0 | 1 | 1 | NaN | 0 | 0 | NaN | 0.0 |
| 19270 | 0 | 0.0 | 2 | 17 | NaN | 0 | 3 | 1.0 | 1.0 |
| 18683 | 0 | 0.0 | 2 | 0 | NaN | 0 | 0 | NaN | 0.0 |
| ... | .. | ... | .. | .. | ... | ... | .. | ... | ... |
| 13868 | 0 | 0.0 | 0 | 10 | NaN | 0 | 4 | NaN | 1.0 |
| 18682 | 0 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 19843 | 0 | 0.0 | 0 | 1 | NaN | 0 | 0 | 0.0 | 0.0 |

```
19844    0  1.0   1  14  NaN    0   9  0.0   2.0
18050    0  0.0   0   0  0.0    0   0  0.0   0.0
```

```
[19845 rows x 22 columns]
```

And we have 19845 total rows up to the year 1954. That means...

```
[134]: 19159 / 19845
```

```
[134]: 0.9654320987654321
```

Over 96% of the data before 1955 is missing IBB. I think this gives justification to just setting all of those NaNs to 0.

```
[135]: df['IBB'].fillna(value=0, inplace=True)
```

```
[136]: 100 * df.isnull().sum() / len(df)
```

```
[136]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID      0.000000
      lgID        0.000000
      G           0.000000
      AB          0.000000
      R           0.000000
      H           0.000000
      2B          0.000000
      3B          0.000000
      HR          0.000000
      RBI         0.000000
      SB          0.000000
      CS          8.221708
      BB          0.000000
      SO          0.000000
      IBB         0.000000
      HBP         0.000000
      SH          0.000000
      SF          21.090864
      GIDP        9.839985
      dtype: float64
```

Our IBB issue is solved. Let's move on to SF (sacrifice flies). We'll check the years and rows again to see if we're justified in using the same method to eliminate nulls.

Handling missing SF data

```
[137]: df[(df['SF'].isnull())]['yearID'].max()
```



```
[137]: 1953
```

```
[138]: df[(df['SF'].isnull())]['yearID'].min()
```

```
[138]: 1919
```

```
[139]: df[(df['SF'].isnull())].shape[0]
```

```
[139]: 18611
```

```
[140]: df[(df['yearID'] < 1954)].shape[0]
```

```
[140]: 19269
```

```
[141]: 18611/19269
```

```
[141]: 0.965851886449738
```

Almost the same percentage, and one less year covered. I think we can fill those missing values with 0.

```
[142]: df['SF'].fillna(value=0, inplace=True)
```

```
[143]: 100 * df.isnull().sum() / len(df)
```

```
[143]: retroID      0.000000
      yearID      0.000000
      stint      0.000000
      teamID      0.000000
      lgID        0.000000
      G           0.000000
      AB          0.000000
      R           0.000000
      H           0.000000
      2B          0.000000
      3B          0.000000
      HR          0.000000
      RBI         0.000000
      SB          0.000000
      CS          8.221708
      BB          0.000000
      SO          0.000000
      IBB         0.000000
      HBP         0.000000
      SH          0.000000
      SF          0.000000
      GIDP        9.839985
      dtype: float64
```

Two more to go, let's move on to CS (caught stealing)

Handling missing CS data

```
[144]: df[(df['CS'].isnull())]['yearID'].max()
```

```
[144]: 1950
```

```
[145]: df[(df['CS'].isnull())]['yearID'].min()
```

```
[145]: 1919
```

```
[146]: df[(df['CS'].isnull())]
```

```
[146]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|----------|--------|-------|--------|------|-----|-----|-----|------|-----|-----|-----|---|
| 14448 | abewr101 | 1946 | 1 | NY1 | NL | 15 | 8 | 0 | 0 | 0 | ... | 0 | |
| 16285 | aberc101 | 1949 | 1 | CHN | NL | 4 | 7 | 0 | 0 | 0 | ... | 0 | |
| 15712 | aberc101 | 1948 | 1 | CHN | NL | 12 | 32 | 1 | 6 | 1 | ... | 6 | |
| 15131 | aberc101 | 1947 | 1 | CHN | NL | 47 | 140 | 24 | 39 | 6 | ... | 20 | |
| 16286 | abrac101 | 1949 | 1 | BRO | NL | 8 | 24 | 6 | 2 | 1 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | |
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | |
| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP | | | | |
| 14448 | 0 | NaN | 0 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 | | | | |
| 16285 | 0 | NaN | 0 | 2 | 0.0 | 0 | 0 | 0.0 | 1.0 | | | | |
| 15712 | 0 | NaN | 5 | 10 | 0.0 | 0 | 0 | 0.0 | 0.0 | | | | |
| 15131 | 0 | NaN | 20 | 32 | 0.0 | 0 | 0 | 0.0 | 5.0 | | | | |
| 16286 | 1 | NaN | 7 | 6 | 0.0 | 0 | 0 | 0.0 | 1.0 | | | | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | | | | |
| 532 | 2 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN | | | | |
| 4782 | 9 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN | | | | |
| 5312 | 13 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN | | | | |
| 533 | 0 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN | | | | |
| 5842 | 4 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN | | | | |

[7255 rows x 22 columns]

```
[147]: df[(df['retroID'] == 'zitzb101')]
```

```
[147]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|------|----------|--------|-------|--------|------|-----|-----|----|----|----|-----|-----|----|---|
| 4241 | zitzb101 | 1926 | 1 | CIN | NL | 53 | 94 | 21 | 23 | 2 | ... | 3 | 3 | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | 2 | |
| 3715 | zitzb101 | 1925 | 1 | CIN | NL | 104 | 301 | 53 | 76 | 13 | ... | 21 | 11 | |

| | | | | | | | | | | | | | |
|------|----------|------|---|-----|----|-----|-----|----|----|----|-----|----|----|
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | 9 |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | 13 |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | 0 |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | 4 |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|------|------|----|----|-----|-----|----|-----|------|
| 4241 | NaN | 6 | 7 | 0.0 | 2 | 3 | 0.0 | NaN |
| 532 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN |
| 3715 | 11.0 | 35 | 22 | 0.0 | 6 | 2 | 0.0 | NaN |
| 4782 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN |
| 5312 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN |
| 533 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN |
| 5842 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN |

[7 rows x 22 columns]

```
[148]: df[(df['CS'].isnull())].shape[0]
```

```
[148]: 7255
```

```
[149]: df[(df['yearID'] < 1951)].shape[0]
```

```
[149]: 17435
```

```
[150]: 7255/17435
```

```
[150]: 0.4161170060223688
```

There isn't a great solution for this. If we drop all missing rows with NaN for CS, we're going to lose over 41% of the data prior to 1951. It doesn't encompass enough of the data to just fill in values like we did before, we can't drop rows, and we don't want to drop the column since it isn't missing any data after 1950. One idea, and this may be controversial, is to find the average ratio between SB (stolen bases) and CS and fill in with values based on that ratio.

First, we'll get all data without missing CS values

```
[151]: df_temp = df[(df['CS'].notnull())]
# df_temp
```

```
[152]: total_sb = df_temp['SB'].sum()
total_sb
```

```
[152]: 182622
```

```
[153]: total_cs = df_temp['CS'].sum()
total_cs
```

```
[153]: 94186.0
```

```
[154]: total_sb/total_cs
```

```
[154]: 1.9389505871360924
```

So on average, players are almost twice as likely to steal a base as they are to get caught. This is easy math that we're going to round to make it even easier. It's probably not the best method of solving this issue but at least we still have over 60 years of clean data!

```
[155]: df[(df['CS'].isnull())]
```

```
[155]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | \ |
|-------|-----------|--------|-------|--------|------|-----|-----|-----|-----|-----|-----|-----|---|
| 14448 | aberrw101 | 1946 | 1 | NY1 | NL | 15 | 8 | 0 | 0 | 0 | ... | 0 | |
| 16285 | aberc101 | 1949 | 1 | CHN | NL | 4 | 7 | 0 | 0 | 0 | ... | 0 | |
| 15712 | aberc101 | 1948 | 1 | CHN | NL | 12 | 32 | 1 | 6 | 1 | ... | 6 | |
| 15131 | aberc101 | 1947 | 1 | CHN | NL | 47 | 140 | 24 | 39 | 6 | ... | 20 | |
| 16286 | abrac101 | 1949 | 1 | BRO | NL | 8 | 24 | 6 | 2 | 1 | ... | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | |
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | |

| | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|-----|-----|-----|-----|-----|-----|-----|------|
| 14448 | 0 | NaN | 0 | 4 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 16285 | 0 | NaN | 0 | 2 | 0.0 | 0 | 0 | 0.0 | 1.0 |
| 15712 | 0 | NaN | 5 | 10 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 15131 | 0 | NaN | 20 | 32 | 0.0 | 0 | 0 | 0.0 | 5.0 |
| 16286 | 1 | NaN | 7 | 6 | 0.0 | 0 | 0 | 0.0 | 1.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 532 | 2 | NaN | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN |
| 4782 | 9 | NaN | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN |
| 5312 | 13 | NaN | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN |
| 533 | 0 | NaN | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN |
| 5842 | 4 | NaN | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN |

```
[7255 rows x 22 columns]
```

```
[156]: df[(df['CS'].isnull())].apply(lambda x: x['SB'] / 2, axis=1)
```

```
[156]:
```

| | |
|-------|-----|
| 14448 | 0.0 |
| 16285 | 0.0 |
| 15712 | 0.0 |
| 15131 | 0.0 |
| 16286 | 0.5 |
| ... | ... |
| 532 | 1.0 |

```
4782      4.5
5312      6.5
533       0.0
5842      2.0
Length: 7255, dtype: float64
```

```
[157]: df[(df['CS']).isnull()].apply(lambda x: x['SB'] / 2, axis=1).value_counts()
```

```
[157]: 0.0      4494
      0.5      794
      1.0      438
      1.5      303
      2.0      257
      2.5      165
      3.0      139
      3.5      131
      4.0       91
      4.5       81
      5.0       51
      5.5       50
      6.5       38
      6.0       36
      7.5       28
      7.0       22
      8.0       21
      9.0       19
      8.5       16
      9.5       11
     11.5        8
     10.0        8
     10.5        8
     11.0        7
     13.0        6
     14.0        5
     12.0        5
     14.5        3
     18.5        3
     12.5        2
     17.5        2
     13.5        2
     16.5        2
     16.0        2
     20.0        1
     15.0        1
     15.5        1
     24.0        1
     18.0        1
```

```
17.0      1
21.5      1
dtype: int64
```

I don't love the max of 24, but overall these values look good and we definitely don't have many of the higher values. So we're going to apply this to our missing CS data

First I'm going to test it out on a copy

```
[158]: df_temp = df[(df['CS']).isnull()]
```

```
[159]: df_temp['CS'] = df_temp.apply(lambda x: x['SB'] / 2, axis=1)
```

```
[160]: df_temp[(df_temp['retroID'] == 'zitzb101')]
```

```
[160]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|------|----------|--------|-------|--------|------|-----|-----|----|----|----|-----|-----|----|---|
| 4241 | zitzb101 | 1926 | 1 | CIN | NL | 53 | 94 | 21 | 23 | 2 | ... | 3 | 3 | |
| 532 | zitzb101 | 1919 | 1 | PIT | NL | 11 | 26 | 5 | 5 | 1 | ... | 2 | 2 | |
| 4782 | zitzb101 | 1927 | 1 | CIN | NL | 88 | 232 | 47 | 66 | 10 | ... | 24 | 9 | |
| 5312 | zitzb101 | 1928 | 1 | CIN | NL | 101 | 266 | 53 | 79 | 9 | ... | 33 | 13 | |
| 533 | zitzb101 | 1919 | 2 | CIN | NL | 2 | 1 | 0 | 0 | 0 | ... | 0 | 0 | |
| 5842 | zitzb101 | 1929 | 1 | CIN | NL | 47 | 84 | 18 | 19 | 3 | ... | 6 | 4 | |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|------|-----|----|----|-----|-----|----|-----|------|
| 4241 | 1.5 | 6 | 7 | 0.0 | 2 | 3 | 0.0 | NaN |
| 532 | 1.0 | 0 | 6 | 0.0 | 0 | 1 | 0.0 | NaN |
| 4782 | 4.5 | 20 | 18 | 0.0 | 4 | 17 | 0.0 | NaN |
| 5312 | 6.5 | 13 | 22 | 0.0 | 3 | 14 | 0.0 | NaN |
| 533 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | NaN |
| 5842 | 2.0 | 9 | 10 | 0.0 | 1 | 2 | 0.0 | NaN |

```
[6 rows x 22 columns]
```

We know from before that this guy had NaNs for his CS and now it's all filled in, so our plan worked. Let's do it for the actual data

I don't know how to reassign values to a subset of a DataFrame based on a predicate (or if it's possible), so we'll get a little hacky and apply a function with a conditional. Here's what I tried originally:

```
df[(df['CS']).isnull()]['CS'] = df.apply(lambda x: x['SB'] / 2, axis=1)
```

```
[161]: def fill_cs(data):
        if math.isnan(data['CS']):
            return data['SB'] / 2
        else:
            return data['CS']
```

```
[162]: df['CS'] = df.apply(lambda x: fill_cs(x), axis=1)
```

```
[163]: df[(df['retroID'] == 'zitzb101')]
```

```
[163]:      retroID  yearID  stint teamID lgID   G  AB  R  H  2B  ...  RBI  SB  \
4241  zitzb101   1926     1    CIN   NL   53  94  21  23   2  ...   3   3
532   zitzb101   1919     1    PIT   NL   11  26   5   5   1  ...   2   2
3715  zitzb101   1925     1    CIN   NL  104 301  53  76  13  ...  21  11
4782  zitzb101   1927     1    CIN   NL   88 232  47  66  10  ...  24   9
5312  zitzb101   1928     1    CIN   NL  101 266  53  79   9  ...  33  13
533   zitzb101   1919     2    CIN   NL    2   1   0   0   0  ...   0   0
5842  zitzb101   1929     1    CIN   NL   47  84  18  19   3  ...   6   4
```

```
      CS  BB  SO  IBB  HBP  SH  SF  GIDP
4241  1.5   6   7  0.0    2   3  0.0   NaN
532   1.0   0   6  0.0    0   1  0.0   NaN
3715 11.0  35  22  0.0    6   2  0.0   NaN
4782  4.5  20  18  0.0    4  17  0.0   NaN
5312  6.5  13  22  0.0    3  14  0.0   NaN
533   0.0   0   0  0.0    0   0  0.0   NaN
5842  2.0   9  10  0.0    1   2  0.0   NaN
```

```
[7 rows x 22 columns]
```

```
[164]: 100 * df.isnull().sum() / len(df)
```

```
[164]: retroID    0.000000
yearID      0.000000
stint       0.000000
teamID      0.000000
lgID        0.000000
G           0.000000
AB          0.000000
R           0.000000
H           0.000000
2B          0.000000
3B          0.000000
HR          0.000000
RBI         0.000000
SB          0.000000
CS          0.000000
BB          0.000000
SO          0.000000
IBB         0.000000
HBP         0.000000
SH          0.000000
SF          0.000000
GIDP        9.839985
dtype: float64
```

Handling missing GIDP data

```
[165]: df[(df['GIDP'].isnull())]
```

```
[165]:      retroID  yearID  stint  teamID lgID   G  AB  R  H  2B  ...  RBI  \
2082  abrag101   1923     1    CIN   NL    3   1  0  1  0  ...    0
534   acosj101   1920     1    WS1   AL   17  25  2  6  1  ...    1
1569  acosj101   1922     1    CHA   AL    5   5  0  1  0  ...    0
1049  acosj101   1921     1    WS1   AL   33  30  2  2  0  ...    0
6374  adaij102   1931     1    CHN   NL   18  76  9  21  3  ...    3
...      ...      ...    ...    ...   ...  ...  ...  ...  ...  ...  ...
5312  zitzb101   1928     1    CIN   NL  101 266 53 79  9  ...   33
533   zitzb101   1919     2    CIN   NL    2   1  0  0  0  ...    0
5842  zitzb101   1929     1    CIN   NL   47  84 18 19  3  ...    6
9445  zubeb101   1936     1    CLE   AL    2   5  1  1  0  ...    0
10501 zubeb101   1938     1    CLE   AL   15   7  0  0  0  ...    0
```

```
      SB  CS  BB  SO  IBB  HBP  SH  SF  GIDP
2082   0  0.0  0   0  0.0   0   0  0.0  NaN
534    0  0.0  4   7  0.0   0   2  0.0  NaN
1569   0  0.0  1   1  0.0   0   0  0.0  NaN
1049   1  0.0  6  14  0.0   0   1  0.0  NaN
6374   1  0.5  1   8  0.0   0   2  0.0  NaN
...    ..  ...  ..  ..  ...  ...  ..  ...  ...
5312  13  6.5 13  22  0.0   3  14  0.0  NaN
533    0  0.0  0   0  0.0   0   0  0.0  NaN
5842   4  2.0  9  10  0.0   1   2  0.0  NaN
9445   0  0.0  0   1  0.0   0   0  0.0  NaN
10501  0  0.0  0   1  0.0   0   1  0.0  NaN
```

[8683 rows x 22 columns]

```
[166]: df[(df['GIDP'].isnull())]['yearID'].max()
```

```
[166]: 1938
```

```
[167]: df[(df['yearID'] < 1939)].shape[0]
```

```
[167]: 10502
```

```
[168]: df[(df['GIDP'].isnull())].shape[0]
```

```
[168]: 8683
```

```
[169]: 8683/10502
```

```
[169]: 0.8267948962102457
```


Over 82% of records before 1939 are missing GIDP, but it doesn't extend beyond that. I think we can once again just fill the values in with 0

```
[170]: df['GIDP'].fillna(value=0, inplace=True)
```

```
[171]: 100 * df.isnull().sum() / len(df)
```

```
[171]: retroID    0.0
      yearID    0.0
      stint    0.0
      teamID    0.0
      lgID     0.0
      G        0.0
      AB       0.0
      R        0.0
      H        0.0
      2B       0.0
      3B       0.0
      HR       0.0
      RBI      0.0
      SB       0.0
      CS       0.0
      BB       0.0
      SO       0.0
      IBB      0.0
      HBP      0.0
      SH       0.0
      SF       0.0
      GIDP     0.0
      dtype: float64
```

We've handled all missing data in the batting database

Data Integration

Now we need to eliminate an columns that we don't want (if any) and convert the ones we keep to numerical values.

```
[172]: df.head()
```

```
[172]:
```

| | retroID | yearID | stint | teamID | lgID | G | AB | R | H | 2B | ... | RBI | SB | \ |
|-------|----------|--------|-------|--------|------|----|----|---|---|----|-----|-----|----|---|
| 79400 | aardd001 | 2013 | 1 | NYN | NL | 43 | 0 | 0 | 0 | 0 | ... | 0 | 0 | |
| 82244 | aardd001 | 2015 | 1 | ATL | NL | 33 | 1 | 0 | 0 | 0 | ... | 0 | 0 | |
| 69712 | aardd001 | 2006 | 1 | CHN | NL | 45 | 2 | 0 | 0 | 0 | ... | 0 | 0 | |
| 73859 | aardd001 | 2009 | 1 | SEA | AL | 73 | 0 | 0 | 0 | 0 | ... | 0 | 0 | |
| 71089 | aardd001 | 2007 | 1 | CHA | AL | 25 | 0 | 0 | 0 | 0 | ... | 0 | 0 | |

| | CS | BB | SO | IBB | HBP | SH | SF | GIDP |
|-------|-----|----|----|-----|-----|----|-----|------|
| 79400 | 0.0 | 0 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |

```

82244  0.0  0  1  0.0  0  0  0.0  0.0
69712  0.0  0  0  0.0  0  1  0.0  0.0
73859  0.0  0  0  0.0  0  0  0.0  0.0
71089  0.0  0  0  0.0  0  0  0.0  0.0

```

[5 rows x 22 columns]

```
[173]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 88242 entries, 79400 to 86706
Data columns (total 22 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   retroID    88242 non-null  object
 1   yearID     88242 non-null  int64
 2   stint      88242 non-null  int64
 3   teamID     88242 non-null  object
 4   lgID       88242 non-null  object
 5   G          88242 non-null  int64
 6   AB         88242 non-null  int64
 7   R          88242 non-null  int64
 8   H          88242 non-null  int64
 9   2B         88242 non-null  int64
10  3B         88242 non-null  int64
11  HR         88242 non-null  int64
12  RBI        88242 non-null  int64
13  SB         88242 non-null  int64
14  CS         88242 non-null  float64
15  BB         88242 non-null  int64
16  SO         88242 non-null  int64
17  IBB        88242 non-null  float64
18  HBP        88242 non-null  int64
19  SH         88242 non-null  int64
20  SF         88242 non-null  float64
21  GIDP       88242 non-null  float64
dtypes: float64(4), int64(15), object(3)
memory usage: 15.5+ MB

```

We will handle the metadata columns later and only worry about numerical columns for now

```
[174]: df['lgID'].value_counts()
```

```

[174]: NL      44129
      AL      44113
      Name: lgID, dtype: int64

```

```
[175]: pd.get_dummies(df['lgID'], drop_first=True)
```

```
[175]:      NL
79400    1
82244    1
69712    1
73859    0
71089    0
...     ..
20499    0
18050    0
83729    0
85212    0
86706    0

[88242 rows x 1 columns]
```

This one will be easy - there are only two leagues in the dataset, so we can just transform that into a single boolean column. Of course that column will be NL, the superior league.

```
[176]: df['NL'] = pd.get_dummies(df['lgID'], drop_first=True)
df.drop(columns=['lgID'], inplace=True)
```

```
[177]: df
```

```
[177]:      retroID  yearID  stint  teamID   G  AB  R  H  2B  3B  ...  SB  CS  BB  \
79400  aardd001   2013      1    NYN   43   0  0  0   0   0  ...   0  0.0  0
82244  aardd001   2015      1    ATL   33   1  0  0   0   0  ...   0  0.0  0
69712  aardd001   2006      1    CHN   45   2  0  0   0   0  ...   0  0.0  0
73859  aardd001   2009      1    SEA   73   0  0  0   0   0  ...   0  0.0  0
71089  aardd001   2007      1    CHA   25   0  0  0   0   0  ...   0  0.0  0
...     ...     ...     ...     ...   ..  ..  ..  ..  ..  ..  ...  ..  ...  ..
20499  zuveg101   1955      2    BAL   28  23  1  5   1   0  ...   0  0.0  1
18050  zuveg101   1951      1    CLE   16   0  0  0   0   0  ...   0  0.0  0
83729  zycht001   2015      1    SEA   13   0  0  0   0   0  ...   0  0.0  0
85212  zycht001   2016      1    SEA   12   0  0  0   0   0  ...   0  0.0  0
86706  zycht001   2017      1    SEA   45   0  0  0   0   0  ...   0  0.0  0

      SO  IBB  HBP  SH  SF  GIDP  NL
79400  0  0.0   0   0  0.0   0.0   1
82244  1  0.0   0   0  0.0   0.0   1
69712  0  0.0   0   1  0.0   0.0   1
73859  0  0.0   0   0  0.0   0.0   0
71089  0  0.0   0   0  0.0   0.0   0
...     ..  ...  ...  ..  ...  ...  ..
20499  5  0.0   0   1  0.0   1.0   0
18050  0  0.0   0   0  0.0   0.0   0
83729  0  0.0   0   0  0.0   0.0   0
85212  0  0.0   0   0  0.0   0.0   0
```

```
86706    0  0.0    0  0  0.0    0.0    0
```

```
[88242 rows x 22 columns]
```

Now we need to figure out how to handle the teamID column.

```
[178]: df['teamID'].nunique()
```

```
[178]: 45
```

Since we have more than 30 team IDs, to keep things consistent I'm just going to map them to franchise ID.

```
[179]: # This will be exported to a separate module
teams = pd.read_csv('../data/lahman/mlb_data/Teams.csv')
teams = teams[['teamID', 'franchID']]
team_dict = teams.set_index('teamID').to_dict()['franchID']

def get_team(team):
    return team_dict[team] if id_dict is not None else team
```

```
[180]: df['teamID'] = df['teamID'].apply(get_team)
```

```
[181]: df['teamID'].nunique()
```

```
[181]: 30
```

We're now all set with team IDs as strings

```
[182]: df.head()
```

```
[182]:
```

| | retroID | yearID | stint | teamID | G | AB | R | H | 2B | 3B | ... | SB | CS | BB | \ |
|-------|----------|--------|-------|--------|----|----|---|---|----|----|-----|----|-----|----|---|
| 79400 | aardd001 | 2013 | 1 | NYM | 43 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 82244 | aardd001 | 2015 | 1 | ATL | 33 | 1 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 69712 | aardd001 | 2006 | 1 | CHC | 45 | 2 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 73859 | aardd001 | 2009 | 1 | SEA | 73 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |
| 71089 | aardd001 | 2007 | 1 | CHW | 25 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | |

| | SO | IBB | HBP | SH | SF | GIDP | NL |
|-------|----|-----|-----|----|-----|------|----|
| 79400 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 82244 | 1 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 69712 | 0 | 0.0 | 0 | 1 | 0.0 | 0.0 | 1 |
| 73859 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 71089 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

```
[5 rows x 22 columns]
```

```
[183]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 88242 entries, 79400 to 86706
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  -
0   retroID     88242 non-null  object
1   yearID      88242 non-null  int64
2   stint       88242 non-null  int64
3   teamID      88242 non-null  object
4   G           88242 non-null  int64
5   AB          88242 non-null  int64
6   R           88242 non-null  int64
7   H           88242 non-null  int64
8   2B          88242 non-null  int64
9   3B          88242 non-null  int64
10  HR          88242 non-null  int64
11  RBI         88242 non-null  int64
12  SB          88242 non-null  int64
13  CS          88242 non-null  float64
14  BB          88242 non-null  int64
15  SO          88242 non-null  int64
16  IBB         88242 non-null  float64
17  HBP         88242 non-null  int64
18  SH          88242 non-null  int64
19  SF          88242 non-null  float64
20  GIDP        88242 non-null  float64
21  NL          88242 non-null  uint8
dtypes: float64(4), int64(15), object(2), uint8(1)
memory usage: 14.9+ MB

```

```
[184]: df = df.sort_index()
```

```
[185]: df.head()
```

```

[185]:   retroID  yearID  stint  teamID   G  AB  R  H  2B  3B  ...  SB  CS  BB  \
0  adamb104    1919      1    PIT   34  92  2  17  2   1  ...  0  0.0  6
1  adamb106    1919      1    PHI   78 232 14  54  7   2  ...  4  2.0  6
2  adamw101    1919      1    OAK    1   2  0   0  0   0  ...  0  0.0  0
3  agnes101    1919      1    MIN   42  98  6  23  7   0  ...  1  0.5 10
4  ainse101    1919      1    DET  114 364 42  99 17  12  ...  9  4.5 45

      SO  IBB  HBP  SH  SF  GIDP  NL
0   13  0.0   0   3  0.0  0.0   1
1   27  0.0   0   3  0.0  0.0   1
2    1  0.0   0   0  0.0  0.0   0
3    8  0.0   1   9  0.0  0.0   0
4   30  0.0   1  12  0.0  0.0   0

```

[5 rows x 22 columns]

We need some sort of dictionary to associate a player's retroID with an index. The following steps care of that. This is so we can later associate the correct retroID with our data.

```
[186]: df.reset_index(inplace=True)
```

```
[187]: metadata_column_labels = ['index', 'yearID', 'stint', 'teamID']
```

```
[188]: metadata = df[metadata_column_labels].set_index(df['retroID']).reset_index()
```

```
[189]: metadata.head()
```

```
[189]:
```

| | retroID | index | yearID | stint | teamID |
|---|----------|-------|--------|-------|--------|
| 0 | adamb104 | 0 | 1919 | 1 | PIT |
| 1 | adamb106 | 1 | 1919 | 1 | PHI |
| 2 | adamw101 | 2 | 1919 | 1 | OAK |
| 3 | agnes101 | 3 | 1919 | 1 | MIN |
| 4 | ainse101 | 4 | 1919 | 1 | DET |

The metadata table will eventually be expanded with information from Players.csv to hold all relevant player information that isn't used for the neural network.

```
[190]: indexer = metadata.drop_duplicates('retroID').set_index('index').T.  
      →to_dict('retroID')[0]
```

```
[191]: df = df.drop(columns=metadata_column_labels)
```

```
[192]: df.head()
```

```
[192]:
```

| | retroID | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | SO | IBB | HBP | SH | \ |
|---|----------|-----|-----|----|----|----|----|----|-----|----|-----|----|----|-----|-----|----|---|
| 0 | adamb104 | 34 | 92 | 2 | 17 | 2 | 1 | 0 | 4 | 0 | 0.0 | 6 | 13 | 0.0 | 0 | 3 | |
| 1 | adamb106 | 78 | 232 | 14 | 54 | 7 | 2 | 1 | 17 | 4 | 2.0 | 6 | 27 | 0.0 | 0 | 3 | |
| 2 | adamw101 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 1 | 0.0 | 0 | 0 | |
| 3 | agnes101 | 42 | 98 | 6 | 23 | 7 | 0 | 0 | 10 | 1 | 0.5 | 10 | 8 | 0.0 | 1 | 9 | |
| 4 | ainse101 | 114 | 364 | 42 | 99 | 17 | 12 | 3 | 32 | 9 | 4.5 | 45 | 30 | 0.0 | 1 | 12 | |

| | SF | GIDP | NL |
|---|-----|------|----|
| 0 | 0.0 | 0.0 | 1 |
| 1 | 0.0 | 0.0 | 1 |
| 2 | 0.0 | 0.0 | 0 |
| 3 | 0.0 | 0.0 | 0 |
| 4 | 0.0 | 0.0 | 0 |

Now that the metadata is gone, we just have the ID and the numerical batting information. We can group by the ID and just sum every other column to get player career totals.

```
[193]: df = df.groupby('retroID').sum().reset_index()
```

```
[194]: df
```

```
[194]:
```

| | retroID | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | \ |
|-------|----------|------|-------|------|------|-----|-----|-----|------|-----|------|------|---|
| 0 | aardd001 | 331 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |
| 1 | aaroh101 | 3298 | 12364 | 2174 | 3771 | 624 | 98 | 755 | 2297 | 240 | 73.0 | 1402 | |
| 2 | aarot101 | 437 | 944 | 102 | 216 | 42 | 6 | 13 | 94 | 9 | 8.0 | 86 | |
| 3 | aased001 | 448 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |
| 4 | abada001 | 15 | 21 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 1.0 | 4 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15187 | zupcb001 | 319 | 795 | 99 | 199 | 47 | 4 | 7 | 80 | 7 | 5.0 | 57 | |
| 15188 | zupof101 | 16 | 18 | 3 | 3 | 1 | 0 | 0 | 0 | 0 | 0.0 | 2 | |
| 15189 | zuveg101 | 266 | 142 | 5 | 21 | 2 | 1 | 0 | 7 | 0 | 1.0 | 9 | |
| 15190 | zuvep001 | 209 | 491 | 41 | 109 | 17 | 2 | 2 | 20 | 2 | 0.0 | 34 | |
| 15191 | zycht001 | 70 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | |

| | S0 | IBB | HBP | SH | SF | GIDP | NL |
|-------|------|-------|-----|-----|-------|-------|-----|
| 0 | 2 | 0.0 | 0 | 1 | 0.0 | 0.0 | 4 |
| 1 | 1383 | 293.0 | 32 | 21 | 121.0 | 328.0 | 21 |
| 2 | 145 | 3.0 | 0 | 9 | 6.0 | 36.0 | 7 |
| 3 | 3 | 0.0 | 0 | 0 | 0.0 | 0.0 | 2 |
| 4 | 5 | 0.0 | 0 | 0 | 0.0 | 1.0 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 15187 | 137 | 3.0 | 6 | 20 | 8.0 | 15.0 | 0 |
| 15188 | 6 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 15189 | 39 | 0.0 | 0 | 16 | 0.0 | 3.0 | 1 |
| 15190 | 50 | 1.0 | 2 | 18 | 0.0 | 8.0 | 4 |
| 15191 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

```
[15192 rows x 19 columns]
```

Since we summed everything, we just need to change the NL column back. We can divide each value by itself to get either 1 or 0 like we had before.

```
[195]: df['NL'] = np.where(df['NL'] > 0, 1, 0)
```

```
[196]: tensor = df.drop(columns=['retroID'])
```

```
[197]: tensor
```

```
[197]:
```

| | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | S0 | \ |
|---|------|-------|------|------|-----|----|-----|------|-----|------|------|------|---|
| 0 | 331 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 2 | |
| 1 | 3298 | 12364 | 2174 | 3771 | 624 | 98 | 755 | 2297 | 240 | 73.0 | 1402 | 1383 | |
| 2 | 437 | 944 | 102 | 216 | 42 | 6 | 13 | 94 | 9 | 8.0 | 86 | 145 | |
| 3 | 448 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 3 | |
| 4 | 15 | 21 | 1 | 2 | 0 | 0 | 0 | 0 | 0 | 1.0 | 4 | 5 | |

| | | | | | | | | | | | | |
|-------|-----|-----|-----|-----|-----|----|-----|-----|-----|-----|-----|-----|
| ... | ... | ... | ... | ... | ... | .. | ... | ... | ... | ... | ... | ... |
| 15187 | 319 | 795 | 99 | 199 | 47 | 4 | 7 | 80 | 7 | 5.0 | 57 | 137 |
| 15188 | 16 | 18 | 3 | 3 | 1 | 0 | 0 | 0 | 0 | 0.0 | 2 | 6 |
| 15189 | 266 | 142 | 5 | 21 | 2 | 1 | 0 | 7 | 0 | 1.0 | 9 | 39 |
| 15190 | 209 | 491 | 41 | 109 | 17 | 2 | 2 | 20 | 2 | 0.0 | 34 | 50 |
| 15191 | 70 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 |

| | | | | | | |
|-------|-------|-----|----|-------|-------|----|
| | IBB | HBP | SH | SF | GIDP | NL |
| 0 | 0.0 | 0 | 1 | 0.0 | 0.0 | 1 |
| 1 | 293.0 | 32 | 21 | 121.0 | 328.0 | 1 |
| 2 | 3.0 | 0 | 9 | 6.0 | 36.0 | 1 |
| 3 | 0.0 | 0 | 0 | 0.0 | 0.0 | 1 |
| 4 | 0.0 | 0 | 0 | 0.0 | 1.0 | 1 |
| ... | ... | ... | .. | ... | ... | .. |
| 15187 | 3.0 | 6 | 20 | 8.0 | 15.0 | 0 |
| 15188 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |
| 15189 | 0.0 | 0 | 16 | 0.0 | 3.0 | 1 |
| 15190 | 1.0 | 2 | 18 | 0.0 | 8.0 | 1 |
| 15191 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0 |

[15192 rows x 18 columns]

```
[198]: tensor.to_csv('../output/tensor.csv')
        metadata.to_csv('../output/metadata.csv')
```

We now have a tensor with only relevant information, an indexing dictionary to get the player for each row, and a (soon to be expanded) metadata table to get more information on each player.

fielding_pre

March 9, 2020

```
[47]: import numpy as np
import pandas as pd
pd.options.mode.chained_assignment = None # default='warn'

[48]: df = pd.read_csv('../data/lahman/mlb_data/Fielding.csv').sort_values('playerID')

[49]: # This will be exported to a separate module
ids = pd.read_csv('../data/lahman/mlb_data/People.csv')
ids = ids[['playerID', 'retroID']]
id_dict = ids.set_index('playerID').to_dict()['retroID']

def get_retroid(id):
    return id_dict[id] if id_dict is not None else id

[50]: df['playerID'] = df['playerID'].apply(get_retroid)
df.rename(columns={'playerID': 'retroID'}, inplace=True)
```

Exploration

```
[51]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 112837 entries, 85308 to 106797
Data columns (total 18 columns):
#   Column      Non-Null Count  Dtype
---  -
0   retroID     112837 non-null  object
1   yearID      112837 non-null  int64
2   stint       112837 non-null  int64
3   teamID      112837 non-null  object
4   lgID        112837 non-null  object
5   POS         112837 non-null  object
6   G           112837 non-null  int64
7   GS          89431 non-null   float64
8   InnOuts     89431 non-null   float64
9   PO          112837 non-null  int64
10  A           112837 non-null  int64
11  E           112836 non-null  float64
```

```

12  DP      112837 non-null  int64
13  PB      8538 non-null   float64
14  WP      1169 non-null   float64
15  SB      6389 non-null   float64
16  CS      6389 non-null   float64
17  ZR      1169 non-null   float64
dtypes: float64(8), int64(6), object(4)
memory usage: 16.4+ MB

```

```
[52]: df.shape
```

```
[52]: (112837, 18)
```

```
[53]: df.columns
```

```
[53]: Index(['retroID', 'yearID', 'stint', 'teamID', 'lgID', 'POS', 'G', 'GS',
        'InnOuts', 'PO', 'A', 'E', 'DP', 'PB', 'WP', 'SB', 'CS', 'ZR'],
        dtype='object')
```

We want to get rid of columns which already exist in the Batting DataFrame (with which we will be merging this)

```
[54]: columns_to_drop = ['stint', 'teamID', 'lgID', 'G']
```

```
[55]: df.drop(columns=columns_to_drop, inplace=True)
```

```
[56]: df.head()
```

```
[56]:
```

| | retroID | yearID | POS | GS | InnOuts | PO | A | E | DP | PB | WP | SB | CS | ZR |
|--------|----------|--------|-----|-----|---------|----|---|-----|----|-----|-----|-----|-----|-----|
| 85308 | aardd001 | 2004 | P | 0.0 | 32.0 | 0 | 0 | 0.0 | 0 | NaN | NaN | NaN | NaN | NaN |
| 101187 | aardd001 | 2013 | P | 0.0 | 119.0 | 1 | 5 | 0.0 | 0 | NaN | NaN | NaN | NaN | NaN |
| 99344 | aardd001 | 2012 | P | 0.0 | 3.0 | 0 | 0 | 0.0 | 0 | NaN | NaN | NaN | NaN | NaN |
| 95793 | aardd001 | 2010 | P | 0.0 | 149.0 | 2 | 3 | 1.0 | 0 | NaN | NaN | NaN | NaN | NaN |
| 104866 | aardd001 | 2015 | P | 0.0 | 92.0 | 0 | 1 | 1.0 | 0 | NaN | NaN | NaN | NaN | NaN |

Cleaning and Preprocessing

We see a lot of NaNs in the last 5 columns. According to the Lahman readme, these are:

- PB - Passed Balls (by catchers)
- WP - Wild Pitches (by catchers)
- SB - Opponent Stolen Bases (by catchers)
- CS - Opponents Caught Stealing (by catchers)
- ZR - Zone Rating

It looks like the data demands that we treat catchers separately from other position players. This intuitively makes sense from what we know about baseball, and it saves us from getting rid of a lot of data. First, though, let's look at how much of that data is missing if we JUST look at catchers.

```
[57]: df_catchers = df[df['POS'] == 'C']
```

```
[58]: # Get missing data in the catchers category as a percentage
100 * df_catchers.isnull().sum() / len(df)
```

```
[58]: retroID      0.000000
yearID      0.000000
POS         0.000000
GS          1.901858
InnOuts     1.901858
PO          0.000000
A           0.000000
E           0.000000
DP          0.000000
PB          0.000000
WP          6.530659
SB          1.904517
CS          1.904517
ZR          6.530659
dtype: float64
```

Most of the percentages are negligible, but we can take a look at WP and ZR and see if the missing data is from early years.

```
[59]: early_catchers = df_catchers[df_catchers['yearID'] < 1955]
```

```
[60]: 100 * early_catchers.isnull().sum() / len(df)
```

```
[60]: retroID      0.000000
yearID      0.000000
POS         0.000000
GS          1.901858
InnOuts     1.901858
PO          0.000000
A           0.000000
E           0.000000
DP          0.000000
PB          0.000000
WP          1.901858
SB          1.901858
CS          1.901858
ZR          1.901858
dtype: float64
```

Definitely not the case. Let's try to narrow down where the issue is.

```
[61]: post1985_catchers = df_catchers[df_catchers['yearID'] > 1985]
```

```
[62]: 100 * post1985_catchers.isnull().sum() / len(df)
```

```
[62]: retroID    0.000000
      yearID    0.000000
      POS      0.000000
      GS       0.000000
      InnOuts   0.000000
      PO       0.000000
      A        0.000000
      E        0.000000
      DP       0.000000
      PB       0.000000
      WP       3.265773
      SB       0.000000
      CS       0.000000
      ZR       3.265773
      dtype: float64
```

```
[63]: df_1955_to_1986_catchers = df_catchers[(df_catchers['yearID'] >= 1955) &
      ↪(df_catchers['yearID'] <= 1985)]
```

```
[64]: 100 * df_1955_to_1986_catchers.isnull().sum() / len(df)
```

```
[64]: retroID    0.000000
      yearID    0.000000
      POS      0.000000
      GS       0.000000
      InnOuts   0.000000
      PO       0.000000
      A        0.000000
      E        0.000000
      DP       0.000000
      PB       0.000000
      WP       1.363028
      SB       0.002659
      CS       0.002659
      ZR       1.363028
      dtype: float64
```

```
[65]: pre_1930_catchers = df_catchers[df_catchers['yearID'] < 1930]
```

```
[66]: 100 * pre_1930_catchers.isnull().sum() / len(df)
```

```
[66]: retroID    0.000000
      yearID    0.000000
      POS      0.000000
      GS       0.591118
```

```

InnOuts    0.591118
PO         0.000000
A          0.000000
E          0.000000
DP         0.000000
PB         0.000000
WP         0.591118
SB         0.591118
CS         0.591118
ZR         0.591118
dtype: float64

```

We see that the issue is mainly in the very early years, and we are fine with dropping that information by just filling it in as we did in the Batters table.

So with that, we are fine with filling all NA values with 0.

```

[67]: df_catchers['GS'].fillna(value=0, inplace=True)
      df_catchers['InnOuts'].fillna(value=0, inplace=True)
      df_catchers['WP'].fillna(value=0, inplace=True)
      df_catchers['SB'].fillna(value=0, inplace=True)
      df_catchers['CS'].fillna(value=0, inplace=True)
      df_catchers['ZR'].fillna(value=0, inplace=True)

```

```

[68]: df['GS'].fillna(value=0, inplace=True)
      df['InnOuts'].fillna(value=0, inplace=True)
      #We can just drop the catcher-related columns from the original dataframe, as we
      →will also drop all catcher rows
      catcher_columns = ['PB', 'WP', 'SB', 'CS', 'ZR']
      df.drop(columns=catcher_columns, inplace=True)

```

Now drop all catcher rows so we have two separate dataframes, and get rid of the yearID column which we're done with and will be useless after aggregation.

```

[69]: df = df[df['POS'] != 'C']

```

```

[70]: df.drop(columns=['yearID'], inplace=True)
      df_catchers.drop(columns=['yearID'], inplace=True)

```

```

[71]: df.shape

```

```

[71]: (104299, 8)

```

```

[72]: df_catchers.shape

```

```

[72]: (8538, 13)

```

```

[73]: 100 * df.isnull().sum() / len(df)

```

```
[73]: retroID    0.000000
      POS       0.000000
      GS        0.000000
      InnOuts    0.000000
      PO        0.000000
      A         0.000000
      E         0.000959
      DP        0.000000
      dtype: float64
```

Now we just see a little bit of information missing from Errors, so we can fill that with 0s no problem.

```
[74]: df['E'].fillna(value=0, inplace=True)
```

```
[75]: 100 * df.isnull().sum() / len(df)
```

```
[75]: retroID    0.0
      POS       0.0
      GS        0.0
      InnOuts    0.0
      PO        0.0
      A         0.0
      E         0.0
      DP        0.0
      dtype: float64
```

```
[76]: 100 * df_catchers.isnull().sum() / len(df)
```

```
[76]: retroID    0.0
      POS       0.0
      GS        0.0
      InnOuts    0.0
      PO        0.0
      A         0.0
      E         0.0
      DP        0.0
      PB        0.0
      WP        0.0
      SB        0.0
      CS        0.0
      ZR        0.0
      dtype: float64
```

Aggregation

Now we just need to aggregate all stats to get total career numbers for each player.

```
[77]: df = df.groupby('retroID').sum().reset_index()
```

```
[78]: df_catchers = df_catchers.groupby('retroID').sum().reset_index()
```

```
[79]: df
```

```
[79]:
```

| | retroID | GS | InnOuts | P0 | A | E | DP |
|-------|----------|--------|---------|------|-----|-------|-----|
| 0 | aardd001 | 0.0 | 1011.0 | 11 | 29 | 3.0 | 2 |
| 1 | aaroh101 | 2977.0 | 78414.0 | 7436 | 429 | 144.0 | 218 |
| 2 | aarot101 | 206.0 | 6472.0 | 1317 | 113 | 22.0 | 124 |
| 3 | aased001 | 91.0 | 3328.0 | 67 | 135 | 13.0 | 10 |
| 4 | abada001 | 4.0 | 138.0 | 37 | 1 | 1.0 | 3 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 14222 | zumaj001 | 0.0 | 629.0 | 7 | 14 | 2.0 | 1 |
| 14223 | zupcb001 | 198.0 | 5842.0 | 483 | 22 | 12.0 | 5 |
| 14224 | zuveg101 | 31.0 | 1847.0 | 45 | 145 | 7.0 | 10 |
| 14225 | zuvep001 | 136.0 | 3844.0 | 267 | 415 | 23.0 | 84 |
| 14226 | zycht001 | 1.0 | 218.0 | 1 | 6 | 1.0 | 0 |

[14227 rows x 7 columns]

```
[80]: df_catchers
```

```
[80]:
```

| | retroID | GS | InnOuts | P0 | A | E | DP | PB | WP | SB | CS | \ |
|------|----------|-------|---------|------|-----|------|-----|------|------|-------|------|---|
| 0 | adamb105 | 1.0 | 27.0 | 6 | 0 | 0.0 | 0 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 1 | adamb106 | 0.0 | 0.0 | 249 | 90 | 12.0 | 15 | 7.0 | 0.0 | 0.0 | 0.0 | |
| 2 | adamd101 | 3.0 | 78.0 | 9 | 2 | 0.0 | 0 | 1.0 | 0.0 | 0.0 | 0.0 | |
| 3 | adled101 | 65.0 | 1840.0 | 453 | 26 | 4.0 | 2 | 8.0 | 19.0 | 37.0 | 16.0 | |
| 4 | afent001 | 20.0 | 613.0 | 123 | 5 | 1.0 | 3 | 6.0 | 0.0 | 17.0 | 3.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1524 | zimmd101 | 27.0 | 744.0 | 150 | 18 | 6.0 | 1 | 5.0 | 12.0 | 10.0 | 10.0 | |
| 1525 | zimmj101 | 298.0 | 8560.0 | 2131 | 150 | 21.0 | 26 | 19.0 | 84.0 | 110.0 | 80.0 | |
| 1526 | zinta001 | 0.0 | 3.0 | 2 | 0 | 0.0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 1527 | zunim001 | 535.0 | 14489.0 | 4356 | 264 | 21.0 | 22 | 39.0 | 0.0 | 248.0 | 98.0 | |
| 1528 | zupof101 | 1.0 | 114.0 | 31 | 1 | 2.0 | 0 | 1.0 | 1.0 | 2.0 | 1.0 | |
| | ZR | | | | | | | | | | | |
| 0 | 0.0 | | | | | | | | | | | |
| 1 | 0.0 | | | | | | | | | | | |
| 2 | 0.0 | | | | | | | | | | | |
| 3 | 0.0 | | | | | | | | | | | |
| 4 | 0.0 | | | | | | | | | | | |
| ... | ... | | | | | | | | | | | |
| 1524 | 3.0 | | | | | | | | | | | |
| 1525 | 4.0 | | | | | | | | | | | |
| 1526 | 0.0 | | | | | | | | | | | |
| 1527 | 0.0 | | | | | | | | | | | |

```
1528 0.0
```

```
[1529 rows x 12 columns]
```

```
[ ]:
```


add_advanced_batting

April 29, 2020

```
[76]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[30]: df = pd.read_csv('../core/output/batters.csv')
df_adv = pd.read_csv('../core/output/advanced_batting.csv')
```

Adding Advanced Stats

We will use a combination of wOBA, wRC+ and WAR as our overall rating - our Y value.

```
[43]: df_adv.sort_values('retroID')
```

```
[43]:
```

| | retroID | wOBA | wRC+ | WAR |
|-------|----------|-------|--------|-------|
| 9203 | aardd001 | 0.000 | -100.0 | -0.1 |
| 3 | aaroh101 | 0.403 | 153.0 | 136.3 |
| 13920 | aarot101 | 0.282 | 76.0 | -1.7 |
| 9158 | aased001 | 0.000 | -100.0 | -0.1 |
| 11841 | abada001 | 0.184 | 0.0 | -0.4 |
| ... | ... | ... | ... | ... |
| 13227 | zupcb001 | 0.293 | 74.0 | -0.9 |
| 10487 | zupof101 | 0.225 | 37.0 | -0.2 |
| 11591 | zuveg101 | 0.179 | 0.0 | -0.3 |
| 14134 | zuvep001 | 0.254 | 52.0 | -2.2 |
| 6759 | zycht001 | 0.000 | NaN | 0.0 |

[14399 rows x 4 columns]

```
[42]: df
```

```
[42]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | \ |
|-------|----------|----------|--------|-----------|-----------|--------|--------|---|
| 0 | aardd001 | 0.569672 | 0.60 | 2004 | 2015 | 0 | 0 | |
| 1 | aaroh101 | 0.426230 | 0.45 | 1954 | 1976 | 0 | 0 | |
| 2 | aarot101 | 0.467213 | 0.60 | 1962 | 1971 | 1 | 0 | |
| 3 | aased001 | 0.467213 | 0.60 | 1977 | 1990 | 0 | 0 | |
| 4 | abada001 | 0.442623 | 0.50 | 2001 | 2006 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | zupcb001 | 0.590164 | 0.65 | 1991 | 1994 | 0 | 0 | |
| 15289 | zupof101 | 0.434426 | 0.40 | 1957 | 1961 | 0 | 0 | |

| | | | | | | | |
|-------|----------|----------|------|------|------|---|---|
| 15290 | zuveg101 | 0.487705 | 0.65 | 1951 | 1959 | 0 | 0 |
| 15291 | zuvep001 | 0.397541 | 0.45 | 1982 | 1991 | 0 | 0 |
| 15292 | zycht001 | 0.467213 | 0.60 | 2015 | 2017 | 0 | 0 |

| | pos_3B | pos_C | pos_OF | ... | SB | CS | BB | S0 | IBB | HBP | SH | SF | \ |
|-------|--------|-------|--------|-----|-----|------|------|------|-----|-----|-----|-----|---|
| 0 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | 2 | 0 | 0 | 1 | 0 | |
| 1 | 0 | 0 | 1 | ... | 240 | 73.0 | 1402 | 1383 | 293 | 32 | 21 | 121 | |
| 2 | 0 | 0 | 0 | ... | 9 | 8.0 | 86 | 145 | 3 | 0 | 9 | 6 | |
| 3 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | 3 | 0 | 0 | 0 | 0 | |
| 4 | 0 | 0 | 0 | ... | 0 | 1.0 | 4 | 5 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0 | 0 | 1 | ... | 7 | 5.0 | 57 | 137 | 3 | 6 | 20 | 8 | |
| 15289 | 0 | 1 | 0 | ... | 0 | 0.0 | 2 | 6 | 0 | 0 | 0 | 0 | |
| 15290 | 0 | 0 | 0 | ... | 0 | 1.0 | 9 | 39 | 0 | 0 | 16 | 0 | |
| 15291 | 0 | 0 | 0 | ... | 2 | 0.0 | 34 | 50 | 1 | 2 | 18 | 0 | |
| 15292 | 0 | 0 | 0 | ... | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | 0 | |

| | GIDP | NL |
|-------|------|-----|
| 0 | 0 | 1 |
| 1 | 328 | 1 |
| 2 | 36 | 1 |
| 3 | 0 | 1 |
| 4 | 1 | 1 |
| ... | ... | ... |
| 15288 | 15 | 0 |
| 15289 | 0 | 0 |
| 15290 | 3 | 1 |
| 15291 | 8 | 1 |
| 15292 | 0 | 0 |

[15293 rows x 36 columns]

```
[35]: df.shape
```

```
[35]: (15293, 36)
```

```
[36]: df_adv.shape
```

```
[36]: (14399, 4)
```

```
[44]: df = df.merge(df_adv, how='left')
```

```
[46]: df['wOBA'].fillna(0, inplace=True)
df['wRC+'].fillna(0, inplace=True)
df['WAR'].fillna(0, inplace=True)
```

```
[47]: df
```

```
[47]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | \ |
|-------|----------|----------|--------|-----------|-----------|--------|--------|---|
| 0 | aardd001 | 0.569672 | 0.60 | 2004 | 2015 | 0 | 0 | |
| 1 | aaroh101 | 0.426230 | 0.45 | 1954 | 1976 | 0 | 0 | |
| 2 | aarot101 | 0.467213 | 0.60 | 1962 | 1971 | 1 | 0 | |
| 3 | aased001 | 0.467213 | 0.60 | 1977 | 1990 | 0 | 0 | |
| 4 | abada001 | 0.442623 | 0.50 | 2001 | 2006 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | zupcb001 | 0.590164 | 0.65 | 1991 | 1994 | 0 | 0 | |
| 15289 | zupof101 | 0.434426 | 0.40 | 1957 | 1961 | 0 | 0 | |
| 15290 | zuveg101 | 0.487705 | 0.65 | 1951 | 1959 | 0 | 0 | |
| 15291 | zuvep001 | 0.397541 | 0.45 | 1982 | 1991 | 0 | 0 | |
| 15292 | zycht001 | 0.467213 | 0.60 | 2015 | 2017 | 0 | 0 | |

| | pos_3B | pos_C | pos_OF | ... | SO | IBB | HBP | SH | SF | GIDP | NL | wOBA | \ |
|-------|--------|-------|--------|-----|------|-----|-----|-----|-----|------|-----|-------|---|
| 0 | 0 | 0 | 0 | ... | 2 | 0 | 0 | 1 | 0 | 0 | 1 | 0.000 | |
| 1 | 0 | 0 | 1 | ... | 1383 | 293 | 32 | 21 | 121 | 328 | 1 | 0.403 | |
| 2 | 0 | 0 | 0 | ... | 145 | 3 | 0 | 9 | 6 | 36 | 1 | 0.282 | |
| 3 | 0 | 0 | 0 | ... | 3 | 0 | 0 | 0 | 0 | 0 | 1 | 0.000 | |
| 4 | 0 | 0 | 0 | ... | 5 | 0 | 0 | 0 | 0 | 1 | 1 | 0.184 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0 | 0 | 1 | ... | 137 | 3 | 6 | 20 | 8 | 15 | 0 | 0.293 | |
| 15289 | 0 | 1 | 0 | ... | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 0.225 | |
| 15290 | 0 | 0 | 0 | ... | 39 | 0 | 0 | 16 | 0 | 3 | 1 | 0.179 | |
| 15291 | 0 | 0 | 0 | ... | 50 | 1 | 2 | 18 | 0 | 8 | 1 | 0.254 | |
| 15292 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.000 | |

| | wRC+ | WAR |
|-------|--------|-------|
| 0 | -100.0 | -0.1 |
| 1 | 153.0 | 136.3 |
| 2 | 76.0 | -1.7 |
| 3 | -100.0 | -0.1 |
| 4 | 0.0 | -0.4 |
| ... | ... | ... |
| 15288 | 74.0 | -0.9 |
| 15289 | 37.0 | -0.2 |
| 15290 | 0.0 | -0.3 |
| 15291 | 52.0 | -2.2 |
| 15292 | 0.0 | 0.0 |

[15293 rows x 39 columns]

For now, we're just going to take the mean of the three most accepted advanced statistics, giving them equal importance. This will lead to a model that favors offense over defense, as WAR is the only stat that takes defense into account, but that's fine.

```
[50]: df['Batting'] = df[['wOBA', 'wRC+', 'WAR']].mean(axis=1).round(3)
```

```
[51]: df['Batting']
```

```
[51]: 0      -33.367
      1       96.568
      2       24.861
      3     -33.367
      4      -0.072
      ...
     15288    24.464
     15289    12.342
     15290    -0.040
     15291    16.685
     15292     0.000
      Name: Rating, Length: 15293, dtype: float64
```

```
[81]: df['Batting'].mean()
```

```
[81]: 11.660071993722617
```

```
[82]: df['Batting'].min()
```

```
[82]: -33.5
```

```
[83]: df['Batting'].max()
```

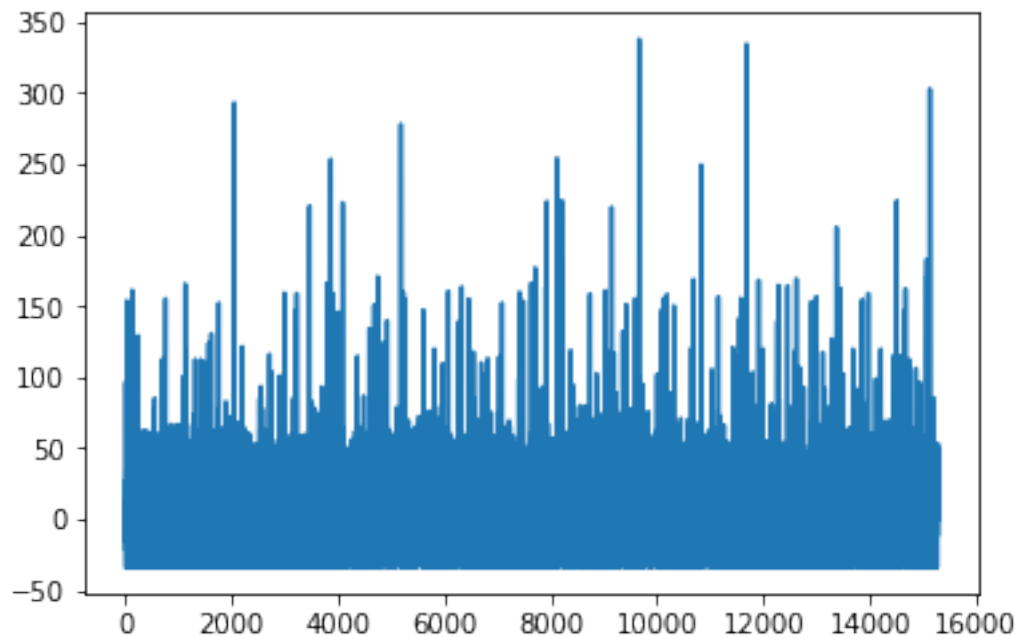
```
[83]: 337.916
```

Normalization

The Batting stat now has a very wide range which seems to trend more toward the lower end. We need to normalize the statistic so that our sigmoid output will be able to accurately predict it. For this reason, we'll use min-max normalization to get a range [0, 1].

```
[77]: plt.plot(df['Batting'])
```

```
[77]: [<matplotlib.lines.Line2D at 0x12a2a5190>]
```

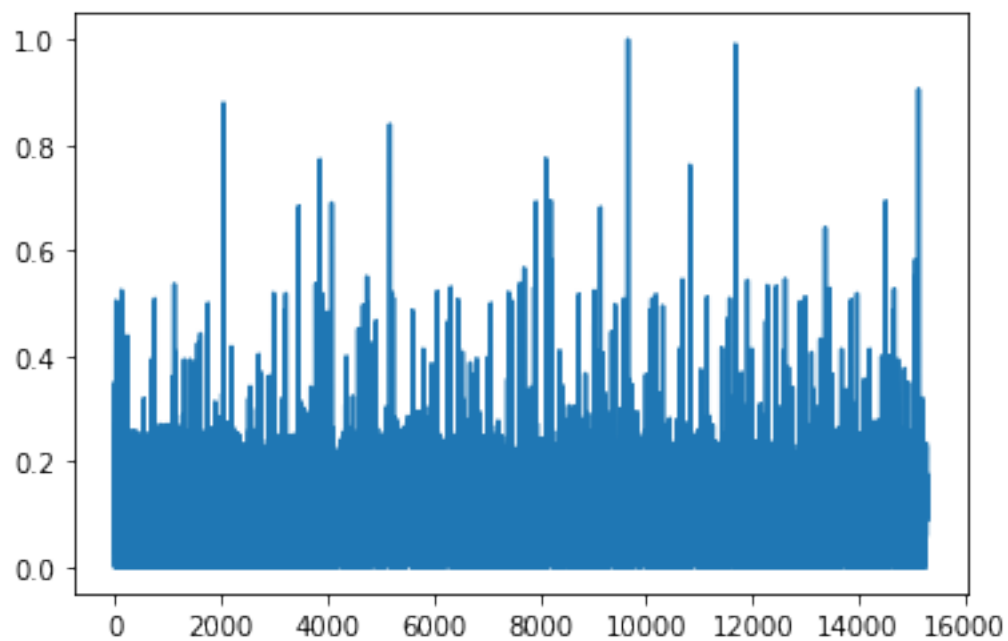


```
[84]: from sklearn.preprocessing import MinMaxScaler
```

```
[85]: scaler = MinMaxScaler()
```

```
[86]: plt.plot(scaler.fit_transform(df[['Batting']]))
```

```
[86]: [<matplotlib.lines.Line2D at 0x121311650>]
```



```
[88]: df['Batting'] = scaler.fit_transform(df[['Batting']])
```

```
[89]: df
```

```
[89]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | \ |
|-------|----------|----------|--------|-----------|-----------|--------|--------|---|
| 0 | aardd001 | 0.569672 | 0.60 | 2004 | 2015 | 0 | 0 | |
| 1 | aaroh101 | 0.426230 | 0.45 | 1954 | 1976 | 0 | 0 | |
| 2 | aarot101 | 0.467213 | 0.60 | 1962 | 1971 | 1 | 0 | |
| 3 | aased001 | 0.467213 | 0.60 | 1977 | 1990 | 0 | 0 | |
| 4 | abada001 | 0.442623 | 0.50 | 2001 | 2006 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | zupcb001 | 0.590164 | 0.65 | 1991 | 1994 | 0 | 0 | |
| 15289 | zupof101 | 0.434426 | 0.40 | 1957 | 1961 | 0 | 0 | |
| 15290 | zuveg101 | 0.487705 | 0.65 | 1951 | 1959 | 0 | 0 | |
| 15291 | zuvep001 | 0.397541 | 0.45 | 1982 | 1991 | 0 | 0 | |
| 15292 | zycht001 | 0.467213 | 0.60 | 2015 | 2017 | 0 | 0 | |

| | pos_3B | pos_C | pos_OF | ... | IBB | HBP | SH | SF | GIDP | NL | wOBA | wRC+ | \ |
|-------|--------|-------|--------|-----|-----|-----|-----|-----|------|-----|-------|--------|---|
| 0 | 0 | 0 | 0 | ... | 0 | 0 | 1 | 0 | 0 | 1 | 0.000 | -100.0 | |
| 1 | 0 | 0 | 1 | ... | 293 | 32 | 21 | 121 | 328 | 1 | 0.403 | 153.0 | |
| 2 | 0 | 0 | 0 | ... | 3 | 0 | 9 | 6 | 36 | 1 | 0.282 | 76.0 | |
| 3 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 1 | 0.000 | -100.0 | |
| 4 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 1 | 1 | 0.184 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0 | 0 | 1 | ... | 3 | 6 | 20 | 8 | 15 | 0 | 0.293 | 74.0 | |
| 15289 | 0 | 1 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0.225 | 37.0 | |
| 15290 | 0 | 0 | 0 | ... | 0 | 0 | 16 | 0 | 3 | 1 | 0.179 | 0.0 | |
| 15291 | 0 | 0 | 0 | ... | 1 | 2 | 18 | 0 | 8 | 1 | 0.254 | 52.0 | |
| 15292 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0.000 | 0.0 | |

| | WAR | Rating |
|-------|-------|----------|
| 0 | -0.1 | 0.000358 |
| 1 | 136.3 | 0.350195 |
| 2 | -1.7 | 0.157131 |
| 3 | -0.1 | 0.000358 |
| 4 | -0.4 | 0.090002 |
| ... | ... | ... |
| 15288 | -0.9 | 0.156062 |
| 15289 | -0.2 | 0.123425 |
| 15290 | -0.3 | 0.090088 |
| 15291 | -2.2 | 0.135118 |
| 15292 | 0.0 | 0.090195 |

```
[15293 rows x 40 columns]
```

We now have the Y value that our NN should attempt to predict. We'll keep wOBA, wRC+ and WAR as columns at this point so we can decide later if they need to come out.

[]:

pitching_pre

April 28, 2020

```
[211]: import pandas as pd
import numpy as np
pd.options.mode.chained_assignment = None # default='warn'
```

```
[212]: df = pd.read_csv('../data/lahman/mlb_data/Pitching.csv')
```

```
[213]: # This will be exported to a separate module
ids = pd.read_csv('../data/lahman/mlb_data/People.csv')
ids = ids[['playerID', 'retroID']]
id_dict = ids.set_index('playerID').to_dict()['retroID']

def get_retroid(id):
    return id_dict[id] if id_dict is not None else id
```

```
[214]: df['playerID'] = df['playerID'].apply(get_retroid)
df.rename(columns={'playerID': 'retroID'}, inplace=True)
```

Exploration

```
[215]: df.head()
```

```
[215]:
```

| | retroID | yearID | stint | teamID | lgID | W | L | G | GS | CG | ... | IBB | WP | HBP | \ |
|---|----------|--------|-------|--------|------|----|----|----|----|----|-----|-----|----|-----|---|
| 0 | adamb104 | 1919 | 1 | PIT | NL | 17 | 10 | 34 | 29 | 23 | ... | NaN | 2 | 3 | |
| 1 | adamw101 | 1919 | 1 | PHA | AL | 0 | 0 | 1 | 0 | 0 | ... | NaN | 0 | 1 | |
| 2 | alexg102 | 1919 | 1 | CHN | NL | 16 | 11 | 30 | 27 | 20 | ... | NaN | 1 | 0 | |
| 3 | altrn101 | 1919 | 1 | WS1 | AL | 0 | 0 | 1 | 0 | 0 | ... | NaN | 0 | 0 | |
| 4 | amesr101 | 1919 | 1 | SLN | NL | 3 | 5 | 23 | 7 | 1 | ... | NaN | 3 | 1 | |

| | BK | BFP | GF | R | SH | SF | GIDP |
|---|----|--------|----|----|-----|-----|------|
| 0 | 0 | 1017.0 | 5 | 66 | NaN | NaN | NaN |
| 1 | 0 | 21.0 | 1 | 2 | NaN | NaN | NaN |
| 2 | 0 | 906.0 | 3 | 51 | NaN | NaN | NaN |
| 3 | 0 | 4.0 | 0 | 4 | NaN | NaN | NaN |
| 4 | 0 | 314.0 | 10 | 44 | NaN | NaN | NaN |

[5 rows x 30 columns]


```
[216]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40372 entries, 0 to 40371
Data columns (total 30 columns):
#   Column      Non-Null Count  Dtype
---  -
0   retroID     40372 non-null   object
1   yearID      40372 non-null   int64
2   stint       40372 non-null   int64
3   teamID      40372 non-null   object
4   lgID        40372 non-null   object
5   W           40372 non-null   int64
6   L           40372 non-null   int64
7   G           40372 non-null   int64
8   GS          40372 non-null   int64
9   CG          40372 non-null   int64
10  SHO         40372 non-null   int64
11  SV          40372 non-null   int64
12  IPouts      40372 non-null   int64
13  H           40372 non-null   int64
14  ER          40372 non-null   int64
15  HR          40372 non-null   int64
16  BB          40372 non-null   int64
17  SO          40372 non-null   int64
18  BAOpp       40360 non-null   float64
19  ERA         40298 non-null   float64
20  IBB         32121 non-null   float64
21  WP          40372 non-null   int64
22  HBP         40372 non-null   int64
23  BK          40372 non-null   int64
24  BFP         40369 non-null   float64
25  GF          40372 non-null   int64
26  R           40372 non-null   int64
27  SH          27512 non-null   float64
28  SF          27512 non-null   float64
29  GIDP        26381 non-null   float64
dtypes: float64(7), int64(20), object(3)
memory usage: 9.2+ MB
```

```
[217]: df.columns
```

```
[217]: Index(['retroID', 'yearID', 'stint', 'teamID', 'lgID', 'W', 'L', 'G', 'GS',
        'CG', 'SHO', 'SV', 'IPouts', 'H', 'ER', 'HR', 'BB', 'SO', 'BAOpp',
        'ERA', 'IBB', 'WP', 'HBP', 'BK', 'BFP', 'GF', 'R', 'SH', 'SF', 'GIDP'],
        dtype='object')
```

```
[218]: columns_to_drop = ['stint', 'teamID', 'lgID']
```

```
[219]: df.drop(columns=columns_to_drop, inplace=True)
```

```
[220]: df.shape
```

```
[220]: (40372, 27)
```

Cleaning and Preprocessing

```
[221]: 100 * df.isnull().sum() / len(df)
```

```
[221]: retroID      0.000000
      yearID      0.000000
      W          0.000000
      L          0.000000
      G          0.000000
      GS         0.000000
      CG         0.000000
      SHO        0.000000
      SV         0.000000
      IPouts     0.000000
      H          0.000000
      ER         0.000000
      HR         0.000000
      BB         0.000000
      SO         0.000000
      BAOpp      0.029724
      ERA        0.183295
      IBB       20.437432
      WP         0.000000
      HBP        0.000000
      BK         0.000000
      BFP        0.007431
      GF         0.000000
      R          0.000000
      SH        31.853760
      SF        31.853760
      GIDP       34.655207
      dtype: float64
```

```
[222]: df_early = df[df['yearID'] <= 1930]
```

```
[223]: 100 * df_early.isnull().sum() / len(df)
```

```
[223]: retroID      0.000000
      yearID      0.000000
```

| | |
|--------|----------|
| W | 0.000000 |
| L | 0.000000 |
| G | 0.000000 |
| GS | 0.000000 |
| CG | 0.000000 |
| SHO | 0.000000 |
| SV | 0.000000 |
| IPouts | 0.000000 |
| H | 0.000000 |
| ER | 0.000000 |
| HR | 0.000000 |
| BB | 0.000000 |
| SO | 0.000000 |
| BAOpp | 0.002477 |
| ERA | 0.042108 |
| IBB | 6.442584 |
| WP | 0.000000 |
| HBP | 0.000000 |
| BK | 0.000000 |
| BFP | 0.007431 |
| GF | 0.000000 |
| R | 0.000000 |
| SH | 6.442584 |
| SF | 6.442584 |
| GIDP | 6.442584 |

dtype: float64

```
[224]: df_modern = df[df['yearID'] >= 1980]
```

```
[225]: 100 * df_modern.isnull().sum() / len(df)
```

```
[225]: retroID    0.000000
yearID    0.000000
W          0.000000
L          0.000000
G          0.000000
GS         0.000000
CG         0.000000
SHO        0.000000
SV          0.000000
IPouts     0.000000
H          0.000000
ER         0.000000
HR         0.000000
BB         0.000000
SO         0.000000
BAOpp      0.019816
```

```

ERA      0.056970
IBB      0.000000
WP       0.000000
HBP      0.000000
BK       0.000000
BFP      0.000000
GF       0.000000
R        0.000000
SH       0.000000
SF       0.000000
GIDP     0.000000
dtype: float64

```

Luckily the more modern data is barely missing any information.

```
[226]: df_mid = df[(df['yearID'] > 1935) & (df['yearID'] < 1975)]
```

```
[227]: 100 * df_mid.isnull().sum() / len(df)
```

```

[227]: retroID      0.000000
yearID      0.000000
W           0.000000
L           0.000000
G           0.000000
GS          0.000000
CG          0.000000
SHO         0.000000
SV          0.000000
IPouts      0.000000
H           0.000000
ER          0.000000
HR          0.000000
BB          0.000000
SO          0.000000
BAOpp       0.007431
ERA         0.066878
IBB         11.428713
WP          0.000000
HBP         0.000000
BK          0.000000
BFP         0.000000
GF          0.000000
R           0.000000
SH          22.845041
SF          22.845041
GIDP        25.646488
dtype: float64

```

We see that much of the lost data comes within this 40-year span. I think that given what the major missing information is - intentional bases on balls, sacrifice hits, sacrifice flies and grounded into double play - and the fact that these statistics are not often used as primary indicators of a pitcher's ability, coupled with the fact that it's mostly localized within less than half of our time frame, I can be forgiven for just filling these values as 0.

```
[228]: df['IBB'].fillna(0, inplace=True)
df['SH'].fillna(0, inplace=True)
df['SF'].fillna(0, inplace=True)
df['GIDP'].fillna(0, inplace=True)
```

```
[229]: 100 * df.isnull().sum() / len(df)
```

```
[229]: retroID      0.000000
yearID      0.000000
W           0.000000
L           0.000000
G           0.000000
GS          0.000000
CG          0.000000
SHO         0.000000
SV          0.000000
IPouts      0.000000
H           0.000000
ER          0.000000
HR          0.000000
BB          0.000000
SO          0.000000
BAOpp       0.029724
ERA         0.183295
IBB         0.000000
WP          0.000000
HBP         0.000000
BK          0.000000
BFP        0.007431
GF          0.000000
R           0.000000
SH          0.000000
SF          0.000000
GIDP        0.000000
dtype: float64
```

We're left with three fields that having missing data: opponents' batting average, earned run average and batters faced by pitcher. We'll have to do some data exploration on these because I don't want to just fill them with 0s.

Missing Values: BAOpp

```
[230]: df_baopp_missing = df[df['BAOpp'].isnull()]
```

```
[231]: df_baopp_missing.shape
```

```
[231]: (12, 27)
```

```
[232]: df_baopp_missing.sort_values('retroID')
```

```
[232]:
```

| | retroID | yearID | W | L | G | GS | CG | SHO | SV | IPouts | ... | IBB | WP | HBP | \ |
|-------|----------|--------|---|---|---|----|----|-----|----|--------|-----|-----|----|-----|---|
| 14000 | apodb101 | 1973 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 19114 | arrof001 | 1986 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 39581 | brotr001 | 2018 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 36447 | dunnj001 | 2014 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 2 | ... | 0.0 | 2 | 0 | |
| 38848 | eschj001 | 2017 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 3709 | fordw103 | 1936 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 297 | glasn101 | 1920 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 7 | ... | 0.0 | 0 | 1 | |
| 26791 | halts001 | 2000 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 27036 | radis001 | 2000 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 36189 | tolls002 | 2013 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 36208 | villb002 | 2013 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 13621 | younl101 | 1971 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |

| | BK | BFP | GF | R | SH | SF | GIDP |
|-------|----|------|----|---|-----|-----|------|
| 14000 | 0 | 2.0 | 0 | 1 | 0.0 | 0.0 | 0.0 |
| 19114 | 0 | 3.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 39581 | 0 | 2.0 | 0 | 1 | 0.0 | 0.0 | 0.0 |
| 36447 | 0 | 2.0 | 0 | 0 | 0.0 | 1.0 | 0.0 |
| 38848 | 0 | 2.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 3709 | 0 | 3.0 | 0 | 2 | 0.0 | 0.0 | 0.0 |
| 297 | 0 | 12.0 | 0 | 4 | 0.0 | 0.0 | 0.0 |
| 26791 | 0 | 1.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 27036 | 0 | 1.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 36189 | 0 | 2.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 36208 | 0 | 1.0 | 1 | 0 | 0.0 | 0.0 | 0.0 |
| 13621 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |

```
[12 rows x 27 columns]
```

Nobody appears in this table more than once. We'll hope that we can get career numbers for them and fill with the average. For anyone who only appears once I'll go with the league average, and I'll add a standard deviation since they probably we're exactly middle-of-the-road. Probably not the best way but it's only a few datapoints.

```
[233]: baopp_checks = df[df['retroID'].isin(df_baopp_missing['retroID'])].  
      ↪sort_values('retroID')
```

```
[234]: baopp_checks['retroID'].value_counts()
```

```
[234]: radis001    11
      arrof001    9
      brotr001    7
      tolls002    5
      apodb101    5
      villb002    4
      dunnj001    2
      eschj001    2
      halts001    2
      younl101    1
      glasn101    1
      fordw103    1
      Name: retroID, dtype: int64
```

We only have three data points with one year of appearances. We'll fill those with the league average.

```
[235]: val_counts = baopp_checks['retroID'].value_counts()
```

```
[236]: from itertools import compress
      # Get list of retroIDs of players who only have one year appearance
```

```
[237]: one_time_players = list(compress(val_counts.index, val_counts.eq(1)))
```

```
[238]: df[df['retroID'].isin(one_time_players)]
```

```
[238]:      retroID  yearID  W  L  G  GS  CG  SHO  SV  IPouts  ...  IBB  WP  HBP  \
297    glasn101    1920  0  0  1   0   0   0   0       7  ...  0.0  0   1
3709   fordw103    1936  0  0  1   0   0   0   0       0  ...  0.0  0   0
13621  younl101    1971  0  0  1   0   0   0   0       0  ...  0.0  0   0

      BK  BFP  GF  R  SH  SF  GIDP
297    0  12.0  0  4  0.0  0.0  0.0
3709    0   3.0  0  2  0.0  0.0  0.0
13621    0   0.0  0  0  0.0  0.0  0.0
```

[3 rows x 27 columns]

```
[239]: df['BAOpp'].mean()
```

```
[239]: 0.27445659068384537
```

```
[240]: df['BAOpp'].std()
```

```
[240]: 0.07751058079199835
```

```
[241]: filled_baopp = df['BAOpp'].mean() + df['BAOpp'].std()
```

```
[242]: filled_baopp
```

```
[242]: 0.3519671714758437
```

```
[243]: df.loc[df['retroID'].isin(one_time_players), ['BAOpp']] = filled_baopp
```

```
[244]: df[df['retroID'].isin(one_time_players)]['BAOpp']
```

```
[244]: 297      0.351967
3709      0.351967
13621     0.351967
Name: BAOpp, dtype: float64
```

```
[245]: df_baopp_missing = df[df['BAOpp'].isnull()].sort_values('retroID')
```

```
[246]: df_baopp_missing
```

```
[246]:      retroID  yearID  W  L  G  GS  CG  SHO  SV  IPouts  ...  IBB  WP  HBP  \
14000  apodb101   1973  0  0  1   0   0   0   0      0  ...  0.0  0   0
19114  arrof001   1986  0  0  1   0   0   0   0      0  ...  0.0  0   0
39581  brotr001   2018  0  0  1   0   0   0   0      0  ...  0.0  0   0
36447  dunnj001   2014  0  0  1   0   0   0   0      2  ...  0.0  2   0
38848  eschj001   2017  0  0  1   0   0   0   0      0  ...  0.0  0   0
26791  halts001   2000  0  0  1   0   0   0   0      0  ...  0.0  0   0
27036  radis001   2000  0  0  1   0   0   0   0      0  ...  0.0  0   0
36189  tolls002   2013  0  0  1   0   0   0   0      0  ...  0.0  0   0
36208  villb002   2013  0  0  1   0   0   0   0      0  ...  0.0  0   0
```

```
      BK  BFP  GF  R  SH  SF  GIDP
14000   0  2.0   0  1  0.0  0.0   0.0
19114   0  3.0   0  0  0.0  0.0   0.0
39581   0  2.0   0  1  0.0  0.0   0.0
36447   0  2.0   0  0  0.0  1.0   0.0
38848   0  2.0   0  0  0.0  0.0   0.0
26791   0  1.0   0  0  0.0  0.0   0.0
27036   0  1.0   0  0  0.0  0.0   0.0
36189   0  2.0   0  0  0.0  0.0   0.0
36208   0  1.0   1  0  0.0  0.0   0.0
```

```
[9 rows x 27 columns]
```

Now we just have to worry about the players with at least two years of appearances.

```
[247]: baopp_checks = df[df['retroID'].isin(df_baopp_missing['retroID'])].
      ↪sort_values('retroID')
```

```
[248]: baopp_checks['retroID'].value_counts()
```



```
[248]: radis001    11
      arrof001    9
      brotr001    7
      tolls002    5
      apodb101    5
      villb002    4
      dunnj001    2
      eschj001    2
      halts001    2
      Name: retroID, dtype: int64
```

```
[249]: one_time_players = list(val_counts.index)
```

```
[250]: one_time_players
```

```
[250]: ['radis001',
      'arrof001',
      'brotr001',
      'tolls002',
      'apodb101',
      'villb002',
      'dunnj001',
      'eschj001',
      'halts001',
      'younl101',
      'glasn101',
      'fordw103']
```

```
[251]: df[df['retroID'] == 'radis001']['BAOpp']
```

```
[251]: 21355    0.241
      21865    0.206
      22335    0.243
      22871    0.268
      23957    0.309
      24557    0.264
      25145    0.236
      25740    0.272
      26377    0.270
      27036     NaN
      27708    0.400
      Name: BAOpp, dtype: float64
```

```
[252]: df[df['retroID'] == 'radis001']['BAOpp'].mean().round(4)
```

```
[252]: 0.2709
```

This looks good, so we'll iterate through and assign each player's missing BAOpp as his career

mean for that state.

```
[253]: df['BAOpp'] = df.groupby("retroID")['BAOpp'].transform(lambda baopp: baopp.  
    ↳ fillna(baopp.mean()))
```

```
[254]: 100 * df.isnull().sum() / len(df)
```

```
[254]: retroID    0.000000  
yearID    0.000000  
W          0.000000  
L          0.000000  
G          0.000000  
GS         0.000000  
CG         0.000000  
SHO        0.000000  
SV         0.000000  
IPouts     0.000000  
H          0.000000  
ER         0.000000  
HR         0.000000  
BB         0.000000  
SO         0.000000  
BAOpp      0.000000  
ERA        0.183295  
IBB        0.000000  
WP         0.000000  
HBP        0.000000  
BK         0.000000  
BFP        0.007431  
GF         0.000000  
R          0.000000  
SH         0.000000  
SF         0.000000  
GIDP       0.000000  
dtype: float64
```

```
[255]: df.head()
```

```
[255]:   retroID  yearID  W  L  G  GS  CG  SHO  SV  IPouts  ...  IBB  WP  HBP  \  
0  adamb104    1919  17 10 34  29  23   6   1    790  ...  0.0  2   3  
1  adamw101    1919   0  0  1   0   0   0   0    14  ...  0.0  0   1  
2  alexg102    1919  16 11 30  27  20   9   1   705  ...  0.0  1   0  
3  altrn101    1919   0  0  1   0   0   0   0     0  ...  0.0  0   0  
4  amesr101    1919   3  5 23   7   1   0   1   210  ...  0.0  3   1  
  
   BK  BFP  GF  R  SH  SF  GIDP  
0   0  1017.0  5  66  0.0  0.0  0.0  
1   0   21.0  1  2  0.0  0.0  0.0
```

```

2  0  906.0  3  51  0.0  0.0  0.0
3  0    4.0  0   4  0.0  0.0  0.0
4  0  314.0 10  44  0.0  0.0  0.0

```

[5 rows x 27 columns]

Missing Values: ERA

```
[256]: df_era_missing = df[df['ERA'].isnull()].sort_values('retroID')
```

```
[257]: df_era_missing
```

```
[257]:
```

| | retroID | yearID | W | L | G | GS | CG | SHO | SV | IPouts | ... | IBB | WP | HBP | \ |
|-------|----------|--------|-----|-----|-----|-----|-----|-----|-----|--------|-----|-----|-----|-----|-----|
| 3 | altrn101 | 1919 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 20491 | alvaw001 | 1989 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 14000 | apodb101 | 1973 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 19114 | arrof001 | 1986 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 663 | bents101 | 1922 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 40336 | weisz001 | 2018 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 6277 | willa103 | 1946 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 10476 | willt102 | 1962 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 18235 | wortr101 | 1983 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |
| 13621 | younl101 | 1971 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0 | 0 | |

| | BK | BFP | GF | R | SH | SF | GIDP |
|-------|-----|-----|-----|-----|-----|-----|------|
| 3 | 0 | 4.0 | 0 | 4 | 0.0 | 0.0 | 0.0 |
| 20491 | 0 | 5.0 | 0 | 3 | 0.0 | 0.0 | 0.0 |
| 14000 | 0 | 2.0 | 0 | 1 | 0.0 | 0.0 | 0.0 |
| 19114 | 0 | 3.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 663 | 0 | 2.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 40336 | 0 | 4.0 | 0 | 4 | 0.0 | 0.0 | 0.0 |
| 6277 | 0 | 2.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |
| 10476 | 0 | 3.0 | 0 | 1 | 0.0 | 0.0 | 0.0 |
| 18235 | 0 | 4.0 | 0 | 1 | 0.0 | 0.0 | 0.0 |
| 13621 | 0 | 0.0 | 0 | 0 | 0.0 | 0.0 | 0.0 |

[74 rows x 27 columns]

```
[258]: df_era_missing['retroID'].nunique()
```

```
[258]: 74
```

74 rows and 74 unique IDs means that each of these players is only missing the ERA stat for one year. We'll first see, like with BAOpp, if they played other years.

```
[259]: era_checks = df[df['retroID'].isin(df_era_missing['retroID'])].  
      ↪sort_values('retroID')
```

```
[260]: val_counts = era_checks['retroID'].value_counts()
```

```
[261]: one_time_players = list(compress(val_counts.index, val_counts.eq(1)))
```

```
[262]: one_time_players
```

```
[262]: ['russr102',  
      'moorb104',  
      'palam101',  
      'musis101',  
      'garda103',  
      'weis001',  
      'koenw101',  
      'bents101',  
      'fordw103',  
      'schej101',  
      'brucf101',  
      'davav101',  
      'hamad101',  
      'walkm101',  
      'sundg101',  
      'younl101',  
      'browj102']
```

```
[263]: df['ERA'].mean()
```

```
[263]: 5.165393567919002
```

```
[264]: df['ERA'].std()
```

```
[264]: 5.2791159271962815
```

Intuitively, 5.17 is a bit of a high ERA. Though the stat can grow infinitely in theory and low numbers are very difficult, I don't want to assign 10 to the missing values. It's just too much. I'll just do mean + std/2.

```
[265]: filled_era = df['ERA'].mean() + (df['ERA'].std())/2
```

```
[266]: df.loc[df['retroID'].isin(one_time_players), ['ERA']] = filled_era  
df[df['retroID'].isin(one_time_players)]['ERA']
```

```
[266]: 101      7.804952  
      173      7.804952  
      663      7.804952
```

```

722      7.804952
931      7.804952
1447     7.804952
1755     7.804952
2154     7.804952
3709     7.804952
4485     7.804952
4513     7.804952
7678     7.804952
8773     7.804952
9903     7.804952
12532    7.804952
13621    7.804952
40336    7.804952
Name: ERA, dtype: float64

```

```
[267]: df_era_missing = df[df['ERA'].isnull()].sort_values('retroID')
```

We'll continue to follow the same method as for BAOpp with the rest of the missing values.

```
[268]: df_era_missing.shape
```

```
[268]: (57, 27)
```

```
[269]: era_checks = df[df['retroID'].isin(df_era_missing['retroID'])].
        ↪sort_values('retroID')
era_checks['retroID'].value_counts()
```

```
[269]: hillr001    16
choar001    16
mclic101    15
alvaw001    15
farme101    14
medid101    14
kosld101    13
coopm101    13
radis001    11
burkb102    10
owchb001    10
milna101    10
arrof001     9
chent101     9
harvb001     9
perim001     9
deanp101     9
ray-j101     9
pennb001     7

```

| | |
|----------|---|
| navaj101 | 7 |
| brotr001 | 7 |
| jonen001 | 7 |
| painp101 | 6 |
| willt102 | 6 |
| moorc101 | 6 |
| tolls002 | 5 |
| luebs101 | 5 |
| reina102 | 5 |
| scarm101 | 5 |
| mccud001 | 5 |
| apodb101 | 5 |
| blake101 | 4 |
| villb002 | 4 |
| kreur101 | 4 |
| wortr101 | 4 |
| tankd001 | 3 |
| pitls101 | 3 |
| stufp101 | 3 |
| sabee001 | 3 |
| geard101 | 3 |
| kochm001 | 3 |
| vaugp101 | 3 |
| roeto101 | 2 |
| urdal001 | 2 |
| willa103 | 2 |
| green002 | 2 |
| dibup101 | 2 |
| uhl-b101 | 2 |
| wardd101 | 2 |
| kella101 | 2 |
| eschj001 | 2 |
| kammb101 | 2 |
| jeant101 | 2 |
| engej101 | 2 |
| smitd105 | 2 |
| halts001 | 2 |
| altrn101 | 2 |

Name: retroID, dtype: int64

```
[270]: df['ERA'] = df.groupby("retroID")['ERA'].transform(lambda era: era.fillna(era.
    ↳mean()))
```

```
[271]: 100 * df.isnull().sum() / len(df)
```

```
[271]: retroID    0.000000
      yearID     0.000000
```

```

W          0.000000
L          0.000000
G          0.000000
GS         0.000000
CG         0.000000
SHO        0.000000
SV         0.000000
IPouts     0.000000
H          0.000000
ER         0.000000
HR         0.000000
BB         0.000000
SO         0.000000
BAOpp      0.000000
ERA        0.000000
IBB        0.000000
WP         0.000000
HBP        0.000000
BK         0.000000
BFP        0.007431
GF         0.000000
R          0.000000
SH         0.000000
SF         0.000000
GIDP       0.000000
dtype: float64

```

Missing Values: BFP

```
[272]: df_bfp_missing = df[df['BFP'].isnull()].sort_values('retroID')
```

```
[273]: df_bfp_missing
```

```
[273]:
```

| | retroID | yearID | W | L | G | GS | CG | SHO | SV | IPouts | ... | IBB | WP | HBP | \ |
|------|----------|--------|---|---|----|----|----|-----|----|--------|-----|-----|----|-----|---|
| 709 | fourj101 | 1922 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 3 | ... | 0.0 | 0 | 0 | |
| 1171 | jamel101 | 1924 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 3 | ... | 0.0 | 0 | 0 | |
| 802 | pierb103 | 1922 | 3 | 9 | 29 | 12 | 7 | 1 | 0 | 364 | ... | 0.0 | 4 | 6 | |

| | BK | BFP | GF | R | SH | SF | GIDP |
|------|----|-----|----|----|-----|-----|------|
| 709 | 0 | NaN | 1 | 0 | 0.0 | 0.0 | 0.0 |
| 1171 | 0 | NaN | 1 | 2 | 0.0 | 0.0 | 0.0 |
| 802 | 0 | NaN | 10 | 77 | 0.0 | 0.0 | 0.0 |

```
[3 rows x 27 columns]
```

These amounts are negligible. Rather than take averages or set to 0, I'm going to get a little clever. A pitcher intuitively faces hitters until he gets an out, and comes out of the game if he can't get one. There's a lot of work I could do to get a good approximation, but since I'm only filling 3 rows

out of over 40,000 I'm just going to set the missing values to (IPouts - G). This gives us the number of outs a pitcher earned minus 1 for each game he appeared in (presumably this 1 represents the final batter, whom the pitcher did not get out).

```
[274]: df['BFP'].fillna(df['IPouts'] - df['G'], inplace=True)
```

```
[275]: 100 * df.isnull().sum() / len(df)
```

```
[275]: retroID      0.0
      yearID      0.0
      W          0.0
      L          0.0
      G          0.0
      GS         0.0
      CG         0.0
      SHO        0.0
      SV         0.0
      IPouts     0.0
      H          0.0
      ER         0.0
      HR         0.0
      BB         0.0
      SO         0.0
      BAOpp      0.0
      ERA        0.0
      IBB        0.0
      WP         0.0
      HBP        0.0
      BK         0.0
      BFP        0.0
      GF         0.0
      R          0.0
      SH         0.0
      SF         0.0
      GIDP       0.0
      dtype: float64
```

Data Aggregation

Now we can group by retroID, but we need to be more careful than we were with fielding, catching and batting. Some of these stats are averages and some are sum totals, so when we group by we need to handle them differently. We'll split them into two dataframes and do a join. The splitting step will require some intuitive knowledge about baseball statistics. But before we do all of this, we can now get rid of the yearID column.

```
[276]: df.drop(columns=['yearID'], inplace=True)
```

```
[277]: df.columns
```



```
[277]: Index(['retroID', 'W', 'L', 'G', 'GS', 'CG', 'SHO', 'SV', 'IPouts', 'H', 'ER',
          'HR', 'BB', 'SO', 'BAOpp', 'ERA', 'IBB', 'WP', 'HBP', 'BK', 'BFP', 'GF',
          'R', 'SH', 'SF', 'GIDP'],
          dtype='object')
```

```
[278]: average_stats = ['BAOpp', 'ERA']
```

It's only two columns that are averages, and we could probably do without ERA since it's a function of batters faced and runs allowed. We'll keep it since it's such a fundamental statistic in the sport and we have to split anyway for BAOpp, which is a very important one to keep track of.

```
[279]: df_avgs = df[['retroID', 'BAOpp', 'ERA']]
```

```
[280]: df_avgs.head()
```

```
[280]:      retroID  BAOpp  ERA
0  adamb104    0.22  1.98
1  adamw101    0.38  3.86
2  alexg102    0.21  1.72
3  altrn101    1.00  0.00
4  amesr101    0.31  4.89
```

```
[281]: df_sums = df.drop(columns=average_stats)
```

```
[282]: df_sums.head()
```

```
[282]:      retroID   W   L   G  GS  CG  SHO  SV  IPouts   H   ...  IBB  WP  HBP  BK  \
0  adamb104   17   10   34  29  23    6    1    790  213  ...   0.0   2    3    0
1  adamw101    0    0    1   0   0    0    0     14    7  ...   0.0   0    1    0
2  alexg102   16   11   30  27  20    9    1   705  180  ...   0.0   1    0    0
3  altrn101    0    0    1   0   0    0    0     0    4  ...   0.0   0    0    0
4  amesr101    3    5   23   7   1    0    1   210   88  ...   0.0   3    1    0

      BFP  GF   R   SH   SF  GIDP
0  1017.0   5  66  0.0  0.0   0.0
1   21.0   1   2  0.0  0.0   0.0
2   906.0   3  51  0.0  0.0   0.0
3    4.0   0   4  0.0  0.0   0.0
4   314.0  10  44  0.0  0.0   0.0
```

[5 rows x 24 columns]

```
[283]: df_avgs.shape
```

```
[283]: (40372, 3)
```

```
[284]: df_sums.shape
```

```
[284]: (40372, 24)
```

```
[289]: df_avgs = df_avgs.groupby('retroID').mean().round(4).reset_index()
```

```
[292]: df_avgs.shape
```

```
[292]: (7835, 3)
```

```
[288]: df_sums = df_sums.groupby('retroID').sum().reset_index()
```

```
[293]: df_sums.shape
```

```
[293]: (7835, 24)
```

```
[291]: pd.merge(df_avgs, df_sums, on='retroID')
```

```
[291]:
```

| | retroID | BAOpp | ERA | W | L | G | GS | CG | SHO | SV | ... | IBB | WP | \ |
|------|----------|--------|--------|-----|-----|-----|-----|-----|-----|-----|-----|------|-----|---|
| 0 | aardd001 | 0.2574 | 5.1944 | 16 | 18 | 331 | 0 | 0 | 0 | 69 | ... | 22.0 | 12 | |
| 1 | aased001 | 0.2508 | 3.4931 | 66 | 60 | 448 | 91 | 22 | 5 | 82 | ... | 45.0 | 22 | |
| 2 | abadf001 | 0.2501 | 4.0733 | 8 | 27 | 363 | 6 | 0 | 0 | 2 | ... | 10.0 | 9 | |
| 3 | abbog001 | 0.2786 | 4.3317 | 62 | 83 | 248 | 206 | 37 | 5 | 0 | ... | 28.0 | 18 | |
| 4 | abboj001 | 0.2804 | 4.4964 | 87 | 108 | 263 | 254 | 31 | 6 | 0 | ... | 30.0 | 53 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 7830 | zolds101 | 0.2700 | 3.6890 | 43 | 53 | 250 | 93 | 30 | 5 | 8 | ... | 0.0 | 8 | |
| 7831 | zubeb101 | 0.2717 | 5.3617 | 43 | 42 | 224 | 65 | 23 | 3 | 6 | ... | 0.0 | 28 | |
| 7832 | zumaj001 | 0.2286 | 3.4420 | 13 | 12 | 171 | 0 | 0 | 0 | 5 | ... | 11.0 | 16 | |
| 7833 | zuveg101 | 0.2760 | 4.1280 | 32 | 36 | 265 | 31 | 9 | 2 | 40 | ... | 29.0 | 10 | |
| 7834 | zycht001 | 0.2183 | 2.8000 | 7 | 3 | 70 | 1 | 0 | 0 | 1 | ... | 5.0 | 2 | |

| | HBP | BK | BFP | GF | R | SH | SF | GIDP |
|------|-----|-----|--------|-----|-----|------|------|-------|
| 0 | 16 | 1 | 1475.0 | 141 | 169 | 17.0 | 11.0 | 21.0 |
| 1 | 7 | 3 | 4730.0 | 235 | 503 | 50.0 | 34.0 | 106.0 |
| 2 | 12 | 2 | 1350.0 | 96 | 137 | 7.0 | 12.0 | 22.0 |
| 3 | 32 | 5 | 5508.0 | 13 | 707 | 60.0 | 39.0 | 111.0 |
| 4 | 32 | 11 | 7211.0 | 5 | 880 | 70.0 | 47.0 | 200.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7830 | 3 | 4 | 3946.0 | 78 | 423 | 0.0 | 0.0 | 0.0 |
| 7831 | 4 | 1 | 3476.0 | 90 | 418 | 0.0 | 0.0 | 0.0 |
| 7832 | 4 | 0 | 911.0 | 35 | 80 | 6.0 | 10.0 | 10.0 |
| 7833 | 27 | 1 | 2746.0 | 139 | 296 | 0.0 | 0.0 | 0.0 |
| 7834 | 8 | 1 | 309.0 | 14 | 24 | 1.0 | 3.0 | 6.0 |

```
[7835 rows x 26 columns]
```

This gives us the appropriate amount of rows and columns, so the merge worked. We'll send this as our final output.

```
[294]: df = pd.merge(df_avgs, df_sums, on='retroID')
```

```
[296]: df.shape
```

```
[296]: (7835, 26)
```

We're ready to export the resulting table by saving to a csv.

add_advanced_pitching_stats

April 29, 2020

```
[53]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[54]: df = pd.read_csv('../core/output/pitchers.csv')
df_adv = pd.read_csv('../core/output/advanced_pitching.csv')
```

Adding Advanced Stats

```
[55]: df_adv.sort_values('retroID')
```

```
[55]:
```

| | retroID | IP | K/9 | BB/9 | HR/9 | BABIP | LOB% | ERA | FIP | WAR |
|------|----------|----------|------|------|------|-------|------|------|------|------|
| 2866 | aardd001 | 0.062360 | 9.08 | 4.89 | 1.09 | 0.285 | 74.5 | 4.27 | 4.45 | 1.1 |
| 841 | aased001 | 0.205233 | 5.20 | 3.71 | 0.72 | 0.282 | 73.4 | 3.80 | 3.85 | 11.7 |
| 3237 | abadf001 | 0.061102 | 7.62 | 3.16 | 1.14 | 0.281 | 77.7 | 3.67 | 4.24 | 0.6 |
| 949 | abbog001 | 0.237967 | 3.39 | 2.46 | 1.13 | 0.278 | 69.3 | 4.39 | 4.46 | 10.2 |
| 394 | abboj001 | 0.309765 | 4.77 | 3.33 | 0.83 | 0.295 | 70.0 | 4.25 | 4.25 | 22.7 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1030 | zolds101 | 0.171925 | 2.00 | 2.91 | 0.52 | 0.267 | 70.7 | 3.54 | 3.80 | 9.3 |
| 1934 | zubeb101 | 0.145445 | 4.39 | 5.36 | 0.40 | 0.283 | 69.0 | 4.28 | 3.96 | 3.3 |
| 2098 | zumaj001 | 0.038711 | 9.01 | 4.89 | 0.77 | 0.267 | 78.7 | 3.00 | 3.94 | 2.7 |
| 2399 | zuveg101 | 0.118817 | 3.12 | 2.84 | 0.78 | 0.270 | 73.2 | 3.54 | 3.93 | 1.9 |
| 2808 | zycht001 | 0.013360 | 9.91 | 4.21 | 0.37 | 0.293 | 79.1 | 2.72 | 3.22 | 1.1 |

[8025 rows x 10 columns]

```
[56]: df
```

```
[56]:
```

| | retroID | BAOpp | ERA | CG | SHO | IPouts | H | ER | HR | BB | ... | WP | \ |
|------|----------|--------|--------|-----|-----|--------|------|-----|-----|-----|-----|-----|-----|
| 0 | aardd001 | 0.2574 | 5.1944 | 0 | 0 | 1011 | 296 | 160 | 41 | 183 | ... | 12 | |
| 1 | aased001 | 0.2508 | 3.4931 | 22 | 5 | 3328 | 1085 | 468 | 89 | 457 | ... | 22 | |
| 2 | abadf001 | 0.2447 | 4.0810 | 0 | 0 | 992 | 309 | 135 | 42 | 116 | ... | 10 | |
| 3 | abbog001 | 0.2786 | 4.3317 | 37 | 5 | 3858 | 1405 | 627 | 162 | 352 | ... | 18 | |
| 4 | abboj001 | 0.2804 | 4.4964 | 31 | 6 | 5022 | 1779 | 791 | 154 | 620 | ... | 53 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8020 | zolds101 | 0.2700 | 3.6890 | 30 | 5 | 2788 | 956 | 366 | 54 | 301 | ... | 8 | |
| 8021 | zubeb101 | 0.2717 | 5.3617 | 23 | 3 | 2358 | 767 | 374 | 35 | 468 | ... | 28 | |
| 8022 | zumaj001 | 0.2286 | 3.4420 | 0 | 0 | 629 | 169 | 71 | 18 | 114 | ... | 16 | |

| | | | | | | | | | | | | |
|------|----------|--------|--------|---|---|------|-----|-----|----|-----|-----|----|
| 8023 | zuveg101 | 0.2760 | 4.1280 | 9 | 2 | 1927 | 660 | 253 | 56 | 203 | ... | 10 |
| 8024 | zycht001 | 0.2183 | 2.8000 | 0 | 0 | 218 | 57 | 22 | 3 | 34 | ... | 2 |

| | HBP | BK | BFP | GF | R | SH | SF | GIDP | K% |
|------|-----|----|------|-----|-----|----|----|------|----------|
| 0 | 16 | 1 | 1475 | 141 | 169 | 17 | 11 | 21 | 0.230508 |
| 1 | 7 | 3 | 4730 | 235 | 503 | 50 | 34 | 106 | 0.135518 |
| 2 | 12 | 2 | 1399 | 97 | 143 | 7 | 12 | 25 | 0.200143 |
| 3 | 32 | 5 | 5508 | 13 | 707 | 60 | 39 | 111 | 0.087872 |
| 4 | 32 | 11 | 7211 | 5 | 880 | 70 | 47 | 200 | 0.123145 |
| ... | ... | .. | ... | ... | ... | .. | .. | ... | ... |
| 8020 | 3 | 4 | 3946 | 78 | 423 | 0 | 0 | 0 | 0.052458 |
| 8021 | 4 | 1 | 3476 | 90 | 418 | 0 | 0 | 0 | 0.110184 |
| 8022 | 4 | 0 | 911 | 35 | 80 | 6 | 10 | 10 | 0.230516 |
| 8023 | 27 | 1 | 2746 | 139 | 296 | 0 | 0 | 0 | 0.081209 |
| 8024 | 8 | 1 | 309 | 14 | 24 | 1 | 3 | 6 | 0.258900 |

[8025 rows x 22 columns]

```
[57]: df = df.drop(columns=['ERA'])
```

```
[58]: df = df.merge(df_adv, on='retroID', how='left')
```

```
[59]: 100 * df.isnull().sum() / len(df)
```

```
[59]: retroID      0.0
      BAOpp       0.0
      CG         0.0
      SHO        0.0
      IPouts     0.0
      H          0.0
      ER         0.0
      HR         0.0
      BB         0.0
      SO         0.0
      IBB        0.0
      WP         0.0
      HBP        0.0
      BK         0.0
      BFP        0.0
      GF         0.0
      R          0.0
      SH         0.0
      SF         0.0
      GIDP       0.0
      K%         0.0
      IP         0.0
      K/9        0.0
```

```

BB/9      0.0
HR/9      0.0
BABIP     0.0
LOB%      0.0
ERA        0.0
FIP        0.0
WAR        0.0
dtype: float64

```

```
[60]: df['Pitching'] = df[['K%', 'ERA', 'FIP', 'WAR']].mean(axis=1).round(3)
```

```
[61]: df['Pitching']
```

```

[61]: 0      2.513
      1      4.871
      2      2.178
      3      4.784
      4      7.831
      ...
      8020    4.173
      8021    2.913
      8022    2.468
      8023    2.363
      8024    1.825
      Name: Pitching, Length: 8025, dtype: float64

```

Finalizing the new Pitching statistic

```
[64]: df['Pitching'].mean()
```

```
[64]: 3.627380436137072
```

```
[65]: df['Pitching'].min()
```


```
[65]: -0.654
```

```
[66]: df['Pitching'].max()
```

```
[66]: 53.138
```

```
[63]: plt.plot(df['Pitching'])
```

```
[63]: [<matplotlib.lines.Line2D at 0x127615250>]
```



```
[67]: df[df['Pitching'] == df['Pitching'].max()]
```

```
[67]:      retroID  BAOpp  CG  SHO  IPouts  H  ER  HR  BB  SO  ...      IP  \
1333  cleaj101   0.83   0    0         1  5   7   0   3   1  ...  0.000019

      K/9  BB/9  HR/9  BABIP  LOB%    ERA    FIP  WAR  Pitching
1333  27.0  81.0   0.0    1.0  12.5  189.0  23.54 -0.1    53.138
```


```
[1 rows x 31 columns]
```

This immediately demonstrates an issue with our Pitching stat - a player with very few appearances, but who did well in those appearances, will be skewed too high. We need to also take innings pitched into account.

```
[91]: df['Pitching'] = df[['K%', 'ERA', 'FIP', 'WAR']].mean(axis=1).round(3) * df['IP']
```

```
[92]: plt.plot(df['Pitching'])
```

```
[92]: [<matplotlib.lines.Line2D at 0x12793b850>]
```



```
[93]: df[df['Pitching'] == df['Pitching'].max()]
```

```
[93]:      retroID  BAOpp  CG  SHO  IPouts  H  ER  HR  BB  SO  ...  \
1341  clemr001  0.2308  118  46  14750  4185  1707  363  1580  4672  ...

      IP  K/9  BB/9  HR/9  BABIP  LOB%  ERA  FIP  WAR  Pitching
1341  0.909717  8.55  2.89  0.66  0.284  74.6  3.12  3.09  133.7  31.871935
```

```
[1 rows x 31 columns]
```

This looks like it worked, but let's explore deeper.

```
[94]: df['Pitching'].mean()
```

```
[94]: 0.49352977097414313
```

```
[95]: df['Pitching'].min()
```

```
[95]: -5.0887375e-05
```

```
[96]: df['Pitching'].max()
```

```
[96]: 31.871935094999998
```

```
[97]: df.sort_values('Pitching').tail(10)
```



```
[97]:      retroID  BAOpp  CG  SHO  IPouts      H    ER   HR   BB    SO  ...  \
5219  niekp001  0.2570  245   45   16213  5044  2012  482  1809  3342  ...
6529  seavt001  0.2285  231   61   14348  3971  1521  380  1390  3640  ...
3588  johnr005  0.2252  100   37   12406  3346  1513  411  1497  4875  ...
7049  suttd001  0.2376  178   58   15847  4692  1914  472  1343  3574  ...
1121  carls001  0.2546  254   55   15652  4672  1864  414  1833  4136  ...
645   blylb001  0.2487  242   60   14910  4632  1830  430  1322  3701  ...
5607  perrg101  0.2540  303   53   16051  4938  1846  399  1379  3534  ...
6311  ryann001  0.2088  222   61   16158  3923  1911  321  2795  5714  ...
4347  maddg002  0.2553  109   35   15025  4726  1756  353   999  3371  ...
1341  clemr001  0.2308  118   46   14750  4185  1707  363  1580  4672  ...
```

```
      IP    K/9  BB/9  HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching
5219  1.000000  5.57  3.01  0.80  0.270  73.6  3.35  3.62  78.5  21.404000
6529  0.884921  6.85  2.62  0.72  0.259  76.7  2.86  3.04  92.7  21.854894
3588  0.765178 10.61  3.26  0.89  0.291  74.7  3.29  3.19 110.4  22.412829
7049  0.977406  6.09  2.29  0.80  0.261  73.6  3.26  3.24  85.9  22.618152
1121  0.965397  7.13  3.16  0.71  0.279  74.1  3.22  3.15  96.9  24.969993
645   0.919672  6.70  2.39  0.78  0.282  74.1  3.31  3.19 103.3  25.286382
5607  0.990008  5.94  2.32  0.67  0.275  73.3  3.11  3.06 100.5  26.441134
6311  0.996651  9.55  4.67  0.54  0.265  73.1  3.19  2.97 107.2  28.307878
4347  0.926722  6.06  1.80  0.63  0.281  72.3  3.16  3.26 116.7  28.562499
1341  0.909717  8.55  2.89  0.66  0.284  74.6  3.12  3.09 133.7  31.871935
```

[10 rows x 31 columns]

```
[98]: df[df['retroID'] == 'kersc001']
```

```
[98]:      retroID  BAOpp  CG  SHO  IPouts      H    ER   HR   BB    SO  ...  \
3761  kersc001  0.2105  25   15   6824  1715  617  173  577  2464  ...

      IP    K/9  BB/9  HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching
3761  0.420829  9.75  2.28  0.68  0.27  79.4  2.44  2.74  64.5  7.359878
```

[1 rows x 31 columns]

This looks good, but we intuitively see a problem with the Pitching stat. ERA and FIP are part of the average, but a low ERA/FIP is better than a high one. We need to subtract them rather than add. With this change, I'm going to see how the stat looks without taking IP into account.

```
[132]: df['-ERA'] = 0 - df['ERA']
df['-FIP'] = 0 - df['FIP']
```

```
[133]: df
```

```
[133]:      retroID  BAOpp  CG  SHO  IPouts      H    ER   HR   BB    SO  ...  BB/9  \
0         aardd001  0.2574   0    0   1011   296  160   41  183  340  ...  4.89
```

| | | | | | | | | | | | | |
|------|----------|--------|-----|-----|------|------|-----|-----|-----|-----|-----|------|
| 1 | aased001 | 0.2508 | 22 | 5 | 3328 | 1085 | 468 | 89 | 457 | 641 | ... | 3.71 |
| 2 | abadf001 | 0.2447 | 0 | 0 | 992 | 309 | 135 | 42 | 116 | 280 | ... | 3.16 |
| 3 | abbog001 | 0.2786 | 37 | 5 | 3858 | 1405 | 627 | 162 | 352 | 484 | ... | 2.46 |
| 4 | abboj001 | 0.2804 | 31 | 6 | 5022 | 1779 | 791 | 154 | 620 | 888 | ... | 3.33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8020 | zolds101 | 0.2700 | 30 | 5 | 2788 | 956 | 366 | 54 | 301 | 207 | ... | 2.91 |
| 8021 | zubeb101 | 0.2717 | 23 | 3 | 2358 | 767 | 374 | 35 | 468 | 383 | ... | 5.36 |
| 8022 | zumaj001 | 0.2286 | 0 | 0 | 629 | 169 | 71 | 18 | 114 | 210 | ... | 4.89 |
| 8023 | zuveg101 | 0.2760 | 9 | 2 | 1927 | 660 | 253 | 56 | 203 | 223 | ... | 2.84 |
| 8024 | zycht001 | 0.2183 | 0 | 0 | 218 | 57 | 22 | 3 | 34 | 80 | ... | 4.21 |

| | HR/9 | BABIP | LOB% | ERA | FIP | WAR | Pitching | -ERA | -FIP |
|------|------|-------|------|------|------|------|----------|-------|-------|
| 0 | 1.09 | 0.285 | 74.5 | 4.27 | 4.45 | 1.1 | -1.847 | -4.27 | -4.45 |
| 1 | 0.72 | 0.282 | 73.4 | 3.80 | 3.85 | 11.7 | 1.046 | -3.80 | -3.85 |
| 2 | 1.14 | 0.281 | 77.7 | 3.67 | 4.24 | 0.6 | -1.777 | -3.67 | -4.24 |
| 3 | 1.13 | 0.278 | 69.3 | 4.39 | 4.46 | 10.2 | 0.359 | -4.39 | -4.46 |
| 4 | 0.83 | 0.295 | 70.0 | 4.25 | 4.25 | 22.7 | 3.581 | -4.25 | -4.25 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8020 | 0.52 | 0.267 | 70.7 | 3.54 | 3.80 | 9.3 | 0.503 | -3.54 | -3.80 |
| 8021 | 0.40 | 0.283 | 69.0 | 4.28 | 3.96 | 3.3 | -1.207 | -4.28 | -3.96 |
| 8022 | 0.77 | 0.267 | 78.7 | 3.00 | 3.94 | 2.7 | -1.002 | -3.00 | -3.94 |
| 8023 | 0.78 | 0.270 | 73.2 | 3.54 | 3.93 | 1.9 | -1.372 | -3.54 | -3.93 |
| 8024 | 0.37 | 0.293 | 79.1 | 2.72 | 3.22 | 1.1 | -1.145 | -2.72 | -3.22 |

[8025 rows x 33 columns]

```
[134]: df['Pitching'] = df[['K%', '-ERA', '-FIP', 'WAR']].mean(axis=1).round(3)
```

```
[135]: df.sort_values('Pitching').tail(10)
```

```
[135]:
```

| | retroID | BAOpp | CG | SHO | IPouts | H | ER | HR | BB | SO | ... | \ |
|------|----------|--------|-----|-----|--------|------|------|-----|------|------|-----|---|
| 7049 | suttd001 | 0.2376 | 178 | 58 | 15847 | 4692 | 1914 | 472 | 1343 | 3574 | ... | |
| 2824 | grovl101 | 0.2535 | 298 | 35 | 11822 | 3849 | 1339 | 162 | 1187 | 2266 | ... | |
| 6529 | seavt001 | 0.2285 | 231 | 61 | 14348 | 3971 | 1521 | 380 | 1390 | 3640 | ... | |
| 1121 | carls001 | 0.2546 | 254 | 55 | 15652 | 4672 | 1864 | 414 | 1833 | 4136 | ... | |
| 5607 | perrg101 | 0.2540 | 303 | 53 | 16051 | 4938 | 1846 | 399 | 1379 | 3534 | ... | |
| 645 | blylb001 | 0.2487 | 242 | 60 | 14910 | 4632 | 1830 | 430 | 1322 | 3701 | ... | |
| 6311 | ryann001 | 0.2088 | 222 | 61 | 16158 | 3923 | 1911 | 321 | 2795 | 5714 | ... | |
| 3588 | johnr005 | 0.2252 | 100 | 37 | 12406 | 3346 | 1513 | 411 | 1497 | 4875 | ... | |
| 4347 | maddg002 | 0.2553 | 109 | 35 | 15025 | 4726 | 1756 | 353 | 999 | 3371 | ... | |
| 1341 | clemr001 | 0.2308 | 118 | 46 | 14750 | 4185 | 1707 | 363 | 1580 | 4672 | ... | |

| | BB/9 | HR/9 | BABIP | LOB% | ERA | FIP | WAR | Pitching | -ERA | -FIP |
|------|------|------|-------|------|------|------|------|----------|-------|-------|
| 7049 | 2.29 | 0.80 | 0.261 | 73.6 | 3.26 | 3.24 | 85.9 | 19.891 | -3.26 | -3.24 |
| 2824 | 2.71 | 0.37 | 0.284 | 71.8 | 3.06 | 3.36 | 88.8 | 20.629 | -3.06 | -3.36 |
| 6529 | 2.62 | 0.72 | 0.259 | 76.7 | 2.86 | 3.04 | 92.7 | 21.747 | -2.86 | -3.04 |
| 1121 | 3.16 | 0.71 | 0.279 | 74.1 | 3.22 | 3.15 | 96.9 | 22.680 | -3.22 | -3.15 |

| | | | | | | | | | | |
|------|------|------|-------|------|------|------|-------|--------|-------|-------|
| 5607 | 2.32 | 0.67 | 0.275 | 73.3 | 3.11 | 3.06 | 100.5 | 23.623 | -3.11 | -3.06 |
| 645 | 2.39 | 0.78 | 0.282 | 74.1 | 3.31 | 3.19 | 103.3 | 24.245 | -3.31 | -3.19 |
| 6311 | 4.67 | 0.54 | 0.265 | 73.1 | 3.19 | 2.97 | 107.2 | 25.323 | -3.19 | -2.97 |
| 3588 | 3.26 | 0.89 | 0.291 | 74.7 | 3.29 | 3.19 | 110.4 | 26.051 | -3.29 | -3.19 |
| 4347 | 1.80 | 0.63 | 0.281 | 72.3 | 3.16 | 3.26 | 116.7 | 27.611 | -3.16 | -3.26 |
| 1341 | 2.89 | 0.66 | 0.284 | 74.6 | 3.12 | 3.09 | 133.7 | 31.930 | -3.12 | -3.09 |

[10 rows x 33 columns]

```
[136]: df[df['retroID'] == 'kersc001']
```

```
[136]:      retroID  BAOpp  CG  SHO  IPouts    H  ER  HR  BB  SO  ...  BB/9  \
3761  kersc001  0.2105  25   15    6824 1715  617  173  577 2464  ...  2.28

      HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching  -ERA  -FIP
3761  0.68   0.27  79.4   2.44  2.74  64.5    14.899 -2.44 -2.74
```

[1 rows x 33 columns]

```
[137]: df[df['retroID'] == 'johnr005']
```

```
[137]:      retroID  BAOpp  CG  SHO  IPouts    H  ER  HR  BB  SO  ...  \
3588  johnr005  0.2252 100   37   12406 3346 1513  411 1497 4875  ...

      BB/9  HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching  -ERA  -FIP
3588  3.26  0.89  0.291  74.7  3.29  3.19 110.4    26.051 -3.29 -3.19
```

[1 rows x 33 columns]

```
[138]: df[df['retroID'] == 'bumgm001']
```

```
[138]:      retroID  BAOpp  CG  SHO  IPouts    H  ER  HR  BB  SO  ...  BB/9  \
942  bumgm001  0.2358  15    6    5538 1622  642  192  428 1794  ...  2.09

      HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching  -ERA  -FIP
942  0.94  0.284  76.4  3.13  3.32  31.3    6.272 -3.13 -3.32
```

[1 rows x 33 columns]

```
[139]: df[df['retroID'] == 'mahop001']
```

```
[139]:      retroID  BAOpp  CG  SHO  IPouts    H  ER  HR  BB  SO  ...  BB/9  \
4377  mahop001  0.2751   0    0    2127  738  431  116  392  452  ...  4.98

      HR/9  BABIP  LOB%   ERA   FIP  WAR  Pitching  -ERA  -FIP
4377  1.47  0.284  69.5  5.47  5.62 -3.0    -3.487 -5.47 -5.62
```

[1 rows x 33 columns]

```
[140]: df[df['retroID'] == 'coleg001']
```


```
[140]:      retroID  BAOpp  CG  SHO  IPouts    H  ER  HR  BB    SO  ...  BB/9  \
1383  coleg001  0.2381  2    1   3585  1034  427  115  315  1336  ...  2.37

      HR/9  BABIP  LOB%   ERA   FIP   WAR  Pitching  -ERA  -FIP
1383  0.87  0.303  75.6  3.22  3.06  28.8    5.699 -3.22 -3.06
```

[1 rows x 33 columns]

```
[141]: plt.plot(df['Pitching'])
```

```
[141]: [<matplotlib.lines.Line2D at 0x12a1537d0>]
```



```
[131]: df.sort_values('Pitching').head(10)
```

```
[131]:      retroID  BAOpp  CG  SHO  IPouts  H  ER  HR  BB  SO  ...  BB/9  HR/9  \
1333  cleaj101   0.83   0    0         1  5   7   0   3   1  ...   81.0   0.0
7932  wurmf101   0.50   0    0         1  1   4   0   5   1  ...  135.0   0.0
3074  heart001   0.75   0    0         1  3   4   0   4   0  ...  108.0   0.0
3159  herne001   0.50   0    0         1  1   3   1   2   0  ...   54.0  27.0
2272  fishf101   0.66   0    0         1  2   4   0   2   1  ...   54.0   0.0
2666  gomec002   0.00   0    0         1  0   3   0   4   0  ...  108.0   0.0
7819  wilst104   0.00   0    0         1  0   3   0   2   0  ...   54.0   0.0
```

| | | | | | | | | | | | | | |
|------|----------|------|---|---|---|---|---|---|---|---|-----|------|------|
| 88 | alexm001 | 0.50 | 0 | 0 | 2 | 1 | 5 | 1 | 4 | 0 | ... | 54.0 | 13.5 |
| 2981 | harll101 | 0.50 | 0 | 0 | 2 | 2 | 5 | 1 | 4 | 1 | ... | 54.0 | 13.5 |
| 5674 | pickr001 | 0.60 | 0 | 0 | 2 | 3 | 6 | 0 | 4 | 2 | ... | 54.0 | 0.0 |

| | BABIP | LOB% | ERA | FIP | WAR | Pitching | -ERA | -FIP |
|------|-------|------|-------|-------|------|----------|--------|--------|
| 1333 | 1.00 | 12.5 | 189.0 | 23.54 | -0.1 | -53.132 | -189.0 | -23.54 |
| 7932 | 1.00 | 33.3 | 108.0 | 41.54 | -0.1 | -37.374 | -108.0 | -41.54 |
| 3074 | 0.75 | 28.6 | 108.0 | 39.21 | -0.1 | -36.828 | -108.0 | -39.21 |
| 3159 | 0.00 | 0.0 | 81.0 | 60.16 | -0.3 | -35.365 | -81.0 | -60.16 |
| 2272 | 1.00 | 0.0 | 108.0 | 14.60 | 0.0 | -30.600 | -108.0 | -14.60 |
| 2666 | 0.00 | 25.0 | 81.0 | 39.16 | -0.1 | -30.065 | -81.0 | -39.16 |
| 7819 | 0.00 | 0.0 | 81.0 | 29.57 | -0.1 | -27.667 | -81.0 | -29.57 |
| 88 | 0.00 | 0.0 | 67.5 | 40.67 | -0.1 | -27.068 | -67.5 | -40.67 |
| 2981 | 0.50 | 21.7 | 67.5 | 37.08 | -0.1 | -26.139 | -67.5 | -37.08 |
| 5674 | 1.00 | 14.3 | 81.0 | 15.14 | 0.0 | -23.979 | -81.0 | -15.14 |

[10 rows x 33 columns]

I'm happy with this. The new Pitching stat somewhat reflects its component parts but doesn't immediately align with WAR. We don't just want to recreate WAR so that's a good thing.


Normalization

```
[142]: from sklearn.preprocessing import MinMaxScaler
```

```
[143]: scaler = MinMaxScaler()
```

```
[144]: plt.plot(scaler.fit_transform(df[['Pitching']]))
```

```
[144]: [<matplotlib.lines.Line2D at 0x1337d3f50>]
```



```
[145]: df['Pitching'] = scaler.fit_transform(df[['Pitching']])
```

```
[147]: df.sort_values('Pitching').head(10)
```

```
[147]:
```

| | retroID | BAOpp | CG | SHO | IPouts | H | ER | HR | BB | SO | ... | BB/9 | HR/9 | \ |
|------|----------|-------|----|-----|--------|---|----|----|----|----|-----|-------|------|---|
| 1333 | cleaj101 | 0.83 | 0 | 0 | 1 | 5 | 7 | 0 | 3 | 1 | ... | 81.0 | 0.0 | |
| 7932 | wurmf101 | 0.50 | 0 | 0 | 1 | 1 | 4 | 0 | 5 | 1 | ... | 135.0 | 0.0 | |
| 3074 | heart001 | 0.75 | 0 | 0 | 1 | 3 | 4 | 0 | 4 | 0 | ... | 108.0 | 0.0 | |
| 3159 | herne001 | 0.50 | 0 | 0 | 1 | 1 | 3 | 1 | 2 | 0 | ... | 54.0 | 27.0 | |
| 2272 | fishf101 | 0.66 | 0 | 0 | 1 | 2 | 4 | 0 | 2 | 1 | ... | 54.0 | 0.0 | |
| 2666 | gomec002 | 0.00 | 0 | 0 | 1 | 0 | 3 | 0 | 4 | 0 | ... | 108.0 | 0.0 | |
| 7819 | wilst104 | 0.00 | 0 | 0 | 1 | 0 | 3 | 0 | 2 | 0 | ... | 54.0 | 0.0 | |
| 88 | alexm001 | 0.50 | 0 | 0 | 2 | 1 | 5 | 1 | 4 | 0 | ... | 54.0 | 13.5 | |
| 2981 | harll101 | 0.50 | 0 | 0 | 2 | 2 | 5 | 1 | 4 | 1 | ... | 54.0 | 13.5 | |
| 5674 | pickr001 | 0.60 | 0 | 0 | 2 | 3 | 6 | 0 | 4 | 2 | ... | 54.0 | 0.0 | |

| | BABIP | LOB% | ERA | FIP | WAR | Pitching | -ERA | -FIP |
|------|-------|------|-------|-------|------|----------|--------|--------|
| 1333 | 1.00 | 12.5 | 189.0 | 23.54 | -0.1 | 0.000000 | -189.0 | -23.54 |
| 7932 | 1.00 | 33.3 | 108.0 | 41.54 | -0.1 | 0.185253 | -108.0 | -41.54 |
| 3074 | 0.75 | 28.6 | 108.0 | 39.21 | -0.1 | 0.191672 | -108.0 | -39.21 |
| 3159 | 0.00 | 0.0 | 81.0 | 60.16 | -0.3 | 0.208871 | -81.0 | -60.16 |
| 2272 | 1.00 | 0.0 | 108.0 | 14.60 | 0.0 | 0.264889 | -108.0 | -14.60 |
| 2666 | 0.00 | 25.0 | 81.0 | 39.16 | -0.1 | 0.271179 | -81.0 | -39.16 |
| 7819 | 0.00 | 0.0 | 81.0 | 29.57 | -0.1 | 0.299370 | -81.0 | -29.57 |
| 88 | 0.00 | 0.0 | 67.5 | 40.67 | -0.1 | 0.306412 | -67.5 | -40.67 |
| 2981 | 0.50 | 21.7 | 67.5 | 37.08 | -0.1 | 0.317333 | -67.5 | -37.08 |

```
5674    1.00  14.3   81.0  15.14  0.0  0.342726  -81.0 -15.14
```

```
[10 rows x 33 columns]
```

```
[148]: df.sort_values('Pitching').tail(10)
```

```
[148]:      retroID  BAOpp  CG  SHO  IPouts    H    ER  HR  BB  SO  ...  \
7049  suttd001  0.2376  178   58  15847  4692  1914  472  1343  3574  ...
2824  grovl101  0.2535  298   35  11822  3849  1339  162  1187  2266  ...
6529  seavt001  0.2285  231   61  14348  3971  1521  380  1390  3640  ...
1121  carls001  0.2546  254   55  15652  4672  1864  414  1833  4136  ...
5607  perrg101  0.2540  303   53  16051  4938  1846  399  1379  3534  ...
645   blylb001  0.2487  242   60  14910  4632  1830  430  1322  3701  ...
6311  ryann001  0.2088  222   61  16158  3923  1911  321  2795  5714  ...
3588  johnr005  0.2252  100   37  12406  3346  1513  411  1497  4875  ...
4347  maddg002  0.2553  109   35  15025  4726  1756  353   999  3371  ...
1341  clemr001  0.2308  118   46  14750  4185  1707  363  1580  4672  ...
```

```
      BB/9  HR/9  BABIP  LOB%  ERA  FIP  WAR  Pitching  -ERA  -FIP
7049  2.29  0.80  0.261  73.6  3.26  3.24  85.9  0.858468 -3.26 -3.24
2824  2.71  0.37  0.284  71.8  3.06  3.36  88.8  0.867144 -3.06 -3.36
6529  2.62  0.72  0.259  76.7  2.86  3.04  92.7  0.880287 -2.86 -3.04
1121  3.16  0.71  0.279  74.1  3.22  3.15  96.9  0.891256 -3.22 -3.15
5607  2.32  0.67  0.275  73.3  3.11  3.06  100.5  0.902342 -3.11 -3.06
645   2.39  0.78  0.282  74.1  3.31  3.19  103.3  0.909654 -3.31 -3.19
6311  4.67  0.54  0.265  73.1  3.19  2.97  107.2  0.922327 -3.19 -2.97
3588  3.26  0.89  0.291  74.7  3.29  3.19  110.4  0.930886 -3.29 -3.19
4347  1.80  0.63  0.281  72.3  3.16  3.26  116.7  0.949225 -3.16 -3.26
1341  2.89  0.66  0.284  74.6  3.12  3.09  133.7  1.000000 -3.12 -3.09
```

```
[10 rows x 33 columns]
```

Finally, we should get rid of the -ERA and -FIP columns.

```
[149]: df = df.drop(columns=['-ERA', '-FIP'])
```

```
[150]: df
```

```
[150]:      retroID  BAOpp  CG  SHO  IPouts    H    ER  HR  BB  SO  ...  \
0      aardd001  0.2574   0   0   1011   296   160   41  183  340  ...
1      aased001  0.2508  22   5   3328  1085   468   89  457  641  ...
2      abadf001  0.2447   0   0    992   309   135   42  116  280  ...
3      abbog001  0.2786  37   5   3858  1405   627  162  352  484  ...
4      abboj001  0.2804  31   6   5022  1779   791  154  620  888  ...
...      ...      ...  ..  ...      ...      ...      ...      ...      ...
8020  zolds101  0.2700  30   5   2788   956   366   54  301  207  ...
8021  zubeb101  0.2717  23   3   2358   767   374   35  468  383  ...
```

| | | | | | | | | | | | |
|------|----------|--------|---|---|------|-----|-----|----|-----|-----|-----|
| 8022 | zumaj001 | 0.2286 | 0 | 0 | 629 | 169 | 71 | 18 | 114 | 210 | ... |
| 8023 | zuveg101 | 0.2760 | 9 | 2 | 1927 | 660 | 253 | 56 | 203 | 223 | ... |
| 8024 | zycht001 | 0.2183 | 0 | 0 | 218 | 57 | 22 | 3 | 34 | 80 | ... |

| | IP | K/9 | BB/9 | HR/9 | BABIP | LOB% | ERA | FIP | WAR | Pitching |
|------|----------|------|------|------|-------|------|------|------|------|----------|
| 0 | 0.062360 | 9.08 | 4.89 | 1.09 | 0.285 | 74.5 | 4.27 | 4.45 | 1.1 | 0.602913 |
| 1 | 0.205233 | 5.20 | 3.71 | 0.72 | 0.282 | 73.4 | 3.80 | 3.85 | 11.7 | 0.636924 |
| 2 | 0.061102 | 7.62 | 3.16 | 1.14 | 0.281 | 77.7 | 3.67 | 4.24 | 0.6 | 0.603736 |
| 3 | 0.237967 | 3.39 | 2.46 | 1.13 | 0.278 | 69.3 | 4.39 | 4.46 | 10.2 | 0.628847 |
| 4 | 0.309765 | 4.77 | 3.33 | 0.83 | 0.295 | 70.0 | 4.25 | 4.25 | 22.7 | 0.666725 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8020 | 0.171925 | 2.00 | 2.91 | 0.52 | 0.267 | 70.7 | 3.54 | 3.80 | 9.3 | 0.630540 |
| 8021 | 0.145445 | 4.39 | 5.36 | 0.40 | 0.283 | 69.0 | 4.28 | 3.96 | 3.3 | 0.610437 |
| 8022 | 0.038711 | 9.01 | 4.89 | 0.77 | 0.267 | 78.7 | 3.00 | 3.94 | 2.7 | 0.612847 |
| 8023 | 0.118817 | 3.12 | 2.84 | 0.78 | 0.270 | 73.2 | 3.54 | 3.93 | 1.9 | 0.608497 |
| 8024 | 0.013360 | 9.91 | 4.21 | 0.37 | 0.293 | 79.1 | 2.72 | 3.22 | 1.1 | 0.611166 |

[8025 rows x 31 columns]

[]:

teams_pre

March 9, 2020

```
[4]: import math
import numpy as np
import pandas as pd

[5]: df = pd.read_csv('../data/lahman/mlb_data/Teams.csv')

[6]: df.columns

[6]: Index(['yearID', 'lgID', 'teamID', 'franchID', 'divID', 'Rank', 'G', 'Ghome',
        'W', 'L', 'DivWin', 'WCWin', 'LgWin', 'WSWin', 'R', 'AB', 'H', '2B',
        '3B', 'HR', 'BB', 'SO', 'SB', 'CS', 'HBP', 'SF', 'RA', 'ER', 'ERA',
        'CG', 'SHO', 'SV', 'IPouts', 'HA', 'HRA', 'BBA', 'SOA', 'E', 'DP', 'FP',
        'name', 'park', 'attendance', 'BPF', 'PPF', 'teamIDBR',
        'teamIDlahman45', 'teamIDretro'],
        dtype='object')

[7]: df.head()
```

| | yearID | lgID | teamID | franchID | divID | Rank | G | Ghome | W | L | ... | DP | \ |
|---|--------|------|--------|----------|-------|------|-----|-------|----|----|-----|-----|---|
| 0 | 1919 | AL | BOS | BOS | NaN | 6 | 138 | 66 | 66 | 71 | ... | 118 | |
| 1 | 1919 | NL | BRO | LAD | NaN | 5 | 141 | 70 | 69 | 71 | ... | 84 | |
| 2 | 1919 | NL | BSN | ATL | NaN | 6 | 140 | 68 | 57 | 82 | ... | 111 | |
| 3 | 1919 | AL | CHA | CHW | NaN | 1 | 140 | 70 | 88 | 52 | ... | 116 | |
| 4 | 1919 | NL | CHN | CHC | NaN | 3 | 140 | 71 | 75 | 65 | ... | 87 | |

| | FP | name | park | attendance | BPF | PPF | teamIDBR | \ |
|---|-------|-------------------|---------------|------------|-----|-----|----------|---|
| 0 | 0.975 | Boston Red Sox | Fenway Park I | 417291 | 94 | 94 | BOS | |
| 1 | 0.963 | Brooklyn Robins | Ebbets Field | 360721 | 103 | 103 | BRO | |
| 2 | 0.966 | Boston Braves | Braves Field | 167401 | 95 | 98 | BSN | |
| 3 | 0.969 | Chicago White Sox | Comiskey Park | 627186 | 100 | 99 | CHW | |
| 4 | 0.969 | Chicago Cubs | Wrigley Field | 424430 | 100 | 99 | CHC | |

| | teamIDlahman45 | teamIDretro |
|---|----------------|-------------|
| 0 | BOS | BOS |
| 1 | BRO | BRO |
| 2 | BSN | BSN |
| 3 | CHA | CHA |

[5 rows x 48 columns]

```
[8]: df = df.drop(columns=['teamIDlahman45', 'teamIDBR'])
```

The first step is to ensure we're only using one ID per team. It would be best to just use Retrosheet's values, so our first step is to see where teamID differs from teamIDretro. Once we come up with a way to fix these differences, we'll want to write it as a script that we can use elsewhere - for example, in the batting table where we're using the regular teamID values.

```
[9]: df[(df['teamID'] != df['teamIDretro'])][['yearID', 'teamID', 'teamIDretro', 'name']]
```

```
[9]:
```

| | yearID | teamID | teamIDretro | name |
|------|--------|--------|-------------|-------------------|
| 551 | 1953 | ML1 | MLN | Milwaukee Braves |
| 568 | 1954 | ML1 | MLN | Milwaukee Braves |
| 585 | 1955 | ML1 | MLN | Milwaukee Braves |
| 601 | 1956 | ML1 | MLN | Milwaukee Braves |
| 617 | 1957 | ML1 | MLN | Milwaukee Braves |
| 633 | 1958 | ML1 | MLN | Milwaukee Braves |
| 649 | 1959 | ML1 | MLN | Milwaukee Braves |
| 665 | 1960 | ML1 | MLN | Milwaukee Braves |
| 683 | 1961 | ML1 | MLN | Milwaukee Braves |
| 702 | 1962 | ML1 | MLN | Milwaukee Braves |
| 722 | 1963 | ML1 | MLN | Milwaukee Braves |
| 742 | 1964 | ML1 | MLN | Milwaukee Braves |
| 762 | 1965 | ML1 | MLN | Milwaukee Braves |
| 867 | 1970 | ML4 | MIL | Milwaukee Brewers |
| 891 | 1971 | ML4 | MIL | Milwaukee Brewers |
| 915 | 1972 | ML4 | MIL | Milwaukee Brewers |
| 939 | 1973 | ML4 | MIL | Milwaukee Brewers |
| 963 | 1974 | ML4 | MIL | Milwaukee Brewers |
| 987 | 1975 | ML4 | MIL | Milwaukee Brewers |
| 1011 | 1976 | ML4 | MIL | Milwaukee Brewers |
| 1035 | 1977 | ML4 | MIL | Milwaukee Brewers |
| 1061 | 1978 | ML4 | MIL | Milwaukee Brewers |
| 1087 | 1979 | ML4 | MIL | Milwaukee Brewers |
| 1113 | 1980 | ML4 | MIL | Milwaukee Brewers |
| 1139 | 1981 | ML4 | MIL | Milwaukee Brewers |
| 1165 | 1982 | ML4 | MIL | Milwaukee Brewers |
| 1191 | 1983 | ML4 | MIL | Milwaukee Brewers |
| 1217 | 1984 | ML4 | MIL | Milwaukee Brewers |
| 1243 | 1985 | ML4 | MIL | Milwaukee Brewers |
| 1269 | 1986 | ML4 | MIL | Milwaukee Brewers |
| 1295 | 1987 | ML4 | MIL | Milwaukee Brewers |
| 1321 | 1988 | ML4 | MIL | Milwaukee Brewers |

| | | | | |
|------|------|-----|-----------------|-------------------|
| 1347 | 1989 | ML4 | MIL | Milwaukee Brewers |
| 1373 | 1990 | ML4 | MIL | Milwaukee Brewers |
| 1399 | 1991 | ML4 | MIL | Milwaukee Brewers |
| 1425 | 1992 | ML4 | MIL | Milwaukee Brewers |
| 1453 | 1993 | ML4 | MIL | Milwaukee Brewers |
| 1481 | 1994 | ML4 | MIL | Milwaukee Brewers |
| 1509 | 1995 | ML4 | MIL | Milwaukee Brewers |
| 1537 | 1996 | ML4 | MIL | Milwaukee Brewers |
| 1565 | 1997 | ML4 | MIL | Milwaukee Brewers |
| 1801 | 2005 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1831 | 2006 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1861 | 2007 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1891 | 2008 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1921 | 2009 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1951 | 2010 | LAA | ANA Los Angeles | Angels of Anaheim |
| 1981 | 2011 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2010 | 2012 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2040 | 2013 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2070 | 2014 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2100 | 2015 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2130 | 2016 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2160 | 2017 | LAA | ANA Los Angeles | Angels of Anaheim |
| 2190 | 2018 | LAA | ANA Los Angeles | Angels of Anaheim |

So clearly we have three teams where the IDs differ. We need to ask a few questions though:

Do they differ on those teams every time? We can't just take that for granted.

```
[10]: df[df['franchID'] == 'ANA']['teamID'].value_counts()
```

```
[10]: CAL    32
      LAA    18
      ANA     8
      Name: teamID, dtype: int64
```

```
[11]: df[(df['teamID'] != df['teamIDretro'])[['teamID', 'teamIDretro', 'name']].
      ↪shape[0]
```

```
[11]: 55
```

```
[12]: df[(df['teamID'] == 'ML1').shape[0] + df[(df['teamID'] == 'ML4').shape[0] +
      ↪df[(df['teamID'] == 'LAA').shape[0]
```

```
[12]: 59
```

Unfortunately we have a disparity of 4, so we need to find out where that is.

```
[13]: df[(df['teamID'] == 'ML1') & (df['teamID'] == df['teamIDretro'])]
```

```
[13]: Empty DataFrame
      Columns: [yearID, lgID, teamID, franchID, divID, Rank, G, Ghome, W, L, DivWin,
      WCWin, LgWin, WSWin, R, AB, H, 2B, 3B, HR, BB, SO, SB, CS, HBP, SF, RA, ER, ERA,
      CG, SHO, SV, IPouts, HA, HRA, BBA, SOA, E, DP, FP, name, park, attendance, BPF,
      PPF, teamIDretro]
      Index: []

      [0 rows x 46 columns]
```

```
[14]: df[(df['teamID'] == 'ML4') & (df['teamID'] == df['teamIDretro'])]
```

```
[14]: Empty DataFrame
      Columns: [yearID, lgID, teamID, franchID, divID, Rank, G, Ghome, W, L, DivWin,
      WCWin, LgWin, WSWin, R, AB, H, 2B, 3B, HR, BB, SO, SB, CS, HBP, SF, RA, ER, ERA,
      CG, SHO, SV, IPouts, HA, HRA, BBA, SOA, E, DP, FP, name, park, attendance, BPF,
      PPF, teamIDretro]
      Index: []

      [0 rows x 46 columns]
```

```
[15]: df[(df['teamID'] == 'LAA') & (df['teamID'] == df['teamIDretro'])]
```

```
[15]:
```

| | yearID | lgID | teamID | franchID | divID | Rank | G | Ghome | W | L | ... | SOA | \ |
|-----|--------|------|--------|----------|-------|------|-----|-------|----|----|-----|-----|---|
| 680 | 1961 | AL | LAA | ANA | NaN | 8 | 162 | 82 | 70 | 91 | ... | 973 | |
| 699 | 1962 | AL | LAA | ANA | NaN | 3 | 162 | 81 | 86 | 76 | ... | 858 | |
| 719 | 1963 | AL | LAA | ANA | NaN | 9 | 161 | 81 | 70 | 91 | ... | 889 | |
| 739 | 1964 | AL | LAA | ANA | NaN | 5 | 162 | 81 | 82 | 80 | ... | 965 | |

| | E | DP | FP | | name | | park | attendance | BPF | \ |
|-----|-----|-----|-------|-------------|--------|--------------------|------|------------|-----|---|
| 680 | 192 | 154 | 0.969 | Los Angeles | Angels | Wrigley Field (LA) | | 603510 | 111 | |
| 699 | 175 | 153 | 0.972 | Los Angeles | Angels | Dodger Stadium | | 1144063 | 97 | |
| 719 | 163 | 155 | 0.974 | Los Angeles | Angels | Dodger Stadium | | 821015 | 94 | |
| 739 | 138 | 168 | 0.978 | Los Angeles | Angels | Dodger Stadium | | 760439 | 90 | |

| | PPF | teamIDretro |
|-----|-----|-------------|
| 680 | 112 | LAA |
| 699 | 97 | LAA |
| 719 | 94 | LAA |
| 739 | 90 | LAA |

```

      [4 rows x 46 columns]
```

```
[16]: df[(df['teamID'] == 'LAA') & (df['teamID'] != df['teamIDretro'])]
```

```
[16]:
```

| | yearID | lgID | teamID | franchID | divID | Rank | G | Ghome | W | L | ... | SOA | \ |
|------|--------|------|--------|----------|-------|------|-----|-------|----|----|-----|------|---|
| 1801 | 2005 | AL | LAA | ANA | W | 1 | 162 | 81 | 95 | 67 | ... | 1126 | |
| 1831 | 2006 | AL | LAA | ANA | W | 2 | 162 | 81 | 89 | 73 | ... | 1164 | |

| | | | | | | | | | | | | |
|------|------|----|-----|-----|---|---|-----|----|-----|----|-----|------|
| 1861 | 2007 | AL | LAA | ANA | W | 1 | 162 | 81 | 94 | 68 | ... | 1156 |
| 1891 | 2008 | AL | LAA | ANA | W | 1 | 162 | 81 | 100 | 62 | ... | 1106 |
| 1921 | 2009 | AL | LAA | ANA | W | 1 | 162 | 81 | 97 | 65 | ... | 1062 |
| 1951 | 2010 | AL | LAA | ANA | W | 3 | 162 | 81 | 80 | 82 | ... | 1130 |
| 1981 | 2011 | AL | LAA | ANA | W | 2 | 162 | 81 | 86 | 76 | ... | 1058 |
| 2010 | 2012 | AL | LAA | ANA | W | 3 | 162 | 81 | 89 | 73 | ... | 1157 |
| 2040 | 2013 | AL | LAA | ANA | W | 3 | 162 | 81 | 78 | 84 | ... | 1200 |
| 2070 | 2014 | AL | LAA | ANA | W | 1 | 162 | 81 | 98 | 64 | ... | 1342 |
| 2100 | 2015 | AL | LAA | ANA | W | 3 | 162 | 81 | 85 | 77 | ... | 1221 |
| 2130 | 2016 | AL | LAA | ANA | W | 4 | 162 | 81 | 74 | 88 | ... | 1136 |
| 2160 | 2017 | AL | LAA | ANA | W | 2 | 162 | 81 | 80 | 82 | ... | 1312 |
| 2190 | 2018 | AL | LAA | ANA | W | 4 | 162 | 81 | 80 | 82 | ... | 1386 |

| | E | DP | FP | name \ | | | | | | | | |
|------|-----|-----|-------|-------------------------------|--|--|--|--|--|--|--|--|
| 1801 | 87 | 139 | 0.986 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1831 | 124 | 154 | 0.979 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1861 | 101 | 154 | 0.983 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1891 | 91 | 159 | 0.985 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1921 | 85 | 174 | 0.986 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1951 | 113 | 116 | 0.981 | Los Angeles Angels of Anaheim | | | | | | | | |
| 1981 | 93 | 157 | 0.985 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2010 | 98 | 141 | 0.984 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2040 | 112 | 135 | 0.981 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2070 | 83 | 127 | 0.986 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2100 | 93 | 108 | 0.984 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2130 | 97 | 148 | 0.983 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2160 | 80 | 135 | 0.986 | Los Angeles Angels of Anaheim | | | | | | | | |
| 2190 | 76 | 173 | 0.987 | Los Angeles Angels of Anaheim | | | | | | | | |

| | | park | attendance | BPF | PPF | teamIDretro |
|------|--------------------------|---------------|------------|-----|-----|-------------|
| 1801 | | Angel Stadium | 3404686 | 98 | 97 | ANA |
| 1831 | | Angel Stadium | 3406790 | 100 | 100 | ANA |
| 1861 | | Angel Stadium | 3365632 | 101 | 100 | ANA |
| 1891 | | Angel Stadium | 3336747 | 103 | 102 | ANA |
| 1921 | | Angel Stadium | 3240386 | 99 | 98 | ANA |
| 1951 | | Angel Stadium | 3250816 | 98 | 98 | ANA |
| 1981 | | Angel Stadium | 3166321 | 93 | 93 | ANA |
| 2010 | Angel Stadium of Anaheim | | 3061770 | 92 | 92 | ANA |
| 2040 | Angel Stadium of Anaheim | | 3019505 | 94 | 94 | ANA |
| 2070 | Angel Stadium of Anaheim | | 3095935 | 96 | 95 | ANA |
| 2100 | Angel Stadium of Anaheim | | 3012765 | 94 | 95 | ANA |
| 2130 | Angel Stadium of Anaheim | | 3016142 | 95 | 95 | ANA |
| 2160 | Angel Stadium of Anaheim | | 3019585 | 96 | 96 | ANA |
| 2190 | Angel Stadium of Anaheim | | 3020216 | 97 | 97 | ANA |

[14 rows x 46 columns]

```
[17]: df['franchID'].unique()
```

```
[17]: array(['BOS', 'LAD', 'ATL', 'CHW', 'CHC', 'CIN', 'CLE', 'DET', 'SFG',  
        'NYY', 'OAK', 'PHI', 'PIT', 'BAL', 'STL', 'MIN', 'ANA', 'TEX',  
        'HOU', 'NYM', 'KCR', 'WSN', 'SDP', 'MIL', 'SEA', 'TOR', 'COL',  
        'FLA', 'ARI', 'TBD'], dtype=object)
```

```
[18]: df[(df['franchID'].isnull())]
```

```
[18]: Empty DataFrame  
Columns: [yearID, lgID, teamID, franchID, divID, Rank, G, Ghome, W, L, DivWin,  
WCWin, LgWin, WSWin, R, AB, H, 2B, 3B, HR, BB, SO, SB, CS, HBP, SF, RA, ER, ERA,  
CG, SHO, SV, IPouts, HA, HRA, BBA, SOA, E, DP, FP, name, park, attendance, BPF,  
PPF, teamIDretro]  
Index: []  
  
[0 rows x 46 columns]
```

```
[19]: df['franchID'].nunique()
```

```
[19]: 30
```

It looks like it will be easiest to just use the franchise ID - they stay consistent throughout and there are only ever 30 max. We'll need a way to map to these values from an external script so we can use it in other files.

Building Tables and Tensors

Aggregating preprocessed data and compiling it into tables that are ready to be read
into the models as tensors

build_tables

March 28, 2020

```
[88]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Building the Tables

We've done the major preprocessing in other scripts, and now it's time to get our final tables together for fielders, catchers and pitchers with all appropriate stats.

```
[89]: df_bat = pd.read_csv('../core/output/batting_pre.csv')
df_field = pd.read_csv('../core/output/fielding_pre.csv')
df_catch = pd.read_csv('../core/output/catching_pre.csv')
df_pitch = pd.read_csv('../core/output/pitching_pre.csv')
df_meta = pd.read_csv('../core/output/metadata.csv')
```

```
[90]: df_meta.head()
```

```
[90]:
```

| | retroID | POS | birthYear | bats | throws | weight | height | debutYear | finalYear |
|---|----------|-----|-----------|------|--------|--------|--------|-----------|-----------|
| 0 | aardd001 | P | 1981.0 | R | R | 215.0 | 75.0 | 2004 | 2015 |
| 1 | aaroh101 | OF | 1934.0 | R | R | 180.0 | 72.0 | 1954 | 1976 |
| 2 | aarot101 | 1B | 1939.0 | R | R | 190.0 | 75.0 | 1962 | 1971 |
| 3 | aased001 | P | 1954.0 | R | R | 190.0 | 75.0 | 1977 | 1990 |
| 4 | abada001 | 1B | 1972.0 | L | L | 184.0 | 73.0 | 2001 | 2006 |

Making Metadata Usable

We are interested in all of these fields, so we want to convert POS, bats and throws to numbers and use dummy variables. Note that we won't be using birthYear as-is, but rather subtracting it from current year to get a player's age for a season. This won't matter for the player career stats tensor so we can drop it here.

```
[91]: df_meta.drop(columns=['birthYear'], inplace=True)
```

```
[92]: df_meta_pos = pd.get_dummies(df_meta['POS'], prefix='pos')
df_meta_bats = pd.get_dummies(df_meta['bats'], drop_first=True, prefix='bats')
df_meta_throws = pd.get_dummies(df_meta['throws'], prefix='throws')
```

```
[93]: dropped_meta_cols = ['POS', 'bats', 'throws']
df_meta.drop(columns=dropped_meta_cols, inplace=True)
```



```
df_meta.head()
```

```
[93]:
```

| | retroID | weight | height | debutYear | finalYear |
|---|----------|--------|--------|-----------|-----------|
| 0 | aardd001 | 215.0 | 75.0 | 2004 | 2015 |
| 1 | aaroh101 | 180.0 | 72.0 | 1954 | 1976 |
| 2 | aarot101 | 190.0 | 75.0 | 1962 | 1971 |
| 3 | aased001 | 190.0 | 75.0 | 1977 | 1990 |
| 4 | abada001 | 184.0 | 73.0 | 2001 | 2006 |

```
[94]: df_meta = df_meta.join([df_meta_pos, df_meta_bats, df_meta_throws])
df_meta
```

```
[94]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | pos_3B | \ |
|-------|----------|--------|--------|-----------|-----------|--------|--------|--------|---|
| 0 | aardd001 | 215.0 | 75.0 | 2004 | 2015 | 0 | 0 | 0 | |
| 1 | aaroh101 | 180.0 | 72.0 | 1954 | 1976 | 0 | 0 | 0 | |
| 2 | aarot101 | 190.0 | 75.0 | 1962 | 1971 | 1 | 0 | 0 | |
| 3 | aased001 | 190.0 | 75.0 | 1977 | 1990 | 0 | 0 | 0 | |
| 4 | abada001 | 184.0 | 73.0 | 2001 | 2006 | 1 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15026 | zupcb001 | 220.0 | 76.0 | 1991 | 1994 | 0 | 0 | 0 | |
| 15027 | zupof101 | 182.0 | 71.0 | 1957 | 1961 | 0 | 0 | 0 | |
| 15028 | zuveg101 | 195.0 | 76.0 | 1951 | 1959 | 0 | 0 | 0 | |
| 15029 | zuvep001 | 173.0 | 72.0 | 1982 | 1991 | 0 | 0 | 0 | |
| 15030 | zycht001 | 190.0 | 75.0 | 2015 | 2017 | 0 | 0 | 0 | |

| | pos_C | pos_OF | pos_P | pos_SS | bats_L | bats_R | throws_L | throws_R | \ |
|-------|-------|--------|-------|--------|--------|--------|----------|----------|---|
| 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 1 | |
| 1 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 1 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | |
| 3 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 1 | |
| 4 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15026 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 1 | |
| 15027 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | |
| 15028 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 1 | |
| 15029 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | |
| 15030 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 1 | |

| | throws_S |
|-------|----------|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 0 |
| ... | ... |
| 15026 | 0 |
| 15027 | 0 |

```
15028      0
15029      0
15030      0
```

```
[15031 rows x 17 columns]
```

I didn't drop_first on the 'throws_' columns because I want to get rid of 'throws_S' instead of 'throws_L'

```
[95]: df_meta.drop(columns=['throws_S'], inplace=True)
```

We want to use weight and height but we can normalize them

```
[96]: from sklearn.preprocessing import MinMaxScaler
```

```
[97]: scaler = MinMaxScaler()
```

```
[98]: df_meta[['weight', 'height']] = scaler.fit_transform(df_meta[['weight', 'height']])
df_meta
```

```
[98]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | \ |
|-------|----------|----------|--------|-----------|-----------|--------|--------|---|
| 0 | aardd001 | 0.569672 | 0.60 | 2004 | 2015 | 0 | 0 | |
| 1 | aaroh101 | 0.426230 | 0.45 | 1954 | 1976 | 0 | 0 | |
| 2 | aarot101 | 0.467213 | 0.60 | 1962 | 1971 | 1 | 0 | |
| 3 | aased001 | 0.467213 | 0.60 | 1977 | 1990 | 0 | 0 | |
| 4 | abada001 | 0.442623 | 0.50 | 2001 | 2006 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 15026 | zupcb001 | 0.590164 | 0.65 | 1991 | 1994 | 0 | 0 | |
| 15027 | zupof101 | 0.434426 | 0.40 | 1957 | 1961 | 0 | 0 | |
| 15028 | zuveg101 | 0.487705 | 0.65 | 1951 | 1959 | 0 | 0 | |
| 15029 | zuvep001 | 0.397541 | 0.45 | 1982 | 1991 | 0 | 0 | |
| 15030 | zycht001 | 0.467213 | 0.60 | 2015 | 2017 | 0 | 0 | |

| | pos_3B | pos_C | pos_OF | pos_P | pos_SS | bats_L | bats_R | throws_L | \ |
|-------|--------|-------|--------|-------|--------|--------|--------|----------|---|
| 0 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 1 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 3 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 4 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15026 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | |
| 15027 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 15028 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 15029 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | |
| 15030 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | |

```

        throws_R
0          1
1          1
2          1
3          1
4          0
...
15026      1
15027      1
15028      1
15029      1
15030      1

```

[15031 rows x 16 columns]

The metadata is now ready to go into the final tensor.

Combining Batting Data

```
[99]: df_bat
```

```

[99]:      retroID    G    AB    R    H    2B    3B    HR    RBI    SB    CS    BB  \
0      aarodd001  331     4     0     0     0     0     0     0     0    0.0     0
1      aaroh101 3298 12364 2174 3771  624    98   755  2297  240  73.0 1402
2      aarot101  437   944   102   216   42     6    13    94     9    8.0    86
3      aased001  448     5     0     0     0     0     0     0     0    0.0     0
4      abada001   15    21     1     2     0     0     0     0     0    1.0     4
...
15187  zupcb001  319   795    99   199   47     4     7    80     7    5.0    57
15188  zupof101   16    18     3     3     1     0     0     0     0    0.0     2
15189  zuveg101  266   142     5    21     2     1     0     7     0    1.0     9
15190  zuveg001  209   491    41   109   17     2     2    20     2    0.0    34
15191  zycht001   70     0     0     0     0     0     0     0     0    0.0     0

```

```

        S0  IBB  HBP  SH  SF  GIDP  NL
0         2    0    0   1   0     0   1
1      1383  293   32  21  121   328   1
2      145    3    0   9   6    36   1
3         3    0    0   0   0     0   1
4         5    0    0   0   0     1   1
...
15187   137    3    6  20   8    15   0
15188    6    0    0   0   0     0   0
15189   39    0    0  16   0     3   1
15190   50    1    2  18   0     8   1
15191    0    0    0   0   0     0   0

```

[15192 rows x 19 columns]

```
[100]: df = pd.merge(df_meta, df_bat, how='inner', on=['retroID'])
```

```
[187]: df.shape
```

```
[187]: (15031, 34)
```

```
[101]: df.head(10)
```

```
[101]:
```

| | retroID | weight | height | debutYear | finalYear | pos_1B | pos_2B | pos_3B | \ |
|---|----------|----------|--------|-----------|-----------|--------|--------|--------|---|
| 0 | aardd001 | 0.569672 | 0.60 | 2004 | 2015 | 0 | 0 | 0 | |
| 1 | aaroh101 | 0.426230 | 0.45 | 1954 | 1976 | 0 | 0 | 0 | |
| 2 | aarot101 | 0.467213 | 0.60 | 1962 | 1971 | 1 | 0 | 0 | |
| 3 | aased001 | 0.467213 | 0.60 | 1977 | 1990 | 0 | 0 | 0 | |
| 4 | abada001 | 0.442623 | 0.50 | 2001 | 2006 | 1 | 0 | 0 | |
| 5 | abadf001 | 0.590164 | 0.50 | 2010 | 2017 | 0 | 0 | 0 | |
| 6 | abbog001 | 0.508197 | 0.75 | 1973 | 1984 | 0 | 0 | 0 | |
| 7 | abboj001 | 0.508197 | 0.60 | 1989 | 1999 | 0 | 0 | 0 | |
| 8 | abboj002 | 0.467213 | 0.55 | 1997 | 2001 | 0 | 0 | 0 | |
| 9 | abbok001 | 0.508197 | 0.65 | 1991 | 1996 | 0 | 0 | 0 | |

| | pos_C | pos_OF | ... | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
|---|-------|--------|-----|-----|------|------|------|-----|-----|----|-----|------|----|
| 0 | 0 | 0 | ... | 0 | 0.0 | 0 | 2 | 0 | 0 | 1 | 0 | 0 | 1 |
| 1 | 0 | 1 | ... | 240 | 73.0 | 1402 | 1383 | 293 | 32 | 21 | 121 | 328 | 1 |
| 2 | 0 | 0 | ... | 9 | 8.0 | 86 | 145 | 3 | 0 | 9 | 6 | 36 | 1 |
| 3 | 0 | 0 | ... | 0 | 0.0 | 0 | 3 | 0 | 0 | 0 | 0 | 0 | 1 |
| 4 | 0 | 0 | ... | 0 | 1.0 | 4 | 5 | 0 | 0 | 0 | 0 | 1 | 1 |
| 5 | 0 | 0 | ... | 0 | 0.0 | 0 | 5 | 0 | 0 | 0 | 0 | 1 | 1 |
| 6 | 0 | 0 | ... | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 7 | 0 | 0 | ... | 0 | 0.0 | 0 | 10 | 0 | 0 | 3 | 0 | 0 | 1 |
| 8 | 0 | 1 | ... | 6 | 5.0 | 38 | 91 | 2 | 3 | 5 | 7 | 12 | 1 |
| 9 | 0 | 0 | ... | 0 | 0.0 | 1 | 19 | 0 | 0 | 6 | 0 | 0 | 1 |

```
[10 rows x 34 columns]
```

We noticed that `df_bat` and `df_meta` don't have the same number of rows, so we want to find out what's going on there.

```
[102]: df_bat[~df_bat['retroID'].isin(df_meta['retroID'])]
```

```
[102]:
```

| | retroID | G | AB | R | H | 2B | 3B | HR | RBI | SB | CS | BB | SO | IBB | HBP | SH | \ |
|-------|----------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---|
| 121 | albeb101 | 6 | 18 | 1 | 5 | 1 | 0 | 0 | 0 | 0 | 0.0 | 0 | 2 | 0 | 0 | 0 | |
| 358 | aragj101 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | |
| 445 | atkil101 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | |
| 611 | banij001 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | |
| 633 | barbr101 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 14543 | westj101 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 1 | 0 | 0 | 0 | |

| | | | | | | | | | | | | | | | | |
|-------|----------|----|---|---|---|---|---|---|---|---|-----|---|---|---|---|---|
| 14730 | willh101 | 10 | 9 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0.0 | 1 | 4 | 0 | 0 | 0 |
| 14811 | wilsi101 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 0 | 0 | 0 | 0 |
| 15009 | wrigr002 | 1 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0 | 1 | 0 | 0 | 0 |
| 15050 | yeabb101 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 1 | 0 | 0 | 0 | 0 |

| | | | |
|-------|----|------|----|
| | SF | GIDP | NL |
| 121 | 0 | 0 | 0 |
| 358 | 0 | 0 | 1 |
| 445 | 0 | 0 | 0 |
| 611 | 0 | 0 | 1 |
| 633 | 0 | 0 | 0 |
| ... | .. | ... | .. |
| 14543 | 0 | 0 | 1 |
| 14730 | 0 | 0 | 1 |
| 14811 | 0 | 0 | 0 |
| 15009 | 0 | 1 | 0 |
| 15050 | 0 | 0 | 1 |

[161 rows x 19 columns]

```
[103]: df_bat[~df_bat['retroID'].isin(df_meta['retroID'])]['G'].max()
```

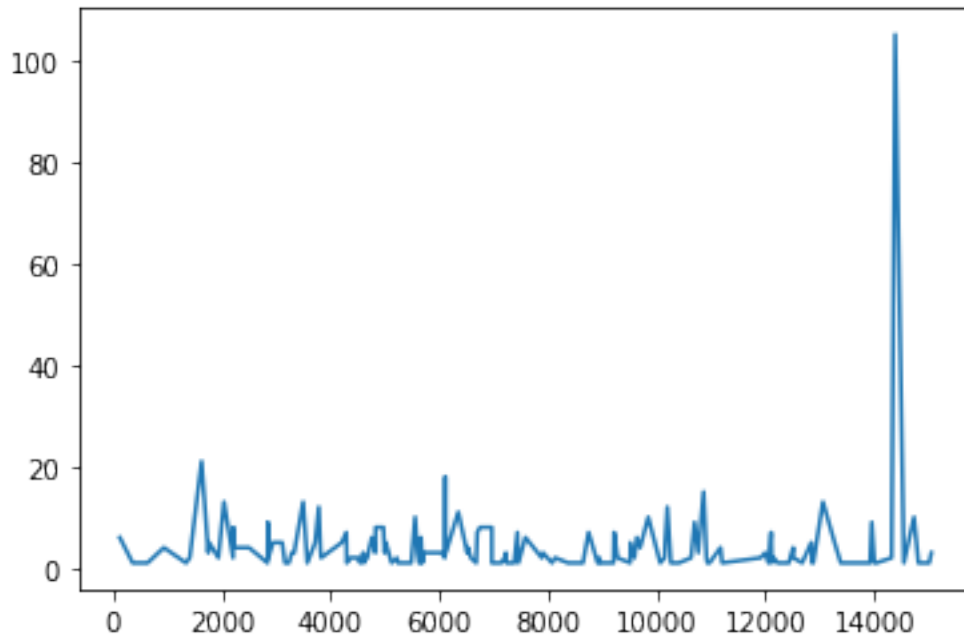
```
[103]: 105
```

```
[104]: df_bat[~df_bat['retroID'].isin(df_meta['retroID'])]['G'].mean()
```

```
[104]: 4.105590062111801
```

```
[105]: plt.plot(df_bat[~df_bat['retroID'].isin(df_meta['retroID'])]['G'])
```

```
[105]: [<matplotlib.lines.Line2D at 0x11ffb5bd0>]
```



For the most part, we're talking about players who have played under 20 total games. We can easily drop these data points and not really affect the overall result.

Combining the Tensors

Catchers

```
[106]: df_catch.shape[0] + df_field.shape[0] + df_pitch.shape[0]
```

```
[106]: 23591
```

```
[107]: df_catch
```

```
[107]:
```

| | retroID | GS | InnOuts | PO | A | E | DP | PB | WP | SB | CS | ZR |
|------|----------|-----|---------|------|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | adamb105 | 1 | 27 | 6 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 1 | adamb106 | 0 | 0 | 249 | 90 | 12 | 15 | 7 | 0 | 0 | 0 | 0 |
| 2 | adamd101 | 3 | 78 | 9 | 2 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 3 | adled101 | 65 | 1840 | 453 | 26 | 4 | 2 | 8 | 19 | 37 | 16 | 0 |
| 4 | afent001 | 20 | 613 | 123 | 5 | 1 | 3 | 6 | 0 | 17 | 3 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1524 | zimmd101 | 27 | 744 | 150 | 18 | 6 | 1 | 5 | 12 | 10 | 10 | 3 |
| 1525 | zimmj101 | 298 | 8560 | 2131 | 150 | 21 | 26 | 19 | 84 | 110 | 80 | 4 |
| 1526 | zinta001 | 0 | 3 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1527 | zunim001 | 535 | 14489 | 4356 | 264 | 21 | 22 | 39 | 0 | 248 | 98 | 0 |
| 1528 | zupof101 | 1 | 114 | 31 | 1 | 2 | 0 | 1 | 1 | 2 | 1 | 0 |

```
[1529 rows x 12 columns]
```

```
[108]: np.intersect1d(df_catch.columns, df.columns)
```

```
[108]: array(['CS', 'SB', 'retroID'], dtype=object)
```

The 'caught stealing' and 'stolen bases' stats appear both offensively and defensively (CS/SB against) for catchers. We need to keep them separate when merging the metadata and we can do so by just adding a prefix to the defensive stats.

```
[109]: df_catch.rename(columns={'CS': 'CS_A', 'SB': 'SB_A'}, inplace=True)
```

```
[110]: catchers = pd.merge(df_catch, df, how='inner', on=['retroID'])
```

```
[111]: catchers
```

```
[111]:
```

| | retroID | GS | InnOuts | PO | A | E | DP | PB | WP | SB_A | ... | SB | CS | \ |
|------|----------|-----|---------|------|-----|-----|-----|-----|-----|------|-----|-----|------|---|
| 0 | adamb105 | 1 | 27 | 6 | 0 | 0 | 0 | 0 | 0 | 1 | ... | 0 | 0.0 | |
| 1 | adamb106 | 0 | 0 | 249 | 90 | 12 | 15 | 7 | 0 | 0 | ... | 4 | 2.0 | |
| 2 | adamd101 | 3 | 78 | 9 | 2 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0.0 | |
| 3 | adled101 | 65 | 1840 | 453 | 26 | 4 | 2 | 8 | 19 | 37 | ... | 0 | 0.0 | |
| 4 | afent001 | 20 | 613 | 123 | 5 | 1 | 3 | 6 | 0 | 17 | ... | 0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1524 | zimmd101 | 27 | 744 | 150 | 18 | 6 | 1 | 5 | 12 | 10 | ... | 45 | 25.0 | |
| 1525 | zimmj101 | 298 | 8560 | 2131 | 150 | 21 | 26 | 19 | 84 | 110 | ... | 1 | 2.0 | |
| 1526 | zinta001 | 0 | 3 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | |
| 1527 | zunim001 | 535 | 14489 | 4356 | 264 | 21 | 22 | 39 | 0 | 248 | ... | 2 | 4.0 | |
| 1528 | zupof101 | 1 | 114 | 31 | 1 | 2 | 0 | 1 | 1 | 2 | ... | 0 | 0.0 | |

| | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
|------|-----|-----|-----|-----|-----|-----|------|-----|
| 0 | 0 | 5 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 6 | 27 | 0 | 0 | 3 | 0 | 0 | 1 |
| 2 | 1 | 3 | 0 | 0 | 0 | 0 | 1 | 0 |
| 3 | 18 | 80 | 5 | 2 | 2 | 1 | 9 | 1 |
| 4 | 5 | 32 | 0 | 0 | 1 | 1 | 1 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1524 | 246 | 678 | 27 | 13 | 37 | 14 | 99 | 1 |
| 1525 | 78 | 154 | 12 | 11 | 31 | 4 | 38 | 1 |
| 1526 | 5 | 34 | 0 | 0 | 0 | 1 | 0 | 1 |
| 1527 | 138 | 714 | 1 | 45 | 8 | 11 | 38 | 0 |
| 1528 | 2 | 6 | 0 | 0 | 0 | 0 | 0 | 0 |

[1529 rows x 45 columns]

```
[112]: catchers.columns
```

```
[112]: Index(['retroID', 'GS', 'InnOuts', 'PO', 'A', 'E', 'DP', 'PB', 'WP', 'SB_A',
        'CS_A', 'ZR', 'weight', 'height', 'debutYear', 'finalYear', 'pos_1B',
        'pos_2B', 'pos_3B', 'pos_C', 'pos_OF', 'pos_P', 'pos_SS', 'bats_L',
```

```
'bats_R', 'throws_L', 'throws_R', 'G', 'AB', 'R', 'H', '2B', '3B', 'HR',
'RBI', 'SB', 'CS', 'BB', 'SO', 'IBB', 'HBP', 'SH', 'SF', 'GIDP', 'NL'],
dtype='object')
```

There's no reason to waste columns on position for the catchers.

```
[113]: catchers.drop(columns=['pos_1B', 'pos_2B', 'pos_3B', 'pos_C',
                             'pos_OF', 'pos_P', 'pos_SS'], inplace=True)
```

```
[114]: catchers.columns
```

```
[114]: Index(['retroID', 'GS', 'InnOuts', 'PO', 'A', 'E', 'DP', 'PB', 'WP', 'SB_A',
             'CS_A', 'ZR', 'weight', 'height', 'debutYear', 'finalYear', 'bats_L',
             'bats_R', 'throws_L', 'throws_R', 'G', 'AB', 'R', 'H', '2B', '3B', 'HR',
             'RBI', 'SB', 'CS', 'BB', 'SO', 'IBB', 'HBP', 'SH', 'SF', 'GIDP', 'NL'],
            dtype='object')
```

Pitchers

```
[147]: np.intersect1d(df_pitch.columns, df.columns)
```

```
[147]: array(['G', 'retroID'], dtype=object)
```

We have quite a few common columns for pitching data and metadata. We'll do what we did for catching and just add 'A' to the end (for 'against'). Since there are quite a few, we'll define a conversion dictionary ahead of time. Before we do that, we see that we can drop the 'G' (games) column as it should be the same between the tables. We can also drop the position information.

```
[148]: batting_data_for_pitchers = df.drop(columns=['G', 'pos_1B', 'pos_2B',
                                                    'pos_3B', 'pos_C', 'pos_OF', 'pos_P', 'pos_SS'])
```

```
[149]: pitching_data_conversion_dict = {
        'BB': 'BB_A',
        'GIDP': 'GIDP_A',
        'H': 'H_A',
        'HBP': 'HBP_A',
        'HR': 'HR_A',
        'IBB': 'IBB_A',
        'R': 'R_A',
        'SF': 'SF_A',
        'SH': 'SH_A',
        'SO': 'SO_A'
    }
```

```
[150]: df_pitch.rename(columns=pitching_data_conversion_dict, inplace=True)
```

```
[151]: pitchers = pd.merge(df_pitch, batting_data_for_pitchers, how='inner',
                           on=['retroID'])
```



```
[152]: pitchers
```

```
[152]:
```

| | retroID | BAOpp | ERA | W | L | G | GS | CG | SHO | SV | ... | SB | CS | \ |
|------|----------|--------|--------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---|
| 0 | aardd001 | 0.2574 | 5.1944 | 16 | 18 | 331 | 0 | 0 | 0 | 69 | ... | 0 | 0.0 | |
| 1 | aased001 | 0.2508 | 3.4931 | 66 | 60 | 448 | 91 | 22 | 5 | 82 | ... | 0 | 0.0 | |
| 2 | abadf001 | 0.2501 | 4.0733 | 8 | 27 | 363 | 6 | 0 | 0 | 2 | ... | 0 | 0.0 | |
| 3 | abbog001 | 0.2786 | 4.3317 | 62 | 83 | 248 | 206 | 37 | 5 | 0 | ... | 0 | 0.0 | |
| 4 | abboj001 | 0.2804 | 4.4964 | 87 | 108 | 263 | 254 | 31 | 6 | 0 | ... | 0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 7830 | zolds101 | 0.2700 | 3.6890 | 43 | 53 | 250 | 93 | 30 | 5 | 8 | ... | 1 | 0.0 | |
| 7831 | zubeb101 | 0.2717 | 5.3617 | 43 | 42 | 224 | 65 | 23 | 3 | 6 | ... | 0 | 0.0 | |
| 7832 | zumaj001 | 0.2286 | 3.4420 | 13 | 12 | 171 | 0 | 0 | 0 | 5 | ... | 0 | 0.0 | |
| 7833 | zuveg101 | 0.2760 | 4.1280 | 32 | 36 | 265 | 31 | 9 | 2 | 40 | ... | 0 | 1.0 | |
| 7834 | zycht001 | 0.2183 | 2.8000 | 7 | 3 | 70 | 1 | 0 | 0 | 1 | ... | 0 | 0.0 | |

| | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
|------|-----|-----|-----|-----|-----|-----|------|-----|
| 0 | 0 | 2 | 0 | 0 | 1 | 0 | 0 | 1 |
| 1 | 0 | 3 | 0 | 0 | 0 | 0 | 0 | 1 |
| 2 | 0 | 5 | 0 | 0 | 0 | 0 | 1 | 1 |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 10 | 0 | 0 | 3 | 0 | 0 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7830 | 10 | 52 | 0 | 1 | 9 | 0 | 4 | 0 |
| 7831 | 10 | 66 | 0 | 0 | 20 | 0 | 8 | 0 |
| 7832 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 7833 | 9 | 39 | 0 | 0 | 16 | 0 | 3 | 1 |
| 7834 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

[7835 rows x 51 columns]

Fielders

```
[32]: np.intersect1d(df_field.columns, df.columns)
```

```
[32]: array(['retroID'], dtype=object)
```

We don't need to worry about common columns between the general fielding stats and metadata.

```
[136]: fielders = pd.merge(df_field, df, how='inner', on=['retroID'])
```

```
[137]: fielders = fielders[~fielders['retroID'].isin(pitchers['retroID'])]
```

```
[138]: fielders
```

```
[138]:
```

| | retroID | GS | InnOuts | PO | A | E | DP | weight | height | \ |
|---|----------|------|---------|------|-----|-----|-----|----------|--------|---|
| 1 | aaroh101 | 2977 | 78414 | 7436 | 429 | 144 | 218 | 0.426230 | 0.45 | |
| 2 | aarot101 | 206 | 6472 | 1317 | 113 | 22 | 124 | 0.467213 | 0.60 | |
| 4 | abada001 | 4 | 138 | 37 | 1 | 1 | 3 | 0.442623 | 0.50 | |

| | | | | | | | | | |
|-------|----------|-----|-------|-----|------|-----|-----|----------|------|
| 8 | abboj002 | 140 | 3688 | 299 | 2 | 8 | 0 | 0.467213 | 0.55 |
| 10 | abbok002 | 504 | 13474 | 938 | 1262 | 79 | 275 | 0.426230 | 0.40 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14218 | zoske001 | 8 | 404 | 16 | 42 | 2 | 8 | 0.405738 | 0.45 |
| 14220 | zubej001 | 26 | 702 | 167 | 12 | 2 | 11 | 0.467213 | 0.50 |
| 14221 | zulej001 | 36 | 1019 | 296 | 15 | 5 | 20 | 0.631148 | 0.75 |
| 14223 | zupcb001 | 198 | 5842 | 483 | 22 | 12 | 5 | 0.590164 | 0.65 |
| 14225 | zuvep001 | 136 | 3844 | 267 | 415 | 23 | 84 | 0.397541 | 0.45 |

| | debutYear | ... | SB | CS | BB | S0 | IBB | HBP | SH | SF | GIDP | NL |
|-------|-----------|-----|-----|------|------|------|-----|-----|-----|-----|------|-----|
| 1 | 1954 | ... | 240 | 73.0 | 1402 | 1383 | 293 | 32 | 21 | 121 | 328 | 1 |
| 2 | 1962 | ... | 9 | 8.0 | 86 | 145 | 3 | 0 | 9 | 6 | 36 | 1 |
| 4 | 2001 | ... | 0 | 1.0 | 4 | 5 | 0 | 0 | 0 | 0 | 1 | 1 |
| 8 | 1997 | ... | 6 | 5.0 | 38 | 91 | 2 | 3 | 5 | 7 | 12 | 1 |
| 10 | 1993 | ... | 22 | 11.0 | 133 | 571 | 11 | 17 | 21 | 12 | 37 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14218 | 1991 | ... | 0 | 0.0 | 1 | 13 | 0 | 0 | 1 | 1 | 1 | 1 |
| 14220 | 1996 | ... | 1 | 0.0 | 12 | 20 | 1 | 1 | 1 | 1 | 4 | 1 |
| 14221 | 2000 | ... | 0 | 2.0 | 10 | 51 | 1 | 6 | 0 | 1 | 5 | 1 |
| 14223 | 1991 | ... | 7 | 5.0 | 57 | 137 | 3 | 6 | 20 | 8 | 15 | 0 |
| 14225 | 1982 | ... | 2 | 0.0 | 34 | 50 | 1 | 2 | 18 | 0 | 8 | 1 |

[6392 rows x 40 columns]

[140]: catchers

[140]:

| | retroID | GS | InnOuts | P0 | A | E | DP | PB | WP | SB_A | ... | SB | CS | \ |
|------|----------|-----|---------|------|-----|-----|-----|-----|-----|------|-----|-----|------|---|
| 0 | adamb105 | 1 | 27 | 6 | 0 | 0 | 0 | 0 | 0 | 1 | ... | 0 | 0.0 | |
| 1 | adamb106 | 0 | 0 | 249 | 90 | 12 | 15 | 7 | 0 | 0 | ... | 4 | 2.0 | |
| 2 | adamd101 | 3 | 78 | 9 | 2 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0.0 | |
| 3 | adled101 | 65 | 1840 | 453 | 26 | 4 | 2 | 8 | 19 | 37 | ... | 0 | 0.0 | |
| 4 | afent001 | 20 | 613 | 123 | 5 | 1 | 3 | 6 | 0 | 17 | ... | 0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1524 | zimmd101 | 27 | 744 | 150 | 18 | 6 | 1 | 5 | 12 | 10 | ... | 45 | 25.0 | |
| 1525 | zimmj101 | 298 | 8560 | 2131 | 150 | 21 | 26 | 19 | 84 | 110 | ... | 1 | 2.0 | |
| 1526 | zinta001 | 0 | 3 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0.0 | |
| 1527 | zunim001 | 535 | 14489 | 4356 | 264 | 21 | 22 | 39 | 0 | 248 | ... | 2 | 4.0 | |
| 1528 | zupof101 | 1 | 114 | 31 | 1 | 2 | 0 | 1 | 1 | 2 | ... | 0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1524 | 246 | 678 | 27 | 13 | 37 | 14 | 99 | 1 | | | | | | |

| | BB | S0 | IBB | HBP | SH | SF | GIDP | NL |
|------|-----|-----|-----|-----|-----|-----|------|-----|
| 0 | 0 | 5 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 6 | 27 | 0 | 0 | 3 | 0 | 0 | 1 |
| 2 | 1 | 3 | 0 | 0 | 0 | 0 | 1 | 0 |
| 3 | 18 | 80 | 5 | 2 | 2 | 1 | 9 | 1 |
| 4 | 5 | 32 | 0 | 0 | 1 | 1 | 1 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1524 | 246 | 678 | 27 | 13 | 37 | 14 | 99 | 1 |

| | | | | | | | | |
|------|-----|-----|----|----|----|----|----|---|
| 1525 | 78 | 154 | 12 | 11 | 31 | 4 | 38 | 1 |
| 1526 | 5 | 34 | 0 | 0 | 0 | 1 | 0 | 1 |
| 1527 | 138 | 714 | 1 | 45 | 8 | 11 | 38 | 0 |
| 1528 | 2 | 6 | 0 | 0 | 0 | 0 | 0 | 0 |

[1529 rows x 38 columns]

```
[169]: fielders[fielders['retroID'].isin(catchers['retroID'])]
```

```
[169]:
```

| | retroID | GS | InnOuts | PO | A | E | DP | weight | height | \ |
|-------|----------|-----|---------|-------|------|-----|------|----------|--------|---|
| 54 | adamb105 | 2 | 54 | 20 | 1 | 0 | 1 | 0.508197 | 0.55 | |
| 55 | adamb106 | 0 | 0 | 0 | 0 | 0 | 0 | 0.446721 | 0.50 | |
| 108 | ainse101 | 0 | 0 | 11 | 0 | 0 | 0 | 0.426230 | 0.40 | |
| 144 | alexg101 | 22 | 500 | 83 | 6 | 3 | 4 | 0.487705 | 0.55 | |
| 151 | alfaj002 | 1 | 31 | 8 | 2 | 0 | 1 | 0.610656 | 0.55 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 14111 | yorkr101 | 0 | 0 | 11425 | 1030 | 136 | 1077 | 0.545082 | 0.50 | |
| 14139 | younj001 | 843 | 23486 | 1679 | 756 | 115 | 113 | 0.426230 | 0.45 | |
| 14183 | zaung001 | 0 | 33 | 3 | 2 | 0 | 1 | 0.385246 | 0.35 | |
| 14199 | zimmd101 | 813 | 21993 | 1491 | 2204 | 150 | 417 | 0.364754 | 0.30 | |
| 14210 | zinta001 | 5 | 198 | 65 | 5 | 1 | 4 | 0.508197 | 0.55 | |

| | debutYear | ... | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
|-------|-----------|-----|-----|------|-----|-----|-----|-----|-----|-----|------|-----|
| 54 | 1977 | ... | 0 | 0.0 | 0 | 5 | 0 | 0 | 0 | 0 | 0 | 0 |
| 55 | 1910 | ... | 4 | 2.0 | 6 | 27 | 0 | 0 | 3 | 0 | 0 | 1 |
| 108 | 1910 | ... | 20 | 11.5 | 125 | 122 | 0 | 3 | 44 | 0 | 0 | 1 |
| 144 | 1975 | ... | 8 | 12.0 | 154 | 381 | 12 | 5 | 4 | 19 | 34 | 1 |
| 151 | 2016 | ... | 3 | 0.0 | 22 | 179 | 8 | 18 | 0 | 1 | 4 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14111 | 1934 | ... | 38 | 26.0 | 792 | 867 | 0 | 12 | 25 | 0 | 155 | 0 |
| 14139 | 1976 | ... | 60 | 55.0 | 332 | 589 | 30 | 36 | 18 | 33 | 80 | 1 |
| 14183 | 1995 | ... | 23 | 19.0 | 479 | 544 | 30 | 29 | 14 | 31 | 87 | 1 |
| 14199 | 1954 | ... | 45 | 25.0 | 246 | 678 | 27 | 13 | 37 | 14 | 99 | 1 |
| 14210 | 2002 | ... | 0 | 0.0 | 5 | 34 | 0 | 0 | 0 | 1 | 0 | 1 |

[655 rows x 40 columns]

We have some catchers that are also in the fielders table.

```
[168]: fielders[(fielders['retroID'].isin(catchers['retroID']) & fielders['pos_C'] == 'C')]
```

```
[168]:
```

| | retroID | GS | InnOuts | PO | A | E | DP | weight | height | debutYear | \ |
|-----|----------|----|---------|----|---|---|----|----------|--------|-----------|---|
| 108 | ainse101 | 0 | 0 | 11 | 0 | 0 | 0 | 0.426230 | 0.40 | 1910 | |
| 144 | alexg101 | 22 | 500 | 83 | 6 | 3 | 4 | 0.487705 | 0.55 | 1975 | |
| 151 | alfaj002 | 1 | 31 | 8 | 2 | 0 | 1 | 0.610656 | 0.55 | 2016 | |
| 156 | allaa001 | 0 | 33 | 15 | 0 | 0 | 2 | 0.590164 | 0.70 | 1986 | |

| | | | | | | | | | | |
|-------|----------|-----|------|------|-----|-----|-----|----------|------|------|
| 169 | alleg001 | 1 | 54 | 4 | 4 | 0 | 2 | 0.446721 | 0.40 | 1979 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 13914 | wingi101 | 0 | 0 | 7 | 2 | 1 | 0 | 0.344262 | 0.35 | 1911 |
| 13956 | wockj001 | 249 | 6120 | 1429 | 90 | 17 | 133 | 0.467213 | 0.45 | 1974 |
| 14066 | wronr001 | 0 | 6 | 2 | 0 | 0 | 0 | 0.446721 | 0.50 | 1988 |
| 14074 | wyneb001 | 0 | 6 | 0 | 0 | 0 | 0 | 0.467213 | 0.50 | 1976 |
| 14183 | zaung001 | 0 | 33 | 3 | 2 | 0 | 1 | 0.385246 | 0.35 | 1995 |

| | | | | | | | | | | | |
|-------|-----|-----|------|-----|-----|-----|-----|-----|-----|------|-----|
| | ... | SB | CS | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
| 108 | ... | 20 | 11.5 | 125 | 122 | 0 | 3 | 44 | 0 | 0 | 1 |
| 144 | ... | 8 | 12.0 | 154 | 381 | 12 | 5 | 4 | 19 | 34 | 1 |
| 151 | ... | 3 | 0.0 | 22 | 179 | 8 | 18 | 0 | 1 | 4 | 1 |
| 156 | ... | 23 | 18.0 | 87 | 223 | 4 | 9 | 35 | 16 | 27 | 1 |
| 169 | ... | 3 | 7.0 | 130 | 192 | 3 | 5 | 15 | 11 | 35 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 13914 | ... | 17 | 16.0 | 121 | 84 | 0 | 4 | 36 | 0 | 0 | 1 |
| 13956 | ... | 5 | 11.0 | 277 | 278 | 14 | 7 | 5 | 12 | 52 | 1 |
| 14066 | ... | 1 | 0.0 | 5 | 41 | 2 | 1 | 2 | 2 | 3 | 1 |
| 14074 | ... | 10 | 13.0 | 626 | 428 | 41 | 17 | 58 | 36 | 119 | 0 |
| 14183 | ... | 23 | 19.0 | 479 | 544 | 30 | 29 | 14 | 31 | 87 | 1 |

[430 rows x 40 columns]

```
[170]: to_inspect = fielders[(fielders['retroID'].isin(catchers['retroID']) &
    ↳fielders['pos_C'] == 1)]['retroID']
```

```
[171]: catchers[catchers['retroID'].isin(to_inspect)]
```

```
[171]:
```

| | | | | | | | | | | | | | | |
|------|----------|------|---------|------|-----|-----|-----|-----|-----|------|-----|-----|------|---|
| | retroID | GS | InnOuts | PO | A | E | DP | PB | WP | SB_A | ... | SB | CS | \ |
| 6 | ainse101 | 0 | 0 | 1528 | 361 | 72 | 31 | 34 | 0 | 0 | ... | 20 | 11.5 | |
| 7 | alexg101 | 205 | 5481 | 1008 | 100 | 35 | 8 | 24 | 0 | 212 | ... | 8 | 12.0 | |
| 8 | alfaj002 | 130 | 3430 | 1135 | 71 | 13 | 9 | 14 | 0 | 72 | ... | 3 | 0.0 | |
| 10 | allaa001 | 453 | 11965 | 2395 | 208 | 52 | 24 | 41 | 0 | 292 | ... | 23 | 18.0 | |
| 11 | alleg001 | 342 | 8959 | 1762 | 146 | 32 | 24 | 27 | 0 | 233 | ... | 3 | 7.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1498 | wingi101 | 0 | 0 | 1716 | 536 | 79 | 53 | 36 | 0 | 0 | ... | 17 | 16.0 | |
| 1502 | wockj001 | 216 | 5957 | 1212 | 119 | 39 | 17 | 27 | 0 | 188 | ... | 5 | 11.0 | |
| 1508 | wronr001 | 47 | 1360 | 296 | 32 | 8 | 3 | 4 | 0 | 25 | ... | 1 | 0.0 | |
| 1509 | wyneb001 | 1164 | 31563 | 6281 | 583 | 75 | 88 | 61 | 0 | 708 | ... | 10 | 13.0 | |
| 1521 | zaung001 | 910 | 24700 | 6134 | 418 | 88 | 59 | 51 | 0 | 651 | ... | 23 | 19.0 | |

| | | | | | | | | |
|----|-----|-----|-----|-----|----|----|------|----|
| | BB | SO | IBB | HBP | SH | SF | GIDP | NL |
| 6 | 125 | 122 | 0 | 3 | 44 | 0 | 0 | 1 |
| 7 | 154 | 381 | 12 | 5 | 4 | 19 | 34 | 1 |
| 8 | 22 | 179 | 8 | 18 | 0 | 1 | 4 | 1 |
| 10 | 87 | 223 | 4 | 9 | 35 | 16 | 27 | 1 |
| 11 | 130 | 192 | 3 | 5 | 15 | 11 | 35 | 0 |

```

...    ...    ...    ...    ...    ..    ..    ...    ..
1498  121    84      0      4    36    0      0    1
1502  277    278    14      7    5    12     52    1
1508    5    41      2      1    2    2      3    1
1509  626    428    41     17   58   36    119    0
1521  479    544    30     29   14   31     87    1

```

[430 rows x 38 columns]

It looks like the information in the catchers table is a better indicator of the player's career.

```
[175]: fielders[fielders['retroID'] == 'alexg101'][['debutYear', 'finalYear']]
```

```
[175]:      debutYear  finalYear
144         1975         1981
```

```
[174]: catchers[catchers['retroID'] == 'alexg101'][['debutYear', 'finalYear']]
```

```
[174]:      debutYear  finalYear
7         1975         1981
```

Three years line up. So for any catcher who appears in the fielders table with his position as catcher, we're going to drop him from the fielders table and only use the catchers information. We'll keep catchers in the fielders table if they're in a different position.

```
[176]: fielders[fielders['pos_C'] == 1]
```

```
[176]:      retroID  GS  InnOuts  PO  A  E  DP  weight  height  debutYear  \
108   ainse101    0         0   11  0  0   0  0.426230   0.40         1910
144   alexg101   22        500   83  6  3   4  0.487705   0.55         1975
151   alfaj002    1         31    8  2  0   1  0.610656   0.55         2016
156   allaa001    0         33   15  0  0   2  0.590164   0.70         1986
169   alleg001    1         54    4  4  0   2  0.446721   0.40         1979
...      ...    ...      ...    ...    ..    ..    ...      ...      ...      ...
13914 wingi101    0         0    7  2  1   0  0.344262   0.35         1911
13956 wockj001  249       6120  1429  90  17  133  0.467213   0.45         1974
14066 wronr001    0         6    2  0  0   0  0.446721   0.50         1988
14074 wyneb001    0         6    0  0  0   0  0.467213   0.50         1976
14183 zaung001    0         33    3  2  0   1  0.385246   0.35         1995

      ...  SB    CS    BB    SO  IBB  HBP  SH  SF  GIDP  NL
108   ...  20  11.5  125  122    0    3  44   0    0    1
144   ...   8  12.0  154  381   12    5   4  19   34    1
151   ...   3   0.0   22  179    8   18   0   1    4    1
156   ...  23  18.0   87  223    4    9  35  16   27    1
169   ...   3   7.0  130  192    3    5  15  11   35    0
...      ...    ..    ...    ...    ...    ...    ..    ..    ...    ..
13914   ...  17  16.0  121   84    0    4  36   0    0    1
```

| | | | | | | | | | | | |
|-------|-----|----|------|-----|-----|----|----|----|----|-----|---|
| 13956 | ... | 5 | 11.0 | 277 | 278 | 14 | 7 | 5 | 12 | 52 | 1 |
| 14066 | ... | 1 | 0.0 | 5 | 41 | 2 | 1 | 2 | 2 | 3 | 1 |
| 14074 | ... | 10 | 13.0 | 626 | 428 | 41 | 17 | 58 | 36 | 119 | 0 |
| 14183 | ... | 23 | 19.0 | 479 | 544 | 30 | 29 | 14 | 31 | 87 | 1 |

[430 rows x 40 columns]

All fielders with a position of 'C' are in the catchers table, so we don't have to worry about leaving any by filtering with the following predicate.

```
[184]: fielders = fielders[~(fielders['retroID'].isin(catchers['retroID']) &
    ↳ fielders['pos_C'] == 1)]
```

```
[185]: fielders.shape
```

```
[185]: (5962, 40)
```

batting_build_tensor

April 30, 2020

Building the Batters Tensor

I've separated this from the model itself so that we could visualize and work through the data, but in the actual script this will just be a short few lines at the beginning of the model file.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Loading the Data

```
[2]: df = pd.read_csv('../core/output/batters.csv')
```

```
[3]: indexer = df.reset_index()[['index', 'retroID']].to_dict()['retroID']
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15293 entries, 0 to 15292
Data columns (total 38 columns):
#   Column      Non-Null Count  Dtype
---  -
0   retroID     15293 non-null  object
1   weight      15285 non-null  float64
2   height      15287 non-null  float64
3   debutYear   15293 non-null  int64
4   finalYear   15293 non-null  int64
5   pos_1B      15293 non-null  int64
6   pos_2B      15293 non-null  int64
7   pos_3B      15293 non-null  int64
8   pos_C       15293 non-null  int64
9   pos_OF      15293 non-null  int64
10  pos_P       15293 non-null  int64
11  pos_SS      15293 non-null  int64
12  bats_L      15293 non-null  int64
13  throws_L    15293 non-null  int64
14  G           15293 non-null  int64
15  AB          15293 non-null  int64
```

```

16 PA          15293 non-null int64
17 R           15293 non-null int64
18 H           15293 non-null int64
19 1B          15293 non-null int64
20 2B          15293 non-null int64
21 3B          15293 non-null int64
22 HR          15293 non-null int64
23 RBI         15293 non-null int64
24 SB          15293 non-null int64
25 CS          15293 non-null float64
26 BB          15293 non-null int64
27 SO          15293 non-null int64
28 IBB         15293 non-null int64
29 HBP         15293 non-null int64
30 SH          15293 non-null int64
31 SF          15293 non-null int64
32 GIDP        15293 non-null int64
33 NL          15293 non-null int64
34 wOBA        15293 non-null float64
35 wRC+        15293 non-null int64
36 WAR         15293 non-null float64
37 Batting     15293 non-null float64
dtypes: float64(6), int64(31), object(1)
memory usage: 4.4+ MB

```

```
[5]: y = df['Batting'].values
```

```
[6]: y
```

```
[6]: array([0.00035809, 0.350195 , 0.157131 , ..., 0.0900877 , 0.135118 ,
          0.0901954 ])
```

```
[7]: to_drop = ['retroID', 'debutYear', 'finalYear', 'Batting']
```

```
[8]: df.drop(columns=to_drop, inplace=True)
```

```
[9]: df
```

```
[9]:
```

| | weight | height | pos_1B | pos_2B | pos_3B | pos_C | pos_OF | pos_P | pos_SS | \ |
|-------|----------|--------|--------|--------|--------|-------|--------|-------|--------|---|
| 0 | 0.569672 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 1 | 0.426230 | 0.45 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 2 | 0.467213 | 0.60 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 4 | 0.442623 | 0.50 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0.590164 | 0.65 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 15289 | 0.434426 | 0.40 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |

| | | | | | | | | | | |
|-------|----------|------|---|---|---|---|---|---|---|---|
| 15290 | 0.487705 | 0.65 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| 15291 | 0.397541 | 0.45 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| 15292 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |

| | bats_L | ... | S0 | IBB | HBP | SH | SF | GIDP | NL | wOBA | wRC+ | WAR |
|-------|--------|-----|------|-----|-----|-----|-----|------|-----|-------|------|-------|
| 0 | 0 | ... | 2 | 0 | 0 | 1 | 0 | 0 | 1 | 0.000 | -100 | -0.1 |
| 1 | 0 | ... | 1383 | 293 | 32 | 21 | 121 | 328 | 1 | 0.403 | 153 | 136.3 |
| 2 | 0 | ... | 145 | 3 | 0 | 9 | 6 | 36 | 1 | 0.282 | 76 | -1.7 |
| 3 | 0 | ... | 3 | 0 | 0 | 0 | 0 | 0 | 1 | 0.000 | -100 | -0.1 |
| 4 | 1 | ... | 5 | 0 | 0 | 0 | 0 | 1 | 1 | 0.184 | 0 | -0.4 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 15288 | 0 | ... | 137 | 3 | 6 | 20 | 8 | 15 | 0 | 0.293 | 74 | -0.9 |
| 15289 | 1 | ... | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 0.225 | 37 | -0.2 |
| 15290 | 0 | ... | 39 | 0 | 0 | 16 | 0 | 3 | 1 | 0.179 | 0 | -0.3 |
| 15291 | 0 | ... | 50 | 1 | 2 | 18 | 0 | 8 | 1 | 0.254 | 52 | -2.2 |
| 15292 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.000 | 0 | 0.0 |


[15293 rows x 34 columns]

We now have a sort of proto-tensor, but maybe we can do some data manipulation to make the resulting model more efficient.

Observing Data Information

```
[10]: plt.figure(figsize=(17,12))
      ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
      bottom, top = ax.get_ylim()
      ax.set_ylim(bottom + 0.5, top - 0.5)
```

[10]: (34.0, 0.0)




We see a high correlation between G (games) and AB/PA (at-bats/plate appearances). It makes sense that we can drop the G column.

```
[11]: df.drop(columns=['G'], inplace=True)
```

```
[12]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
[12]: (33.0, 0.0)
```




There's obviously a high correlation in H (hits) and 1B/2B/3B/HR (singles, doubles, triples and home runs). We added 1B as a column to help with statistics but it's unnecessary now - the relationship between hits and types of hits will be preserved in the model.

```
[13]: df.drop(columns=['1B'], inplace=True)
```

```
[14]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
[14]: (32.0, 0.0)
```




We have total correlation between AB (at-bats) and PA (plate-appearances). This makes sense, because PA is just AB with some other situations added in. PA is more robust and is better related to overall output (since it includes sacrifices, hits-by-pitch and walks) so we'll keep PA and get rid of AB.

```
[15]: df.drop(columns=['AB'], inplace=True)
```

```
[16]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
[16]: (31.0, 0.0)
```




With the high correlation along the PA line, it may seem that we don't need them either. My reasoning for keeping them is a baseball-related one: we have most of the stats that make up a plate appearance, but we're missing the 'flier out' stat. This is an important one, as flyouts are a huge part of producing outs. Because of this, I'm going to keep PA as a stat.

One thing that I think we could drop is the RBI (runs batted in) stat. We don't see it appearing in many advanced stats, primarily wOBA and OPS+ with which we are concerned, and we intuitively see that it encompasses factors beyond the pure output of the hitter. It could be said that the RBI stat tracks a hitter's ability to hit "under pressure", but that's the kind of soft feature we're not going to consider. I think we get more important information from hits, doubles, triples, home runs and even runs than we do from RBIs. We also see extremely high correlation between RBI and R/H/2B/HR, which further supports the decision to remove RBI.

```
[21]: df.drop(columns=['RBI'], inplace=True)
```

```
[22]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
[22]: (30.0, 0.0)
```

We still see a hot spot around PA/R/H/2B, but these are all important stats that we're not going to remove.

One thing I think we SHOULD consider is that we have both wOBA and wRC+. There are a few issues with this:


- They are both stats that essentially only consider offensive output, so do we really need two?
- wRC+ incorporates wRAA, which is built from wOBA, so the information is a bit repeated.
- We want to minimize the model's use of secondary/advanced stats which are extrapolated from the primary stats we have. However, wRC+ and WAR both normalize to league trends and thus offer a nice aggregation of data that might not be intrinsically found by the model.

So I think we can go ahead and get rid of wOBA.

```
[23]: df.drop(columns=['wOBA'], inplace=True)
```

```
[24]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

[24]: (29.0, 0.0)



I said before that I don't want to drop R/H/2B/3B, so ignoring that area I think we now have a good looking model that's ready to run.

[25]: df

```
[25]:
```

| | weight | height | pos_1B | pos_2B | pos_3B | pos_C | pos_OF | pos_P | pos_SS | \ |
|-------|----------|--------|--------|--------|--------|-------|--------|-------|--------|---|
| 0 | 0.569672 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 1 | 0.426230 | 0.45 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 2 | 0.467213 | 0.60 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 4 | 0.442623 | 0.50 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0.590164 | 0.65 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 15289 | 0.434426 | 0.40 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
| 15290 | 0.487705 | 0.65 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 15291 | 0.397541 | 0.45 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | |
| 15292 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |

| | bats_L | ... | BB | SO | IBB | HBP | SH | SF | GIDP | NL | wRC+ | WAR |
|-------|--------|-----|------|------|-----|-----|-----|-----|------|-----|------|-------|
| 0 | 0 | ... | 0 | 2 | 0 | 0 | 1 | 0 | 0 | 1 | -100 | -0.1 |
| 1 | 0 | ... | 1402 | 1383 | 293 | 32 | 21 | 121 | 328 | 1 | 153 | 136.3 |
| 2 | 0 | ... | 86 | 145 | 3 | 0 | 9 | 6 | 36 | 1 | 76 | -1.7 |
| 3 | 0 | ... | 0 | 3 | 0 | 0 | 0 | 0 | 0 | 1 | -100 | -0.1 |
| 4 | 1 | ... | 4 | 5 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | -0.4 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 15288 | 0 | ... | 57 | 137 | 3 | 6 | 20 | 8 | 15 | 0 | 74 | -0.9 |
| 15289 | 1 | ... | 2 | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 37 | -0.2 |
| 15290 | 0 | ... | 9 | 39 | 0 | 0 | 16 | 0 | 3 | 1 | 0 | -0.3 |
| 15291 | 0 | ... | 34 | 50 | 1 | 2 | 18 | 0 | 8 | 1 | 52 | -2.2 |
| 15292 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 |

[15293 rows x 29 columns]

```
[28]: # We'll add the 'Batting' column back in to save our tensor
df.insert(loc=len(df.columns), column='Batting', value=y)
```

```
[29]: df
```

```
[29]:
```

| | weight | height | pos_1B | pos_2B | pos_3B | pos_C | pos_OF | pos_P | pos_SS | \ |
|-------|----------|--------|--------|--------|--------|-------|--------|-------|--------|---|
| 0 | 0.569672 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 1 | 0.426230 | 0.45 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 2 | 0.467213 | 0.60 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 4 | 0.442623 | 0.50 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 15288 | 0.590164 | 0.65 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 15289 | 0.434426 | 0.40 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
| 15290 | 0.487705 | 0.65 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 15291 | 0.397541 | 0.45 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | |
| 15292 | 0.467213 | 0.60 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |

| | bats_L | ... | SO | IBB | HBP | SH | SF | GIDP | NL | wRC+ | WAR | Batting |
|-------|--------|-----|------|-----|-----|-----|-----|------|-----|------|-------|----------|
| 0 | 0 | ... | 2 | 0 | 0 | 1 | 0 | 0 | 1 | -100 | -0.1 | 0.000358 |
| 1 | 0 | ... | 1383 | 293 | 32 | 21 | 121 | 328 | 1 | 153 | 136.3 | 0.350195 |
| 2 | 0 | ... | 145 | 3 | 0 | 9 | 6 | 36 | 1 | 76 | -1.7 | 0.157131 |
| 3 | 0 | ... | 3 | 0 | 0 | 0 | 0 | 0 | 1 | -100 | -0.1 | 0.000358 |
| 4 | 1 | ... | 5 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | -0.4 | 0.090001 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 15288 | 0 | ... | 137 | 3 | 6 | 20 | 8 | 15 | 0 | 74 | -0.9 | 0.156062 |
| 15289 | 1 | ... | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 37 | -0.2 | 0.123425 |
| 15290 | 0 | ... | 39 | 0 | 0 | 16 | 0 | 3 | 1 | 0 | -0.3 | 0.090088 |
| 15291 | 0 | ... | 50 | 1 | 2 | 18 | 0 | 8 | 1 | 52 | -2.2 | 0.135118 |
| 15292 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.0 | 0.090195 |

[15293 rows x 30 columns]

pitching_build_tensor

April 30, 2020

Building the Batters Tensor

I've separated this from the model itself so that we could visualize and work through the data, but in the actual script this will just be a short few lines at the beginning of the model file.

```
[14]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Loading the Data

```
[33]: df = pd.read_csv('../core/output/pitchers.csv')
```

```
[34]: indexer = df.reset_index()[['index', 'retroID']].to_dict()['retroID']
```

```
[35]: y = df['Pitching'].values
```

```
[36]: y
```

```
[36]: array([0.602913, 0.636924, 0.603736, ..., 0.612847, 0.608497, 0.611166])
```

```
[37]: df.columns
```

```
[37]: Index(['retroID', 'BAOpp', 'CG', 'SH0', 'IPouts', 'H', 'ER', 'HR', 'BB', 'SO',
        'IBB', 'WP', 'HBP', 'BK', 'BFP', 'GF', 'R', 'SH', 'SF', 'GIDP', 'K%',
        'IP', 'K/9', 'BB/9', 'HR/9', 'BABIP', 'LOB%', 'ERA', 'FIP', 'WAR',
        'Pitching'],
        dtype='object')
```

```
[38]: to_drop = ['retroID', 'Pitching']
```

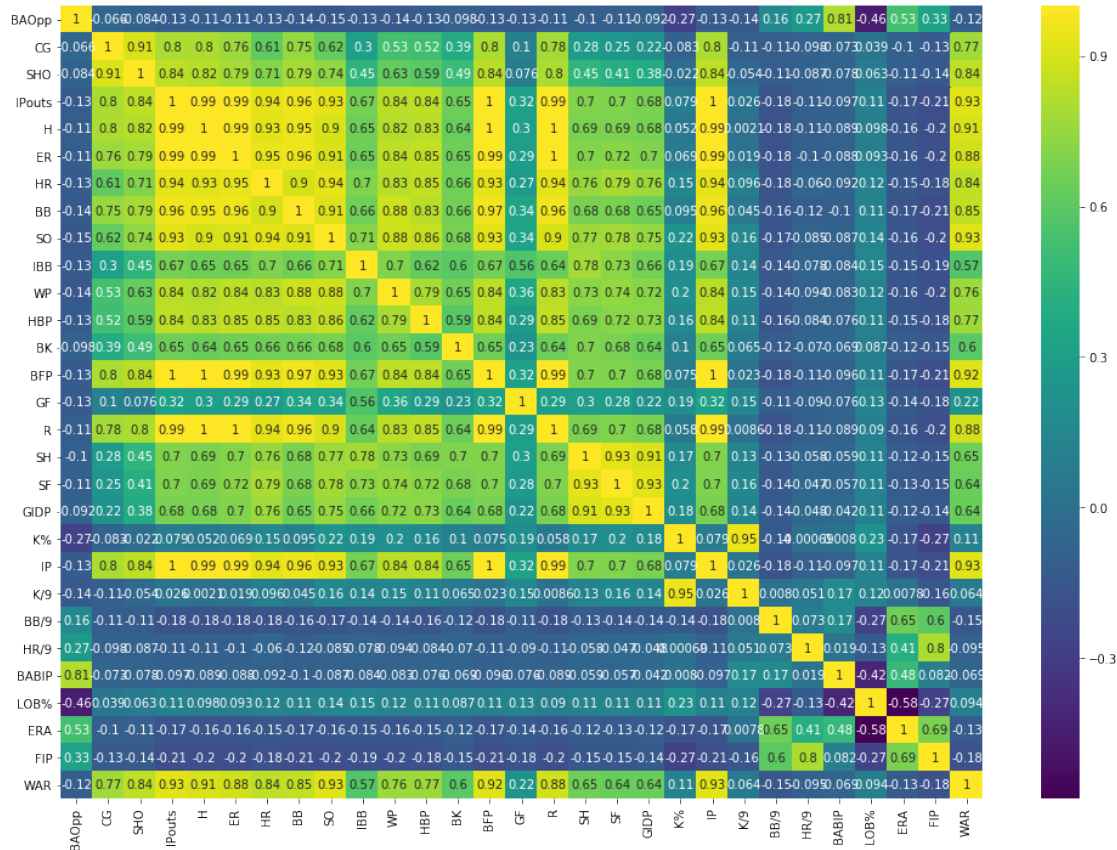
```
[39]: df = df.drop(columns=to_drop)
```

Observing Data Information

```
[40]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
```

```
ax.set_ylim(bottom + 0.5, top - 0.5)
```

[40]: (29.0, 0.0)

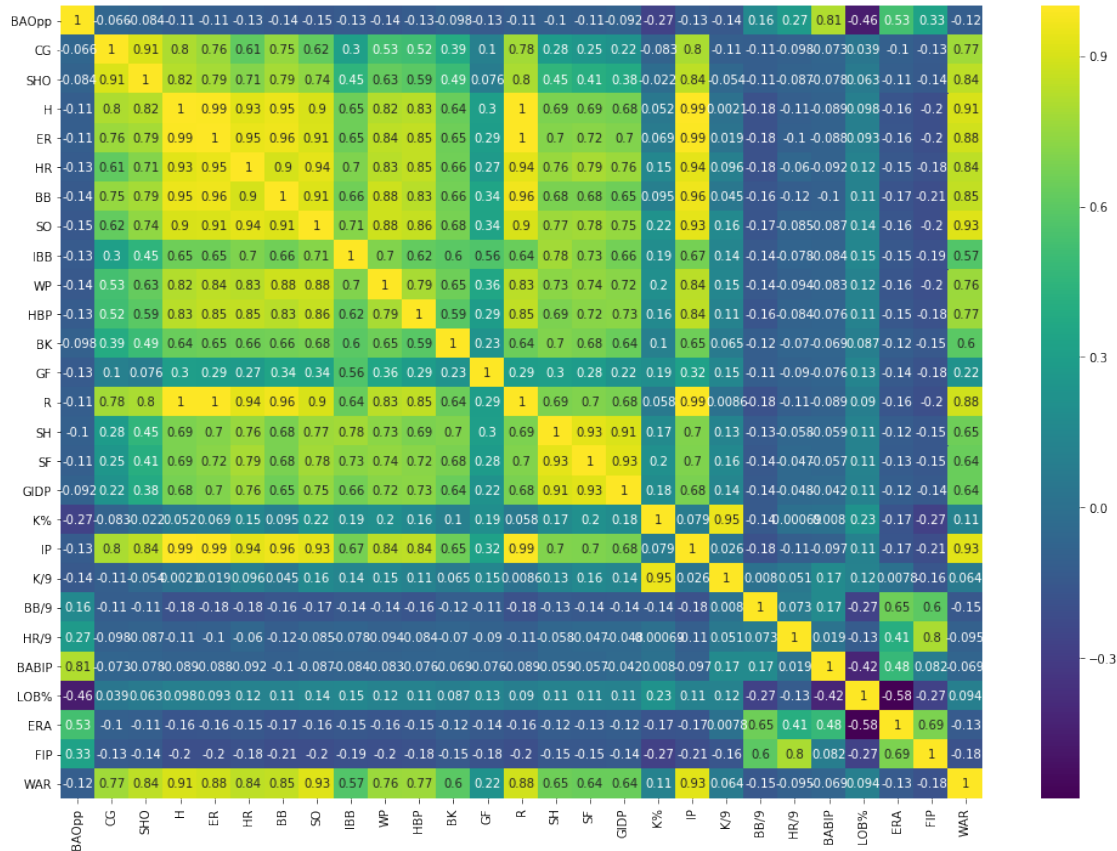


We see a lot of correlations with both IPouts and BFP, so we'll drop those.

```
[41]: df = df.drop(columns=['IPouts', 'BFP'])
```

```
[42]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

[42]: (27.0, 0.0)

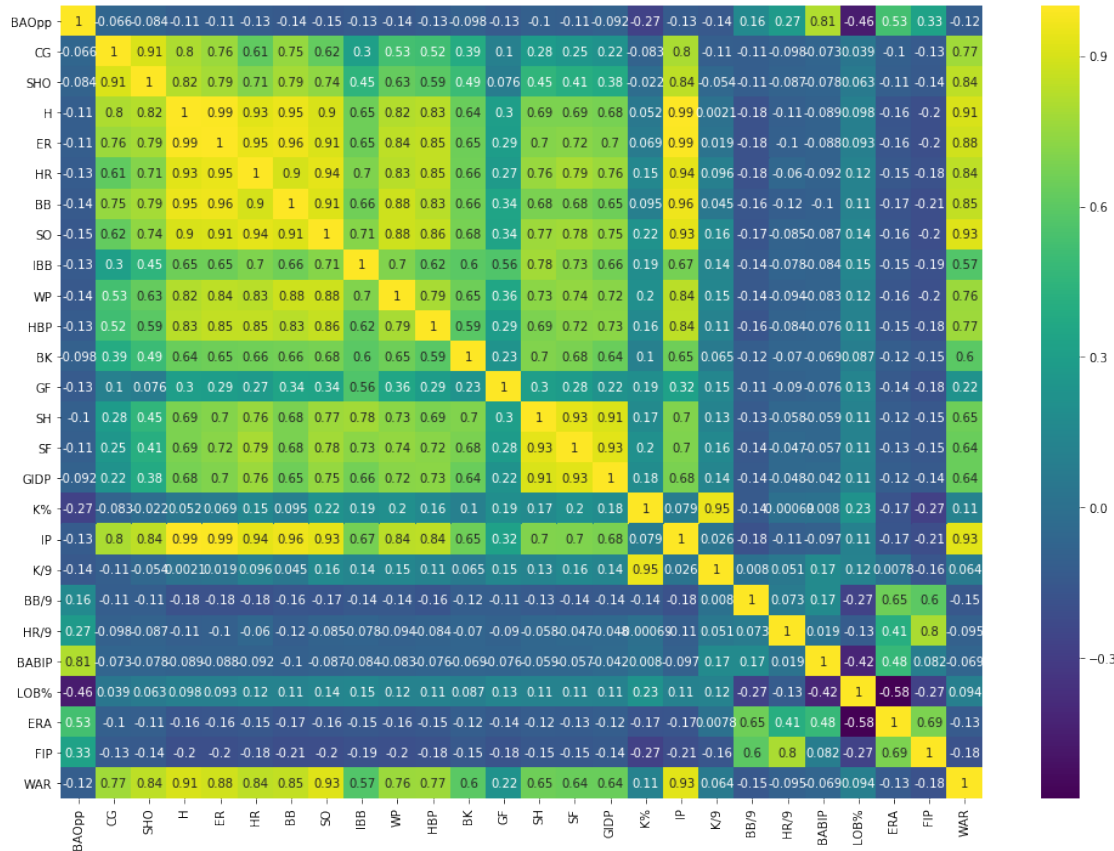


I think R (runs) is sufficiently covered by ER (earned runs) and HR (home runs), so I'll drop it too.

```
[43]: df = df.drop(columns=['R'])
```

```
[44]: plt.figure(figsize=(17,12))
ax = sns.heatmap(df.corr(), annot=True, cmap='viridis')
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
[44]: (26.0, 0.0)
```



I'm happy with this version of the tensor.

```
[45]: # We'll add the 'Pitching' column back in to save our tensor
df.insert(loc=len(df.columns), column='Pitching', value=y)
```

```
[46]: df
```

```
[46]:
```

| | BAOpp | CG | SHO | H | ER | HR | BB | SO | IBB | WP | ... | IP | K/9 | \ |
|------|--------|-----|-----|------|-----|-----|-----|-----|-----|-----|-----|----------|------|---|
| 0 | 0.2574 | 0 | 0 | 296 | 160 | 41 | 183 | 340 | 22 | 12 | ... | 0.062360 | 9.08 | |
| 1 | 0.2508 | 22 | 5 | 1085 | 468 | 89 | 457 | 641 | 45 | 22 | ... | 0.205233 | 5.20 | |
| 2 | 0.2447 | 0 | 0 | 309 | 135 | 42 | 116 | 280 | 10 | 10 | ... | 0.061102 | 7.62 | |
| 3 | 0.2786 | 37 | 5 | 1405 | 627 | 162 | 352 | 484 | 28 | 18 | ... | 0.237967 | 3.39 | |
| 4 | 0.2804 | 31 | 6 | 1779 | 791 | 154 | 620 | 888 | 30 | 53 | ... | 0.309765 | 4.77 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 8020 | 0.2700 | 30 | 5 | 956 | 366 | 54 | 301 | 207 | 0 | 8 | ... | 0.171925 | 2.00 | |
| 8021 | 0.2717 | 23 | 3 | 767 | 374 | 35 | 468 | 383 | 0 | 28 | ... | 0.145445 | 4.39 | |
| 8022 | 0.2286 | 0 | 0 | 169 | 71 | 18 | 114 | 210 | 11 | 16 | ... | 0.038711 | 9.01 | |
| 8023 | 0.2760 | 9 | 2 | 660 | 253 | 56 | 203 | 223 | 29 | 10 | ... | 0.118817 | 3.12 | |
| 8024 | 0.2183 | 0 | 0 | 57 | 22 | 3 | 34 | 80 | 5 | 2 | ... | 0.013360 | 9.91 | |

| | BB/9 | HR/9 | BABIP | LOB% | ERA | FIP | WAR | Pitching |
|------|------|------|-------|------|------|------|------|----------|
| 0 | 4.89 | 1.09 | 0.285 | 74.5 | 4.27 | 4.45 | 1.1 | 0.602913 |
| 1 | 3.71 | 0.72 | 0.282 | 73.4 | 3.80 | 3.85 | 11.7 | 0.636924 |
| 2 | 3.16 | 1.14 | 0.281 | 77.7 | 3.67 | 4.24 | 0.6 | 0.603736 |
| 3 | 2.46 | 1.13 | 0.278 | 69.3 | 4.39 | 4.46 | 10.2 | 0.628847 |
| 4 | 3.33 | 0.83 | 0.295 | 70.0 | 4.25 | 4.25 | 22.7 | 0.666725 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8020 | 2.91 | 0.52 | 0.267 | 70.7 | 3.54 | 3.80 | 9.3 | 0.630540 |
| 8021 | 5.36 | 0.40 | 0.283 | 69.0 | 4.28 | 3.96 | 3.3 | 0.610437 |
| 8022 | 4.89 | 0.77 | 0.267 | 78.7 | 3.00 | 3.94 | 2.7 | 0.612847 |
| 8023 | 2.84 | 0.78 | 0.270 | 73.2 | 3.54 | 3.93 | 1.9 | 0.608497 |
| 8024 | 4.21 | 0.37 | 0.293 | 79.1 | 2.72 | 3.22 | 1.1 | 0.611166 |

[8025 rows x 27 columns]

[]:

Game Log Processing

Taking Retrosheet game log CSV files and converting them into appropriate tensors for the predictive models

convert_gamelogs

May 7, 2020

Processing Game Logs

The information used here was obtained free of charge from and is copyrighted by Retrosheet. Interested parties may contact Retrosheet at “www.retrosheet.org”.

We’re going to use Retrosheet game logs as input data to our predictive model. The first thing we need to do is process them to fit our needs.

```
[135]: import pandas as pd
import os

[156]: df = pd.read_csv('../core/data/lahman/mlb_data/Teams.csv')
df = df[['teamID', 'franchID']]
team_dict = df.set_index('teamID').to_dict()['franchID']
team_dict['MLN'] = 'ATL'

def get_team(team):
    return team_dict[team] if team_dict[team] is not None else team
```

These are the columns of the Retrosheet game logs. This metadata was obtained here: <https://www.retrosheet.org/gamelogs/glfields.txt>

```
[137]: columns = [
    'date',
    'game_number',
    'day_of_week',
    'visit_team',
    'visit_league',
    'visit_game_number',
    'home_team',
    'home_league',
    'home_game_number',
    'visit_score',
    'home_score',
    'game_length_outs',
    'day_night',
    'completion_info',
    'forfeit_info',
```

```
'protest_info',
'park_id',
'attendance',
'time_minutes',
'visit_line_score',
'home_line_score',
'visit_ab',
'visit_h',
'visit_2b',
'visit_3b',
'visit_hr',
'visit_rbi',
'visit_sh',
'visit_sf',
'visit_hbp',
'visit_bb',
'visit_ibt',
'visit_k',
'visit_sb',
'visit_cs',
'visit_gidp',
'visit_ci',
'visit_lob',
'visit_pitchers_used',
'visit_individual_er',
'visit_team_er',
'visit_wp',
'visit_bk',
'visit_po',
'visit_assists',
'visit_e',
'visit_passed_balls',
'visit_double_plays',
'visit_triple_plays',
'home_ab',
'home_h',
'home_2b',
'home_3b',
'home_hr',
'home_rbi',
'home_sh',
'home_sf',
'home_hbp',
'home_bb',
'home_ibt',
'home_k',
'home_sb',
```



```
'home_cs',
'home_gidp',
'home_ci',
'home_lob',
'home_pitchers_used',
'home_individual_er',
'home_team_er',
'home_wp',
'home_bk',
'home_po',
'home_assists',
'home_e',
'home_passed_balls',
'home_double_plays',
'home_triple_plays',
'hp_ump_id',
'hp_ump_name',
'1b_ump_id',
'1b_ump_name',
'2b_ump_id',
'2b_ump_name',
'3b_ump_id',
'3b_ump_name',
'lf_ump_id',
'lf_ump_name',
'rf_ump_id',
'rf_ump_name',
'visit_manager_id',
'visit_manager_name',
'home_manager_id',
'home_manager_name',
'winning_pitcher_id',
'winning_pitcher_name',
'losing_pitcher_id',
'losing_pitcher_name',
'saving_pitcher_id',
'saving_pitcher_name',
'winning_rbi_batter_id',
'winning_rbi_batter_name',
'visit_sp_id',
'visit_sp_name',
'home_sp_id',
'home_sp_name',
'visit_player_1_id',
'visit_player_1_name',
'visit_player_1_pos',
'visit_player_2_id',
```

```
'visit_player_2_name',  
'visit_player_2_pos',  
'visit_player_3_id',  
'visit_player_3_name',  
'visit_player_3_pos',  
'visit_player_4_id',  
'visit_player_4_name',  
'visit_player_4_pos',  
'visit_player_5_id',  
'visit_player_5_name',  
'visit_player_5_pos',  
'visit_player_6_id',  
'visit_player_6_name',  
'visit_player_6_pos',  
'visit_player_7_id',  
'visit_player_7_name',  
'visit_player_7_pos',  
'visit_player_8_id',  
'visit_player_8_name',  
'visit_player_8_pos',  
'visit_player_9_id',  
'visit_player_9_name',  
'visit_player_9_pos',  
'home_player_1_id',  
'home_player_1_name',  
'home_player_1_pos',  
'home_player_2_id',  
'home_player_2_name',  
'home_player_2_pos',  
'home_player_3_id',  
'home_player_3_name',  
'home_player_3_pos',  
'home_player_4_id',  
'home_player_4_name',  
'home_player_4_pos',  
'home_player_5_id',  
'home_player_5_name',  
'home_player_5_pos',  
'home_player_6_id',  
'home_player_6_name',  
'home_player_6_pos',  
'home_player_7_id',  
'home_player_7_name',  
'home_player_7_pos',  
'home_player_8_id',  
'home_player_8_name',  
'home_player_8_pos',
```

```

        'home_player_9_id',
        'home_player_9_name',
        'home_player_9_pos',
        'additional_info',
        'acquisition_info'
    ]

```

The script is broken up here, then I later explore what I need to do to process the data. At the end I combine that all into one loop.

```

[12]: for year in range(1919, 2020):
        file_path = '../core/data/retrosheet/gamelogs/GL{}'.format(year)
        df = pd.read_csv(file_path + '.TXT', delimiter = ',', header = 0, names = _
        ↪columns)
        if os.path.exists(file_path + '.TXT'):
            os.remove(file_path + '.TXT')
        df.to_csv(file_path + '.csv')

```

```

[109]: df = pd.read_csv('../core/data/retrosheet/gamelogs/GL2015.csv')

```

We don't want every column, so we'll specify exactly which ones to use

```

[110]: df = df[[
        'date',
        'visit_team',
        'home_team',
        'visit_score',
        'home_score',
        'game_length_outs',
        'day_night',
        'park_id',
        'visit_manager_id',
        'home_manager_id',
        'visit_sp_id',
        'home_sp_id',
        'visit_player_1_id',
        'visit_player_2_id',
        'visit_player_3_id',
        'visit_player_4_id',
        'visit_player_5_id',
        'visit_player_6_id',
        'visit_player_7_id',
        'visit_player_8_id',
        'visit_player_9_id',
        'home_player_1_id',
        'home_player_2_id',
        'home_player_3_id',

```

```

        'home_player_4_id',
        'home_player_5_id',
        'home_player_6_id',
        'home_player_7_id',
        'home_player_8_id',
        'home_player_9_id'
    ]]

```

```
[111]: df
```

```

[111]:
      date visit_team home_team  visit_score  home_score \
0    20150406      MIN      DET            0            4
1    20150406      CLE      HOU            0            2
2    20150406      CHA      KCA            1           10
3    20150406      TOR      NYA            6            1
4    20150406      TEX      OAK            0            8
...      ...      ...      ...      ...      ...
2423  20151004      CHN      MIL            3            1
2424  20151004      WAS      NYN            0            1
2425  20151004      MIA      PHI            2            7
2426  20151004      CIN      PIT            0            4
2427  20151004      COL      SFN            7            3

      game_length_outs day_night park_id visit_manager_id home_manager_id \
0                    51         D  DET05      molip001      ausmb001
1                    51         N  HOU03      frant001      hinca001
2                    51         D  KAN06      ventr001      yoste001
3                    54         D  NYC21      gibbj001      giraj001
4                    51         N  OAK01      banij001      melvb001
...      ...      ...      ...      ...      ...
2423          54         D  MIL06      maddj801      councl001
2424          51         D  NYC20      willm003      collt801
2425          51         D  PHI13      jennd801      mackp101
2426          51         D  PIT08      pricb801      hurdc001
2427          54         D  SF003      weisw001      bochb002

      ... visit_player_9_id home_player_1_id home_player_2_id \
0      ...      schaj002      davir003      kinsi001
1      ...      ramij003      altuj001      sprig001
2      ...      johnm006      escoa003      mousm001
3      ...      travd001      ellsj001      gardb001
4      ...      odorr001      gentc001      fulds001
...      ...      ...      ...      ...
2423      ...      hared001      genns001      petes002
2424      ...      roart001      granc001      wrigd002
2425      ...      conla001      galvf001      altha001
2426      ...      smitj004      polag001      harrj002

```

| | | | | | |
|------|------------------|------------------|------------------|------------------|---|
| 2427 | ... | bergc001 | pagaa001 | tomlk001 | |
| | home_player_3_id | home_player_4_id | home_player_5_id | home_player_6_id | \ |
| 0 | cabrm001 | martv001 | martj006 | cespy001 | |
| 1 | valbl001 | gatte001 | cartc002 | castj006 | |
| 2 | cainl001 | hosme001 | morak001 | gora001 | |
| 3 | beltc001 | teixm001 | mccab002 | headc001 | |
| 4 | zobrb001 | butlb003 | davii001 | lawrb002 | |
| ... | ... | ... | ... | ... | |
| 2423 | linda001 | davik003 | santd002 | pereh001 | |
| 2424 | murpd006 | cespy001 | dudal001 | darnt001 | |
| 2425 | franm004 | ruf-d001 | fran-j004 | blana001 | |
| 2426 | mccua001 | walkn001 | marts002 | alvap001 | |
| 2427 | duffm002 | poseb001 | parkj002 | willm008 | |

| | | | |
|------|------------------|------------------|------------------|
| | home_player_7_id | home_player_8_id | home_player_9_id |
| 0 | castn001 | avila001 | iglej001 |
| 1 | lowrj001 | rasmc001 | marij002 |
| 2 | rios002 | peres002 | infao001 |
| 3 | rodra001 | drewn001 | gregd001 |
| 4 | vogts001 | semim001 | sogae001 |
| ... | ... | ... | ... |
| 2423 | seguj002 | maldm001 | lopej004 |
| 2424 | confm001 | tejar001 | degrj001 |
| 2425 | krate001 | ruppc001 | buchd001 |
| 2426 | cervf001 | mercj002 | happj001 |
| 2427 | noonn001 | willj005 | cainm001 |

[2428 rows x 33 columns]

```
[112]: df['date'] = df['date'].astype(str)
```

'date' isn't very useful, so we'll export it to three separate columns.

```
[113]: df['year'] = df['date'].str[0:4].astype(int)
df['month'] = df['date'].str[4:6].astype(int)
df['day'] = df['date'].str[6:8].astype(int)
```

```
[114]: df
```

```
[114]:
```

| | | | | | | |
|-----|----------|------------|-----------|-------------|------------|---|
| | date | visit_team | home_team | visit_score | home_score | \ |
| 0 | 20150406 | MIN | DET | 0 | 4 | |
| 1 | 20150406 | CLE | HOU | 0 | 2 | |
| 2 | 20150406 | CHA | KCA | 1 | 10 | |
| 3 | 20150406 | TOR | NYA | 6 | 1 | |
| 4 | 20150406 | TEX | OAK | 0 | 8 | |
| ... | ... | ... | ... | ... | ... | |

| | | | | | |
|------|----------|-----|-----|---|---|
| 2423 | 20151004 | CHN | MIL | 3 | 1 |
| 2424 | 20151004 | WAS | NYN | 0 | 1 |
| 2425 | 20151004 | MIA | PHI | 2 | 7 |
| 2426 | 20151004 | CIN | PIT | 0 | 4 |
| 2427 | 20151004 | COL | SFN | 7 | 3 |

| | game_length_outs | day_night | park_id | visit_manager_id | home_manager_id | \ |
|------|------------------|-----------|---------|------------------|-----------------|---|
| 0 | 51 | D | DET05 | molip001 | ausmb001 | |
| 1 | 51 | N | HOU03 | frant001 | hinca001 | |
| 2 | 51 | D | KAN06 | ventr001 | yoste001 | |
| 3 | 54 | D | NYC21 | gibbj001 | giraj001 | |
| 4 | 51 | N | OAK01 | banij001 | melvb001 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | 54 | D | MIL06 | maddj801 | counc001 | |
| 2424 | 51 | D | NYC20 | willm003 | collt801 | |
| 2425 | 51 | D | PHI13 | jennd801 | mackp101 | |
| 2426 | 51 | D | PIT08 | pricb801 | hurdc001 | |
| 2427 | 54 | D | SF003 | weisw001 | bochb002 | |

| | ... | home_player_3_id | home_player_4_id | home_player_5_id | home_player_6_id | \ |
|------|-----|------------------|------------------|------------------|------------------|---|
| 0 | ... | cabrm001 | martv001 | martj006 | cespy001 | |
| 1 | ... | valbl001 | gatte001 | cartc002 | castj006 | |
| 2 | ... | cainl001 | hosme001 | morak001 | gorda001 | |
| 3 | ... | beltc001 | teixm001 | mccab002 | headc001 | |
| 4 | ... | zobrb001 | butlb003 | davii001 | lawrb002 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | ... | linda001 | davik003 | santd002 | pereh001 | |
| 2424 | ... | murpd006 | cespy001 | dudal001 | darnt001 | |
| 2425 | ... | franm004 | ruf-d001 | franj004 | blana001 | |
| 2426 | ... | mccua001 | walkn001 | marts002 | alvap001 | |
| 2427 | ... | duffm002 | poseb001 | parkj002 | willm008 | |

| | home_player_7_id | home_player_8_id | home_player_9_id | year | month | day |
|------|------------------|------------------|------------------|------|-------|-----|
| 0 | castn001 | avila001 | iglej001 | 2015 | 4 | 6 |
| 1 | lowrj001 | rasmc001 | marij002 | 2015 | 4 | 6 |
| 2 | riosa002 | peres002 | infao001 | 2015 | 4 | 6 |
| 3 | rodra001 | drews001 | gregd001 | 2015 | 4 | 6 |
| 4 | vogts001 | semim001 | sogae001 | 2015 | 4 | 6 |
| ... | ... | ... | ... | ... | ... | ... |
| 2423 | seguj002 | maldm001 | lopej004 | 2015 | 10 | 4 |
| 2424 | confm001 | tejar001 | degrj001 | 2015 | 10 | 4 |
| 2425 | krate001 | ruppc001 | buchd001 | 2015 | 10 | 4 |
| 2426 | cervf001 | mercj002 | happj001 | 2015 | 10 | 4 |
| 2427 | noonn001 | willj005 | cainm001 | 2015 | 10 | 4 |

[2428 rows x 36 columns]

We aren't going to use every column in the final model, but we want to make sure that the ones we will are in the proper format.

```
[115]: night_game = pd.get_dummies(df['day_night'], drop_first=True)
```

```
[116]: night_game
```

```
[116]:
```

| | N |
|------|----|
| 0 | 0 |
| 1 | 1 |
| 2 | 0 |
| 3 | 0 |
| 4 | 1 |
| ... | .. |
| 2423 | 0 |
| 2424 | 0 |
| 2425 | 0 |
| 2426 | 0 |
| 2427 | 0 |

[2428 rows x 1 columns]

```
[117]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2428 entries, 0 to 2427
Data columns (total 36 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  2428 non-null  object
1   visit_team            2428 non-null  object
2   home_team             2428 non-null  object
3   visit_score           2428 non-null  int64
4   home_score            2428 non-null  int64
5   game_length_outs      2428 non-null  int64
6   day_night             2428 non-null  object
7   park_id               2428 non-null  object
8   visit_manager_id      2428 non-null  object
9   home_manager_id       2428 non-null  object
10  winning_pitcher_id    2428 non-null  object
11  losing_pitcher_id     2428 non-null  object
12  saving_pitcher_id     1291 non-null  object
13  visit_sp_id           2428 non-null  object
14  home_sp_id            2428 non-null  object
15  visit_player_1_id     2428 non-null  object
16  visit_player_2_id     2428 non-null  object
17  visit_player_3_id     2428 non-null  object
18  visit_player_4_id     2428 non-null  object
```

```

19 visit_player_5_id    2428 non-null    object
20 visit_player_6_id    2428 non-null    object
21 visit_player_7_id    2428 non-null    object
22 visit_player_8_id    2428 non-null    object
23 visit_player_9_id    2428 non-null    object
24 home_player_1_id     2428 non-null    object
25 home_player_2_id     2428 non-null    object
26 home_player_3_id     2428 non-null    object
27 home_player_4_id     2428 non-null    object
28 home_player_5_id     2428 non-null    object
29 home_player_6_id     2428 non-null    object
30 home_player_7_id     2428 non-null    object
31 home_player_8_id     2428 non-null    object
32 home_player_9_id     2428 non-null    object
33 year                 2428 non-null    int64
34 month                 2428 non-null    int64
35 day                  2428 non-null    int64

```

dtypes: int64(6), object(30)

memory usage: 683.0+ KB

```
[118]: df.insert(loc=6, column='night_game', value=night_game)
```

```
[119]: df
```

```

[119]:
      date visit_team home_team visit_score home_score \
0    20150406      MIN      DET           0           4
1    20150406      CLE      HOU           0           2
2    20150406      CHA      KCA           1          10
3    20150406      TOR      NYA           6           1
4    20150406      TEX      OAK           0           8
...      ...      ...      ...      ...      ...
2423  20151004      CHN      MIL           3           1
2424  20151004      WAS      NYN           0           1
2425  20151004      MIA      PHI           2           7
2426  20151004      CIN      PIT           0           4
2427  20151004      COL      SFN           7           3

      game_length_outs  night_game day_night park_id visit_manager_id ... \
0                   51           0         D   DET05      molip001 ...
1                   51           1         N   HOU03      frant001 ...
2                   51           0         D   KAN06      ventr001 ...
3                   54           0         D   NYC21      gibbj001 ...
4                   51           1         N   OAK01      banij001 ...
...      ...      ...      ...      ...      ...
2423                54           0         D   MIL06      maddj801 ...
2424                51           0         D   NYC20      willm003 ...
2425                51           0         D   PHI13      jennd801 ...

```


| | | | | | | |
|------|----|---|---|-------|----------|-----|
| 2426 | 51 | 0 | D | PIT08 | pricb801 | ... |
| 2427 | 54 | 0 | D | SF003 | weisw001 | ... |

| | home_player_3_id | home_player_4_id | home_player_5_id | home_player_6_id | \ |
|------|------------------|------------------|------------------|------------------|---|
| 0 | cabrm001 | martv001 | martj006 | cespy001 | |
| 1 | valbl001 | gatte001 | cartc002 | castj006 | |
| 2 | cainl001 | hosme001 | morak001 | gorda001 | |
| 3 | beltc001 | teixm001 | mccab002 | headc001 | |
| 4 | zobrb001 | butlb003 | davii001 | lawrb002 | |
| ... | ... | ... | ... | ... | |
| 2423 | linda001 | davik003 | santtd002 | pereh001 | |
| 2424 | murpd006 | cespy001 | dudal001 | darnt001 | |
| 2425 | franm004 | ruf-d001 | franjd004 | blana001 | |
| 2426 | mccua001 | walkn001 | marts002 | alvap001 | |
| 2427 | duffm002 | poseb001 | parkj002 | willm008 | |

| | home_player_7_id | home_player_8_id | home_player_9_id | year | month | day |
|------|------------------|------------------|------------------|------|-------|-----|
| 0 | castn001 | avila001 | iglej001 | 2015 | 4 | 6 |
| 1 | lowrj001 | rasmc001 | marij002 | 2015 | 4 | 6 |
| 2 | riosaa002 | peres002 | infao001 | 2015 | 4 | 6 |
| 3 | rodra001 | drews001 | gregd001 | 2015 | 4 | 6 |
| 4 | vogts001 | semim001 | sogae001 | 2015 | 4 | 6 |
| ... | ... | ... | ... | ... | ... | ... |
| 2423 | seguj002 | maldm001 | lopej004 | 2015 | 10 | 4 |
| 2424 | confm001 | tejar001 | degrj001 | 2015 | 10 | 4 |
| 2425 | krate001 | ruppc001 | buchd001 | 2015 | 10 | 4 |
| 2426 | cervf001 | mercj002 | happj001 | 2015 | 10 | 4 |
| 2427 | noonn001 | willj005 | cainm001 | 2015 | 10 | 4 |

[2428 rows x 37 columns]

```
[120]: to_drop = ['date', 'day_night']
```

```
[121]: df = df.drop(columns=to_drop)
```

```
[122]: df
```

```
[122]:
```

| | visit_team | home_team | visit_score | home_score | game_length_outs | \ |
|------|------------|-----------|-------------|------------|------------------|---|
| 0 | MIN | DET | 0 | 4 | 51 | |
| 1 | CLE | HOU | 0 | 2 | 51 | |
| 2 | CHA | KCA | 1 | 10 | 51 | |
| 3 | TOR | NYA | 6 | 1 | 54 | |
| 4 | TEX | OAK | 0 | 8 | 51 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | CHN | MIL | 3 | 1 | 54 | |
| 2424 | WAS | NYN | 0 | 1 | 51 | |
| 2425 | MIA | PHI | 2 | 7 | 51 | |

| | | | | | |
|------|-----|-----|---|---|----|
| 2426 | CIN | PIT | 0 | 4 | 51 |
| 2427 | COL | SFN | 7 | 3 | 54 |

| | night_game | park_id | visit_manager_id | home_manager_id | winning_pitcher_id | \ |
|------|------------|---------|------------------|-----------------|--------------------|---|
| 0 | 0 | DET05 | molip001 | ausmb001 | pricd001 | |
| 1 | 1 | HOU03 | frant001 | hinca001 | keucd001 | |
| 2 | 0 | KAN06 | ventr001 | yoste001 | venty001 | |
| 3 | 0 | NYC21 | gibbj001 | giraj001 | hutcd001 | |
| 4 | 1 | OAK01 | banij001 | melvb001 | grays001 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | 0 | MIL06 | maddj801 | counc001 | hared001 | |
| 2424 | 0 | NYC20 | willm003 | collt801 | clipt001 | |
| 2425 | 0 | PHI13 | jenn801 | mackp101 | garcl005 | |
| 2426 | 0 | PIT08 | pricb801 | hurdc001 | happj001 | |
| 2427 | 0 | SF003 | weisw001 | bochb002 | brotr001 | |

| | ... | home_player_3_id | home_player_4_id | home_player_5_id | home_player_6_id | \ |
|------|-----|------------------|------------------|------------------|------------------|---|
| 0 | ... | cabrm001 | martv001 | martj006 | cespy001 | |
| 1 | ... | valbl001 | gatte001 | cartc002 | castj006 | |
| 2 | ... | cainl001 | hosme001 | morak001 | gorda001 | |
| 3 | ... | beltc001 | teixm001 | mccab002 | headc001 | |
| 4 | ... | zobrb001 | butlb003 | davii001 | lawrb002 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | ... | linda001 | davik003 | santd002 | pereh001 | |
| 2424 | ... | murpd006 | cespy001 | dudal001 | darnt001 | |
| 2425 | ... | franm004 | ruf-d001 | fran004 | blana001 | |
| 2426 | ... | mccua001 | walkn001 | marts002 | alvap001 | |
| 2427 | ... | duffm002 | poseb001 | parkj002 | willm008 | |

| | home_player_7_id | home_player_8_id | home_player_9_id | year | month | day |
|------|------------------|------------------|------------------|------|-------|-----|
| 0 | castn001 | avila001 | iglej001 | 2015 | 4 | 6 |
| 1 | lowrj001 | rasmc001 | marij002 | 2015 | 4 | 6 |
| 2 | riosa002 | peres002 | infao001 | 2015 | 4 | 6 |
| 3 | rodra001 | drews001 | gregd001 | 2015 | 4 | 6 |
| 4 | vogts001 | semim001 | sogae001 | 2015 | 4 | 6 |
| ... | ... | ... | ... | ... | ... | ... |
| 2423 | seguj002 | maldm001 | lopej004 | 2015 | 10 | 4 |
| 2424 | confm001 | tejar001 | degrj001 | 2015 | 10 | 4 |
| 2425 | krate001 | ruppc001 | buchd001 | 2015 | 10 | 4 |
| 2426 | cervf001 | mercj002 | happj001 | 2015 | 10 | 4 |
| 2427 | noonn001 | willj005 | cainm001 | 2015 | 10 | 4 |

[2428 rows x 35 columns]

```
[123]: df['visit_team'] = df['visit_team'].apply(get_team)
```

```
[124]: df['home_team'] = df['home_team'].apply(get_team)
```

Early games only took place during the day, so we need to handle the effects of using one-hot encoding when dropping first with those.

```
[147]: file_path = '../core/data/retrosheet/gamelogs/GL{}'.format(1919)
df = pd.read_csv(file_path + '.TXT', delimiter = ',', header = 0, names = _
    ↪columns)
```

```
[148]: df['day_night'].nunique()
```

```
[148]: 1
```

```
[149]: pd.get_dummies(df['day_night'])
```

```
[149]:      D
0      1
1      1
2      1
3      1
4      1
...   ..
1112   1
1113   1
1114   1
1115   1
1116   1

[1117 rows x 1 columns]
```

Final Script

```
[172]: for year in range(1919, 2020):
        file_path = '../core/data/retrosheet/gamelogs/GL{}'.format(year)
        df = pd.read_csv(file_path + '.TXT', delimiter = ',', header = 0, names = _
            ↪columns)
        df = df[[
            'date',
            'visit_team',
            'home_team',
            'visit_score',
            'home_score',
            'game_length_outs',
            'day_night',
            'park_id',
            'visit_manager_id',
            'home_manager_id',
            'winning_pitcher_id',
            'losing_pitcher_id',
            'saving_pitcher_id',
```

```

        'visit_sp_id',
        'home_sp_id',
        'visit_player_1_id',
        'visit_player_2_id',
        'visit_player_3_id',
        'visit_player_4_id',
        'visit_player_5_id',
        'visit_player_6_id',
        'visit_player_7_id',
        'visit_player_8_id',
        'visit_player_9_id',
        'home_player_1_id',
        'home_player_2_id',
        'home_player_3_id',
        'home_player_4_id',
        'home_player_5_id',
        'home_player_6_id',
        'home_player_7_id',
        'home_player_8_id',
        'home_player_9_id',
    ]]
    df['date'] = df['date'].astype(str)
    df['year'] = df['date'].str[0:4].astype(int)
    df['month'] = df['date'].str[4:6].astype(int)
    df['day'] = df['date'].str[6:8].astype(int)
    night_game = pd.get_dummies(df['day_night'], drop_first=(df['day_night'].
    ↪nunique() > 1))
    df.insert(loc=6, column='night_game', value=night_game)
    df = df.drop(columns=['date', 'day_night'])
    df['visit_team'] = df['visit_team'].apply(get_team)
    df['home_team'] = df['home_team'].apply(get_team)
    if os.path.exists(file_path + '.TXT'):
        os.remove(file_path + '.TXT')
    df.to_csv(file_path + '.csv', index=False)

```

```
[170]: df = pd.read_csv('../core/data/retrosheet/gamelogs/GL2015.csv')
```

```
[173]: df
```

```
[173]:
```

| | visit_team | home_team | visit_score | home_score | game_length_outs | \ |
|------|------------|-----------|-------------|------------|------------------|---|
| 0 | MIN | DET | 0 | 4 | 51 | |
| 1 | CLE | HOU | 0 | 2 | 51 | |
| 2 | CHW | KCR | 1 | 10 | 51 | |
| 3 | TOR | NYN | 6 | 1 | 54 | |
| 4 | TEX | OAK | 0 | 8 | 51 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | CHC | MIL | 3 | 1 | 54 | |

| | | | | | |
|------|-----|-----|---|---|----|
| 2424 | WSN | NYM | 0 | 1 | 51 |
| 2425 | FLA | PHI | 2 | 7 | 51 |
| 2426 | CIN | PIT | 0 | 4 | 51 |
| 2427 | COL | SFG | 7 | 3 | 54 |

| | night_game | park_id | visit_manager_id | home_manager_id | winning_pitcher_id | \ |
|------|------------|---------|------------------|-----------------|--------------------|---|
| 0 | 0 | DET05 | molip001 | ausmb001 | pricd001 | |
| 1 | 1 | HOU03 | frant001 | hinca001 | keucd001 | |
| 2 | 0 | KAN06 | ventr001 | yoste001 | venty001 | |
| 3 | 0 | NYC21 | gibbj001 | giraj001 | hutcd001 | |
| 4 | 1 | OAK01 | banij001 | melvb001 | grays001 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | 0 | MIL06 | maddj801 | counc001 | hared001 | |
| 2424 | 0 | NYC20 | willm003 | collt801 | clipt001 | |
| 2425 | 0 | PHI13 | jennd801 | mackp101 | garcl005 | |
| 2426 | 0 | PIT08 | pricb801 | hurdc001 | happj001 | |
| 2427 | 0 | SF003 | weisw001 | bochb002 | brotr001 | |

| | ... | home_player_3_id | home_player_4_id | home_player_5_id | home_player_6_id | \ |
|------|-----|------------------|------------------|------------------|------------------|---|
| 0 | ... | cabrm001 | martv001 | martj006 | cespy001 | |
| 1 | ... | valbl001 | gatte001 | cartc002 | castj006 | |
| 2 | ... | cainl001 | hosme001 | morak001 | gorda001 | |
| 3 | ... | beltc001 | teixm001 | mccab002 | headc001 | |
| 4 | ... | zobrb001 | butlb003 | davii001 | lawrb002 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | ... | linda001 | davik003 | santd002 | pereh001 | |
| 2424 | ... | murpd006 | cespy001 | dudal001 | darnt001 | |
| 2425 | ... | franm004 | ruf-d001 | franj004 | blana001 | |
| 2426 | ... | mccua001 | walkn001 | marts002 | alvap001 | |
| 2427 | ... | duffm002 | poseb001 | parkj002 | willm008 | |

| | home_player_7_id | home_player_8_id | home_player_9_id | year | month | day |
|------|------------------|------------------|------------------|------|-------|-----|
| 0 | castn001 | avila001 | iglej001 | 2015 | 4 | 6 |
| 1 | lowrj001 | rasmc001 | marij002 | 2015 | 4 | 6 |
| 2 | rios002 | peres002 | infao001 | 2015 | 4 | 6 |
| 3 | rodra001 | drews001 | gregd001 | 2015 | 4 | 6 |
| 4 | vogts001 | semim001 | sogae001 | 2015 | 4 | 6 |
| ... | ... | ... | ... | ... | ... | ... |
| 2423 | seguj002 | maldm001 | lopej004 | 2015 | 10 | 4 |
| 2424 | confm001 | tejar001 | degrj001 | 2015 | 10 | 4 |
| 2425 | krate001 | ruppc001 | buchd001 | 2015 | 10 | 4 |
| 2426 | cervf001 | mercj002 | happj001 | 2015 | 10 | 4 |
| 2427 | noonn001 | willj005 | cainm001 | 2015 | 10 | 4 |

[2428 rows x 35 columns]

When we do the actual script, we don't want the column names hardcoded into it. So I've pasted

those to .csv files but I need to process them a bit.

```
[252]: gla = pd.read_csv('../core/data/retrosheet/rs_gl_cols_all.csv', header=None)
      gl = pd.read_csv('../core/data/retrosheet/rs_gl_cols.csv', header=None)
```

```
[254]: gl
```

```
[254]:
0
0          'date',
1      'visit_team',
2      'home_team',
3      'visit_score',
4      'home_score',
5      'game_length_outs',
6      'day_night',
7      'park_id',
8      'visit_manager_id',
9      'home_manager_id',
10     'winning_pitcher_id',
11     'losing_pitcher_id',
12     'saving_pitcher_id',
13     'visit_sp_id',
14     'home_sp_id',
15     'visit_player_1_id',
16     'visit_player_2_id',
17     'visit_player_3_id',
18     'visit_player_4_id',
19     'visit_player_5_id',
20     'visit_player_6_id',
21     'visit_player_7_id',
22     'visit_player_8_id',
23     'visit_player_9_id',
24     'home_player_1_id',
25     'home_player_2_id',
26     'home_player_3_id',
27     'home_player_4_id',
28     'home_player_5_id',
29     'home_player_6_id',
30     'home_player_7_id',
31     'home_player_8_id',
32     'home_player_9_id'
```

We need to get rid of whitespace, commas and quotation marks.

```
[189]: gla.iloc[156][0].replace(',', '')
```

```
[189]: "    'home_player_9_id'"
```

```

[190]: from functools import reduce

[241]: def trim_cell(cell):
        replacements = {' ': '', '"': "'", ',': ''}
        string = cell[0]
        return reduce(lambda a, kv: a.replace(*kv), replacements.items(), string)

[200]: print(trim_cell(gla.iloc[156]))

home_player_9_id

[255]: gla = gla.apply(trim_cell, axis=1)

[256]: type(gla)

pandas.core.series.Series

[257]: gl = gl.apply(trim_cell, axis=1)

[258]: gla.to_csv('../core/data/retrosheet/rs_gl_cols_all.csv', header=None)
gl.to_csv('../core/data/retrosheet/rs_gl_cols.csv', header=None)

[259]: gla = pd.read_csv('../core/data/retrosheet/rs_gl_cols_all.csv', header=None)
gl = pd.read_csv('../core/data/retrosheet/rs_gl_cols.csv', header=None)

[260]: gl

[260]:      0      1
0  0      date
1  1  visit_team
2  2  home_team
3  3  visit_score
4  4  home_score
5  5  game_length_outs
6  6  day_night
7  7  park_id
8  8  visit_manager_id
9  9  home_manager_id
10 10 winning_pitcher_id
11 11 losing_pitcher_id
12 12 saving_pitcher_id
13 13 visit_sp_id
14 14 home_sp_id
15 15 visit_player_1_id
16 16 visit_player_2_id
17 17 visit_player_3_id
18 18 visit_player_4_id

```

```

19 19  visit_player_5_id
20 20  visit_player_6_id
21 21  visit_player_7_id
22 22  visit_player_8_id
23 23  visit_player_9_id
24 24  home_player_1_id
25 25  home_player_2_id
26 26  home_player_3_id
27 27  home_player_4_id
28 28  home_player_5_id
29 29  home_player_6_id
30 30  home_player_7_id
31 31  home_player_8_id
32 32  home_player_9_id

```

```
[261]: gla[1].tolist()
```

```

[261]: ['date',
        'game_number',
        'day_of_week',
        'visit_team',
        'visit_league',
        'visit_game_number',
        'home_team',
        'home_league',
        'home_game_number',
        'visit_score',
        'home_score',
        'game_length_outs',
        'day_night',
        'completion_info',
        'forfeit_info',
        'protest_info',
        'park_id',
        'attendance',
        'time_minutes',
        'visit_line_score',
        'home_line_score',
        'visit_ab',
        'visit_h',
        'visit_2b',
        'visit_3b',
        'visit_hr',
        'visit_rbi',
        'visit_sh',
        'visit_sf',
        'visit_hbp',

```


'visit_bb',
'visit_ibt',
'visit_k',
'visit_sb',
'visit_cs',
'visit_gidp',
'visit_ci',
'visit_lob',
'visit_pitchers_used',
'visit_individual_er',
'visit_team_er',
'visit_wp',
'visit_bk',
'visit_po',
'visit_assists',
'visit_e',
'visit_passed_balls',
'visit_double_plays',
'visit_triple_plays',
'home_ab',
'home_h',
'home_2b',
'home_3b',
'home_hr',
'home_rbi',
'home_sh',
'home_sf',
'home_hbp',
'home_bb',
'home_ibt',
'home_k',
'home_sb',
'home_cs',
'home_gidp',
'home_ci',
'home_lob',
'home_pitchers_used',
'home_individual_er',
'home_team_er',
'home_wp',
'home_bk',
'home_po',
'home_assists',
'home_e',
'home_passed_balls',
'home_double_plays',
'home_triple_plays',

'hp_ump_id',
'hp_ump_name',
'1b_ump_id',
'1b_ump_name',
'2b_ump_id',
'2b_ump_name',
'3b_ump_id',
'3b_ump_name',
'lf_ump_id',
'lf_ump_name',
'rf_ump_id',
'rf_ump_name',
'visit_manager_id',
'visit_manager_name',
'home_manager_id',
'home_manager_name',
'winning_pitcher_id',
'winning_pitcher_name',
'losing_pitcher_id',
'losing_pitcher_name',
'saving_pitcher_id',
'saving_pitcher_name',
'winning_rbi_batter_id',
'winning_rbi_batter_name',
'visit_sp_id',
'visit_sp_name',
'home_sp_id',
'home_sp_name',
'visit_player_1_id',
'visit_player_1_name',
'visit_player_1_pos',
'visit_player_2_id',
'visit_player_2_name',
'visit_player_2_pos',
'visit_player_3_id',
'visit_player_3_name',
'visit_player_3_pos',
'visit_player_4_id',
'visit_player_4_name',
'visit_player_4_pos',
'visit_player_5_id',
'visit_player_5_name',
'visit_player_5_pos',
'visit_player_6_id',
'visit_player_6_name',
'visit_player_6_pos',
'visit_player_7_id',

```

'visit_player_7_name',
'visit_player_7_pos',
'visit_player_8_id',
'visit_player_8_name',
'visit_player_8_pos',
'visit_player_9_id',
'visit_player_9_name',
'visit_player_9_pos',
'home_player_1_id',
'home_player_1_name',
'home_player_1_pos',
'home_player_2_id',
'home_player_2_name',
'home_player_2_pos',
'home_player_3_id',
'home_player_3_name',
'home_player_3_pos',
'home_player_4_id',
'home_player_4_name',
'home_player_4_pos',
'home_player_5_id',
'home_player_5_name',
'home_player_5_pos',
'home_player_6_id',
'home_player_6_name',
'home_player_6_pos',
'home_player_7_id',
'home_player_7_name',
'home_player_7_pos',
'home_player_8_id',
'home_player_8_name',
'home_player_8_pos',
'home_player_9_id',
'home_player_9_name',
'home_player_9_pos',
'additional_info',
'acquisition_info']

```

[247]: g1

```

[247]:          0
0      visit_team
1      home_team
2      visit_score
3      home_score
4      game_length_outs
5      day_night

```

```
6         park_id
7     visit_manager_id
8     home_manager_id
9     winning_pitcher_id
10    losing_pitcher_id
11    saving_pitcher_id
12        visit_sp_id
13        home_sp_id
14    visit_player_1_id
15    visit_player_2_id
16    visit_player_3_id
17    visit_player_4_id
18    visit_player_5_id
19    visit_player_6_id
20    visit_player_7_id
21    visit_player_8_id
22    visit_player_9_id
23    home_player_1_id
24    home_player_2_id
25    home_player_3_id
26    home_player_4_id
27    home_player_5_id
28    home_player_6_id
29    home_player_7_id
30    home_player_8_id
31    home_player_9_id
```

[]:

create_gamelog_tensors

May 10, 2020

```
[13]: import os
import pandas as pd
import numpy as np
from tensorflow.keras.models import load_model
from joblib import load
pd.options.mode.chained_assignment = None # default='warn'
```

```
[14]: bat = load_model('../core/models/model_batting.h5')
pitch = load_model('../core/models/model_pitching.h5')
bat_scaler = load('../core/models/batting_scaler.save')
pitch_scaler = load('../core/models/pitching_scaler.save')
```

```
[15]: gl = pd.read_csv('../core/data/retrosheet/gamelogs/GL2015.csv')
```

```
[16]: gl
```

```
[16]:
```

| | visit_team | home_team | visit_score | home_score | game_length_outs | \ |
|------|------------|-----------|-------------|------------|------------------|---|
| 0 | MIN | DET | 0 | 4 | 51 | |
| 1 | CLE | HOU | 0 | 2 | 51 | |
| 2 | CHW | KCR | 1 | 10 | 51 | |
| 3 | TOR | NYY | 6 | 1 | 54 | |
| 4 | TEX | OAK | 0 | 8 | 51 | |
| ... | ... | ... | ... | ... | ... | |
| 2423 | CHC | MIL | 3 | 1 | 54 | |
| 2424 | WSN | NYM | 0 | 1 | 51 | |
| 2425 | FLA | PHI | 2 | 7 | 51 | |
| 2426 | CIN | PIT | 0 | 4 | 51 | |
| 2427 | COL | SFG | 7 | 3 | 54 | |

| | night_game | park_id | visit_manager_id | home_manager_id | visit_sp_id | ... | \ |
|------|------------|---------|------------------|-----------------|-------------|-----|---|
| 0 | 0 | DET05 | molip001 | ausmb001 | hughp001 | ... | |
| 1 | 1 | HOU03 | frant001 | hinca001 | klubc001 | ... | |
| 2 | 0 | KAN06 | ventr001 | yoste001 | samaj001 | ... | |
| 3 | 0 | NYC21 | gibbj001 | giraj001 | hutcd001 | ... | |
| 4 | 1 | OAK01 | banij001 | melvb001 | gally001 | ... | |
| ... | ... | ... | ... | ... | ... | ... | |
| 2423 | 0 | MIL06 | maddj801 | counc001 | hared001 | ... | |

| | | | | | | |
|------|---|-------|----------|----------|----------|-----|
| 2424 | 0 | NYC20 | willm003 | collt801 | roart001 | ... |
| 2425 | 0 | PHI13 | jennd801 | mackp101 | conla001 | ... |
| 2426 | 0 | PIT08 | pricb801 | hurdc001 | smitj004 | ... |
| 2427 | 0 | SF003 | weisw001 | bochb002 | bergc001 | ... |

| | home_player_4_id | home_player_5_id | home_player_6_id | home_player_7_id | \ |
|------|------------------|------------------|------------------|------------------|---|
| 0 | martv001 | martj006 | cespy001 | castn001 | |
| 1 | gatte001 | cartc002 | castj006 | lowrj001 | |
| 2 | hosme001 | morak001 | gorda001 | riosa002 | |
| 3 | teixm001 | mccab002 | headc001 | rodra001 | |
| 4 | butlb003 | davii001 | lawrb002 | vogts001 | |
| ... | ... | ... | ... | ... | |
| 2423 | davik003 | santd002 | pereh001 | seguj002 | |
| 2424 | cespy001 | dudal001 | darnt001 | confm001 | |
| 2425 | ruf-d001 | fran-j004 | blana001 | krate001 | |
| 2426 | walkn001 | marts002 | alvap001 | cervf001 | |
| 2427 | poseb001 | parkj002 | willm008 | noonn001 | |

| | home_player_8_id | home_player_9_id | year | month | day | home_win |
|------|------------------|------------------|------|-------|-----|----------|
| 0 | avila001 | iglej001 | 2015 | 4 | 6 | 1 |
| 1 | rasmc001 | marij002 | 2015 | 4 | 6 | 1 |
| 2 | peres002 | infao001 | 2015 | 4 | 6 | 1 |
| 3 | drews001 | gregd001 | 2015 | 4 | 6 | 0 |
| 4 | semim001 | sogae001 | 2015 | 4 | 6 | 1 |
| ... | ... | ... | ... | ... | ... | ... |
| 2423 | maldm001 | lopej004 | 2015 | 10 | 4 | 0 |
| 2424 | tejar001 | degrj001 | 2015 | 10 | 4 | 1 |
| 2425 | ruppc001 | buchd001 | 2015 | 10 | 4 | 1 |
| 2426 | mercj002 | happj001 | 2015 | 10 | 4 | 1 |
| 2427 | willj005 | cainm001 | 2015 | 10 | 4 | 0 |

[2428 rows x 33 columns]

```
[17]: columns = {
        'batting': [],
        'pitching': []
    }
```

```
[18]: batters = pd.read_csv('../core/output/batters.csv')
batter_years = pd.read_csv('../core/output/batting.csv')
batters_not_counted = list(batter_years[~batter_years['retroID']
                                     .isin(batters['retroID'])]['retroID'].
    →values)
pitchers = pd.read_csv('../core/output/pitchers.csv')
pitcher_years = pd.read_csv('../core/output/pitching.csv')
bat_scaler = load('../core/models/batting_scaler.save')
pitch_scaler = load('../core/models/pitching_scaler.save')
```

```

scalers = {
    'batting': bat_scaler,
    'pitching': pitch_scaler
}
career_features = {
    'batting': [
        'G', 'AB', 'PA', 'R', 'H', '1B', '2B', '3B',
        'HR', 'RBI', 'SB', 'CS', 'BB', 'SO', 'IBB',
        'HBP', 'SH', 'SF', 'GIDP'
    ],
    'pitching': [
        'CG', 'SHO', 'H', 'ER', 'HR', 'BB', 'SO',
        'BAOpp', 'ERA', 'IBB', 'WP', 'HBP', 'BK',
        'BFP', 'GF', 'R', 'SH', 'SF', 'GIDP'
    ]
}
unwanted_features = {
    'batting': ['retroID', 'G', 'AB', '1B', 'RBI', 'wOBA', 'Batting'],
    'pitching': ['IPouts', 'BFP', 'R', 'Pitching']
}
players = {
    'batting': {
        'players': batters,
        'years': batter_years
    },
    'pitching': {
        'players': pitchers,
        'years': pitcher_years
    }
}

```

```

[19]: def to_tensor_input(scaler, player, label):
    scalars[label] = scaler
    return scaler.transform(player.values.reshape(-1, player.shape[0]))[0]

def convert_single_player(retro_id, year, player_type_label):
    scaler = scalars[player_type_label]
    if retro_id in batters_not_counted:
        return np.zeros(shape=(1, 30))
    player_table = players[player_type_label]['players']
    player_so_far_table = players[player_type_label]['years']
    player = player_table[player_table['retroID'] == retro_id]
    player_so_far = player_so_far_table[(player_so_far_table['retroID'] ==
→retro_id)
                                         & (player_so_far_table['yearID'] <=
→year)]

```

```

if not player.size | player_so_far.size:
    print('Handled: {}'.format(retro_id))
    return np.zeros(shape=(1, 30))
player_so_far = player_so_far.groupby('retroID').sum()
features = career_features[player_type_label]
try:
    for column in player[features]:
        player.iloc[0][column] = player_so_far.iloc[0][column]
except:
    print(retro_id)
player_columns_to_drop = unwanted_features[player_type_label]
player = player.drop(columns=player_columns_to_drop)
if not len(list(columns[player_type_label])):
    columns[player_type_label] = player.columns
return to_tensor_input(scaler, player.T, player_type_label)

def get_batter_as_tensor_input(batter, year):
    scaler = scalars['batting']
    player = batters[batters['retroID'] == batter]
    player_so_far = batter_years[(batter_years['retroID'] == batter)
                                & (batter_years['yearID'] <= year)]
    player_so_far = player_so_far.groupby('retroID').sum()
    features = ['G', 'AB', 'PA', 'R', 'H', '1B', '2B', '3B',
                'HR', 'RBI', 'SB', 'CS', 'BB', 'SO', 'IBB',
                'HBP', 'SH', 'SF', 'GIDP']
    for column in player[features]:
        player.iloc[0][column] = player_so_far.iloc[0][column]
    player_columns_to_drop = ['retroID', 'wOBA', 'Batting']
    player = player.drop(columns=player_columns_to_drop)
    return to_tensor_input(scaler, player, 'batting')

```

```
[20]: convert_single_player('bettm001', 2015, 'batting')
```

```
[20]: array([0.42623, 0.3, 0., 0., 0.,
            0., 1., 0., 0., 0.,
            0., 0.16129032, 0.2288002, 0.2671024, 0.22673872,
            0.30697051, 0.13612565, 0.1824147, 0.08961593, 0.07462687,
            0.14503518, 0.17866769, 0.03633721, 0.06666667, 0.01486989,
            0.25, 0.10379747, 0., 0.21133094, 0.26473988])
```

```
[21]: gl.iloc[43]
```

```
[21]: visit_team      SFG
home_team           SDP
visit_score          1
home_score           0
```



```

game_length_outs      72
night_game            0
park_id              SAN02
visit_manager_id      bochb002
home_manager_id       blacb001
visit_sp_id           hudst001
home_sp_id            kenni001
visit_player_1_id     aokin001
visit_player_2_id     panij002
visit_player_3_id     pagaa001
visit_player_4_id     poseb001
visit_player_5_id     crawb001
visit_player_6_id     mcgec001
visit_player_7_id     blang001
visit_player_8_id     ariaj001
visit_player_9_id     hudst001
home_player_1_id      myerw001
home_player_2_id      norrd001
home_player_3_id      kempm001
home_player_4_id      uptoj001
home_player_5_id      middw001
home_player_6_id      alony001
home_player_7_id      gyorj001
home_player_8_id      amara001
home_player_9_id      kenni001
year                  2015
month                  4
day                    9
home_win               0
Name: 43, dtype: object

```

```
[22]: v1 = gl.iloc[0]['visit_player_1_id']
```

```
[23]: v1
```

```
[23]: 'santd001'
```

```
[24]: visit_id = []
      home_id = []
```

```
[25]: for i in range(1, 10):
      visit_id.append(gl.iloc[43]['visit_player_{}_id'.format(i)])
      home_id.append(gl.iloc[43]['home_player_{}_id'.format(i)])
```

```
[26]: visit_id
```

```
[26]: ['aokin001',  
      'panij002',  
      'pagaa001',  
      'poseb001',  
      'crawb001',  
      'mcgec001',  
      'blang001',  
      'ariaj001',  
      'hudst001']
```

```
[27]: gl.iloc[0]['year']
```

```
[27]: 2015
```

```
[28]: visit = []  
home = []  
year = gl.iloc[43]['year']  
for index in range(0, 9):  
    vrid = visit_id[index]  
    # vpos = 'pitching' if vrid == gl.iloc[0]['visit_sp_id'] else 'batting'  
    vplayer = convert_single_player(vrid, year, 'batting')  
    visit.append(vplayer)  
    hrid = home_id[index]  
    # hpos = 'pitching' if hrid == gl.iloc[0]['home_sp_id'] else 'batting'  
    hplayer = convert_single_player(hrid, year, 'batting')  
    home.append(hplayer)
```

```
[29]: visit[0].shape
```

```
[29]: (30,)
```

```
[30]: home[0].shape
```

```
[30]: (30,)
```

```
[31]: gl.columns
```

```
[31]: Index(['visit_team', 'home_team', 'visit_score', 'home_score',  
        'game_length_outs', 'night_game', 'park_id', 'visit_manager_id',  
        'home_manager_id', 'visit_sp_id', 'home_sp_id', 'visit_player_1_id',  
        'visit_player_2_id', 'visit_player_3_id', 'visit_player_4_id',  
        'visit_player_5_id', 'visit_player_6_id', 'visit_player_7_id',  
        'visit_player_8_id', 'visit_player_9_id', 'home_player_1_id',  
        'home_player_2_id', 'home_player_3_id', 'home_player_4_id',  
        'home_player_5_id', 'home_player_6_id', 'home_player_7_id',  
        'home_player_8_id', 'home_player_9_id', 'year', 'month', 'day',  
        'home_win'],
```

```
dtype='object')
```

```
[32]: batters = visit + home
```

```
[33]: dfb = pd.DataFrame(batters, columns=columns['batting'])
```

```
[34]: dfb
```

```
[34]:
```

| | weight | height | pos_1B | pos_2B | pos_3B | pos_C | pos_OF | pos_P | pos_SS | \ |
|----|----------|--------|--------|--------|--------|-------|--------|-------|--------|---|
| 0 | 0.426230 | 0.30 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | |
| 1 | 0.508197 | 0.50 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 2 | 0.508197 | 0.55 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | |
| 3 | 0.549180 | 0.50 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | |
| 4 | 0.618852 | 0.55 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | |
| 5 | 0.590164 | 0.50 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 6 | 0.454918 | 0.35 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | |
| 7 | 0.446721 | 0.50 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 8 | 0.405738 | 0.50 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 9 | 0.528689 | 0.60 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 10 | 0.651639 | 0.45 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | |
| 11 | 0.610656 | 0.65 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | |
| 12 | 0.569672 | 0.50 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | |
| 13 | 0.590164 | 0.60 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 14 | 0.631148 | 0.50 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 15 | 0.569672 | 0.35 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 16 | 0.344262 | 0.15 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 17 | 0.528689 | 0.45 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | |

| | bats_L | ... | BB | SO | IBB | HBP | SH | SF | \ |
|----|--------|-----|----------|----------|----------|----------|----------|----------|---|
| 0 | 1.0 | ... | 0.091478 | 0.099345 | 0.004360 | 0.168421 | 0.111524 | 0.117188 | |
| 1 | 1.0 | ... | 0.085614 | 0.097420 | 0.020349 | 0.073684 | 0.048327 | 0.203125 | |
| 2 | 0.0 | ... | 0.124707 | 0.245668 | 0.026163 | 0.014035 | 0.078067 | 0.265625 | |
| 3 | 0.0 | ... | 0.189210 | 0.244128 | 0.090116 | 0.147368 | 0.003717 | 0.398438 | |
| 4 | 1.0 | ... | 0.155590 | 0.360801 | 0.091570 | 0.129825 | 0.022305 | 0.367188 | |
| 5 | 0.0 | ... | 0.097342 | 0.199076 | 0.020349 | 0.028070 | 0.000000 | 0.242188 | |
| 6 | 1.0 | ... | 0.141908 | 0.254524 | 0.030523 | 0.052632 | 0.096654 | 0.117188 | |
| 7 | 0.0 | ... | 0.014464 | 0.057759 | 0.010174 | 0.028070 | 0.044610 | 0.070312 | |
| 8 | 0.0 | ... | 0.010164 | 0.073161 | 0.000000 | 0.007018 | 0.249071 | 0.015625 | |
| 9 | 0.0 | ... | 0.122361 | 0.322680 | 0.020349 | 0.045614 | 0.003717 | 0.156250 | |
| 10 | 0.0 | ... | 0.076231 | 0.208317 | 0.014535 | 0.066667 | 0.007435 | 0.093750 | |
| 11 | 0.0 | ... | 0.196638 | 0.616095 | 0.098837 | 0.091228 | 0.003717 | 0.562500 | |
| 12 | 0.0 | ... | 0.284988 | 0.692337 | 0.071221 | 0.235088 | 0.011152 | 0.429688 | |
| 13 | 0.0 | ... | 0.025020 | 0.125144 | 0.005814 | 0.031579 | 0.003717 | 0.078125 | |
| 14 | 1.0 | ... | 0.143081 | 0.249519 | 0.042151 | 0.056140 | 0.003717 | 0.218750 | |
| 15 | 0.0 | ... | 0.091087 | 0.243358 | 0.007267 | 0.059649 | 0.000000 | 0.171875 | |
| 16 | 1.0 | ... | 0.042611 | 0.113593 | 0.018895 | 0.014035 | 0.085502 | 0.117188 | |
| 17 | 0.0 | ... | 0.012901 | 0.059299 | 0.000000 | 0.003509 | 0.156134 | 0.015625 | |

| | GIDP | NL | wRC+ | WAR |
|----|----------|-----|----------|----------|
| 0 | 0.124051 | 1.0 | 0.184353 | 0.105202 |
| 1 | 0.129114 | 1.0 | 0.177158 | 0.103468 |
| 2 | 0.136709 | 1.0 | 0.181655 | 0.157803 |
| 3 | 0.374684 | 1.0 | 0.205036 | 0.354335 |
| 4 | 0.235443 | 1.0 | 0.173561 | 0.172254 |
| 5 | 0.291139 | 1.0 | 0.171763 | 0.072832 |
| 6 | 0.088608 | 1.0 | 0.173561 | 0.101734 |
| 7 | 0.060759 | 1.0 | 0.158273 | 0.051445 |
| 8 | 0.030380 | 1.0 | 0.095324 | 0.058382 |
| 9 | 0.179747 | 1.0 | 0.186151 | 0.104624 |
| 10 | 0.106329 | 1.0 | 0.171763 | 0.105202 |
| 11 | 0.453165 | 1.0 | 0.198741 | 0.206358 |
| 12 | 0.313924 | 1.0 | 0.197842 | 0.262428 |
| 13 | 0.081013 | 1.0 | 0.158273 | 0.055491 |
| 14 | 0.237975 | 1.0 | 0.181655 | 0.080925 |
| 15 | 0.184810 | 1.0 | 0.179856 | 0.100000 |
| 16 | 0.081013 | 1.0 | 0.147482 | 0.036416 |
| 17 | 0.007595 | 1.0 | 0.097122 | 0.055491 |

[18 rows x 30 columns]

```
[35]: btensor = [dfb, gl.iloc[43]['home_win']]
```

```
[36]: # btensor
```

```
[37]: gl.shape
```

```
[37]: (2428, 33)
```

```
[38]: gl.shape[0]
```

```
[38]: 2428
```

```
[39]: players['batting']['players']['retroID'].str.contains('aardd001').sum() == 1
```

```
[39]: True
```

Modular script to handle all gamelogs

```
[40]: cols = list(columns['batting'].values) + ['Result']
for year in range(1919, 2020):
    df = pd.DataFrame()
    print('{}'.format(year))
    # gl = pd.read_csv('../core/data/retrosheet/gamelogs/GL{}.csv'.format(year))
    # for index in range(0, gl.shape[0]):
```

```

#         visit_id = []
#         home_id = []
#         for i in range(1, 10):
#             visit_id.append(gl.iloc[index]['visit_player_{}_id'.format(i)])
#             home_id.append(gl.iloc[index]['home_player_{}_id'.format(i)])
#         visit = []
#         home = []
#         for i in range(0, 9):
#             vrid = visit_id[i]
#             vplayer = convert_single_player(vrid, year, 'batting')
#             visit.append(vplayer)
#             hrid = home_id[i]
#             hplayer = convert_single_player(hrid, year, 'batting')
#             home.append(hplayer)
#         batters = list(np.append(np.array(visit + home).flatten(), gl.
→iloc[index]['home_win']))
#         try:
#             bat_df = pd.DataFrame(batters)
#         except:
#             print('{0}\n{1}'.format(vrid))
#             df = df.append(bat_df.T)

#         if not os.path.exists('../core/tensors/games/'):
#             os.mkdir('../core/tensors/games/')
#         df.to_csv('../core/tensors/games/{0}.csv'.format(str(year)), index=False,
→header=None)

```

1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938

1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986

1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019

[]: